


Transfer learning



Dario Garcia Gasulla dario.garcia@bsc.es
Armand Vilalta

Transfer Learning

The idea



Your first driving lesson

Your first driving lesson

- Training a deep neural network typically requires **thousands** if not **millions** of training **examples**.
- **Each example** is passed through the network, an **output** is obtained, and, if the prediction is **wrong**, the weights are **adjusted**.

Your first driving lesson

- Training a deep neural network typically requires **thousands** if not **millions** of training **examples**.
- **Each example** is passed through the network, an **output** is obtained, and, if the prediction is **wrong**, the weights are **adjusted**.
- Imagine **learning to drive** a car this way...

Your first driving lesson

Learning like a
Deep Neural Network



Your first driving lesson

- Training a deep neural network typically requires **thousands** if not **millions** of training examples.
- **Each example** is passed through the network, an **output** is obtained, and, if the prediction is **wrong**, the weights are **adjusted**.
- Imagine learning to drive a car this way...

We naturally reuse what we previously learnt to be able to solve a new task.

Why do we need transfer learning?



Image classification

- 1998 LeNet-5

Gradient-based learning applied to document recognition.
Yann LeCun, Léon Bottou, Yoshua Bengio, Patrick Haffner





Image classification

- 2012 AlexNet

ImageNet Classification with Deep Convolutional Neural Networks
Alex Krizhevsky, Ilya Sutskever and Geoffrey Hinton



1998 LeNet-5



2012 AlexNet



2014 VGG19



2014 GoogLeNet




2015 Inception-V3



2015 ResNet-56



ImageNet classification results



From image-net.org

At what price?

- Data available
 - 1000 images per class
- Computational cost
 - Specific hardware
 - Energy cost
- Human effort
 - Highly skilled professionals
 - Architecture design
 - Hyper-parameter fine tuning

At what price?

- Data available
 - 1,000 images per class
- Computational cost
 - Specific hardware
 - Energy cost
- Human effort
 - Highly skilled professionals
 - Architecture design
 - Hyper-parameter fine tuning

We can't do that for every single problem!!


At what price?

- Data available
 - 1,000 images per class
- Computational cost
 - Specific hardware
 - Energy cost
- Human effort
 - Highly skilled professionals
 - Architecture design
 - Hyper-parameter fine tuning

We can't do that for every single problem!!

→ Transfer Learning to the rescue

What is transfer learning?



What is learning about?



What is learning about?

Train set



What is learning about?

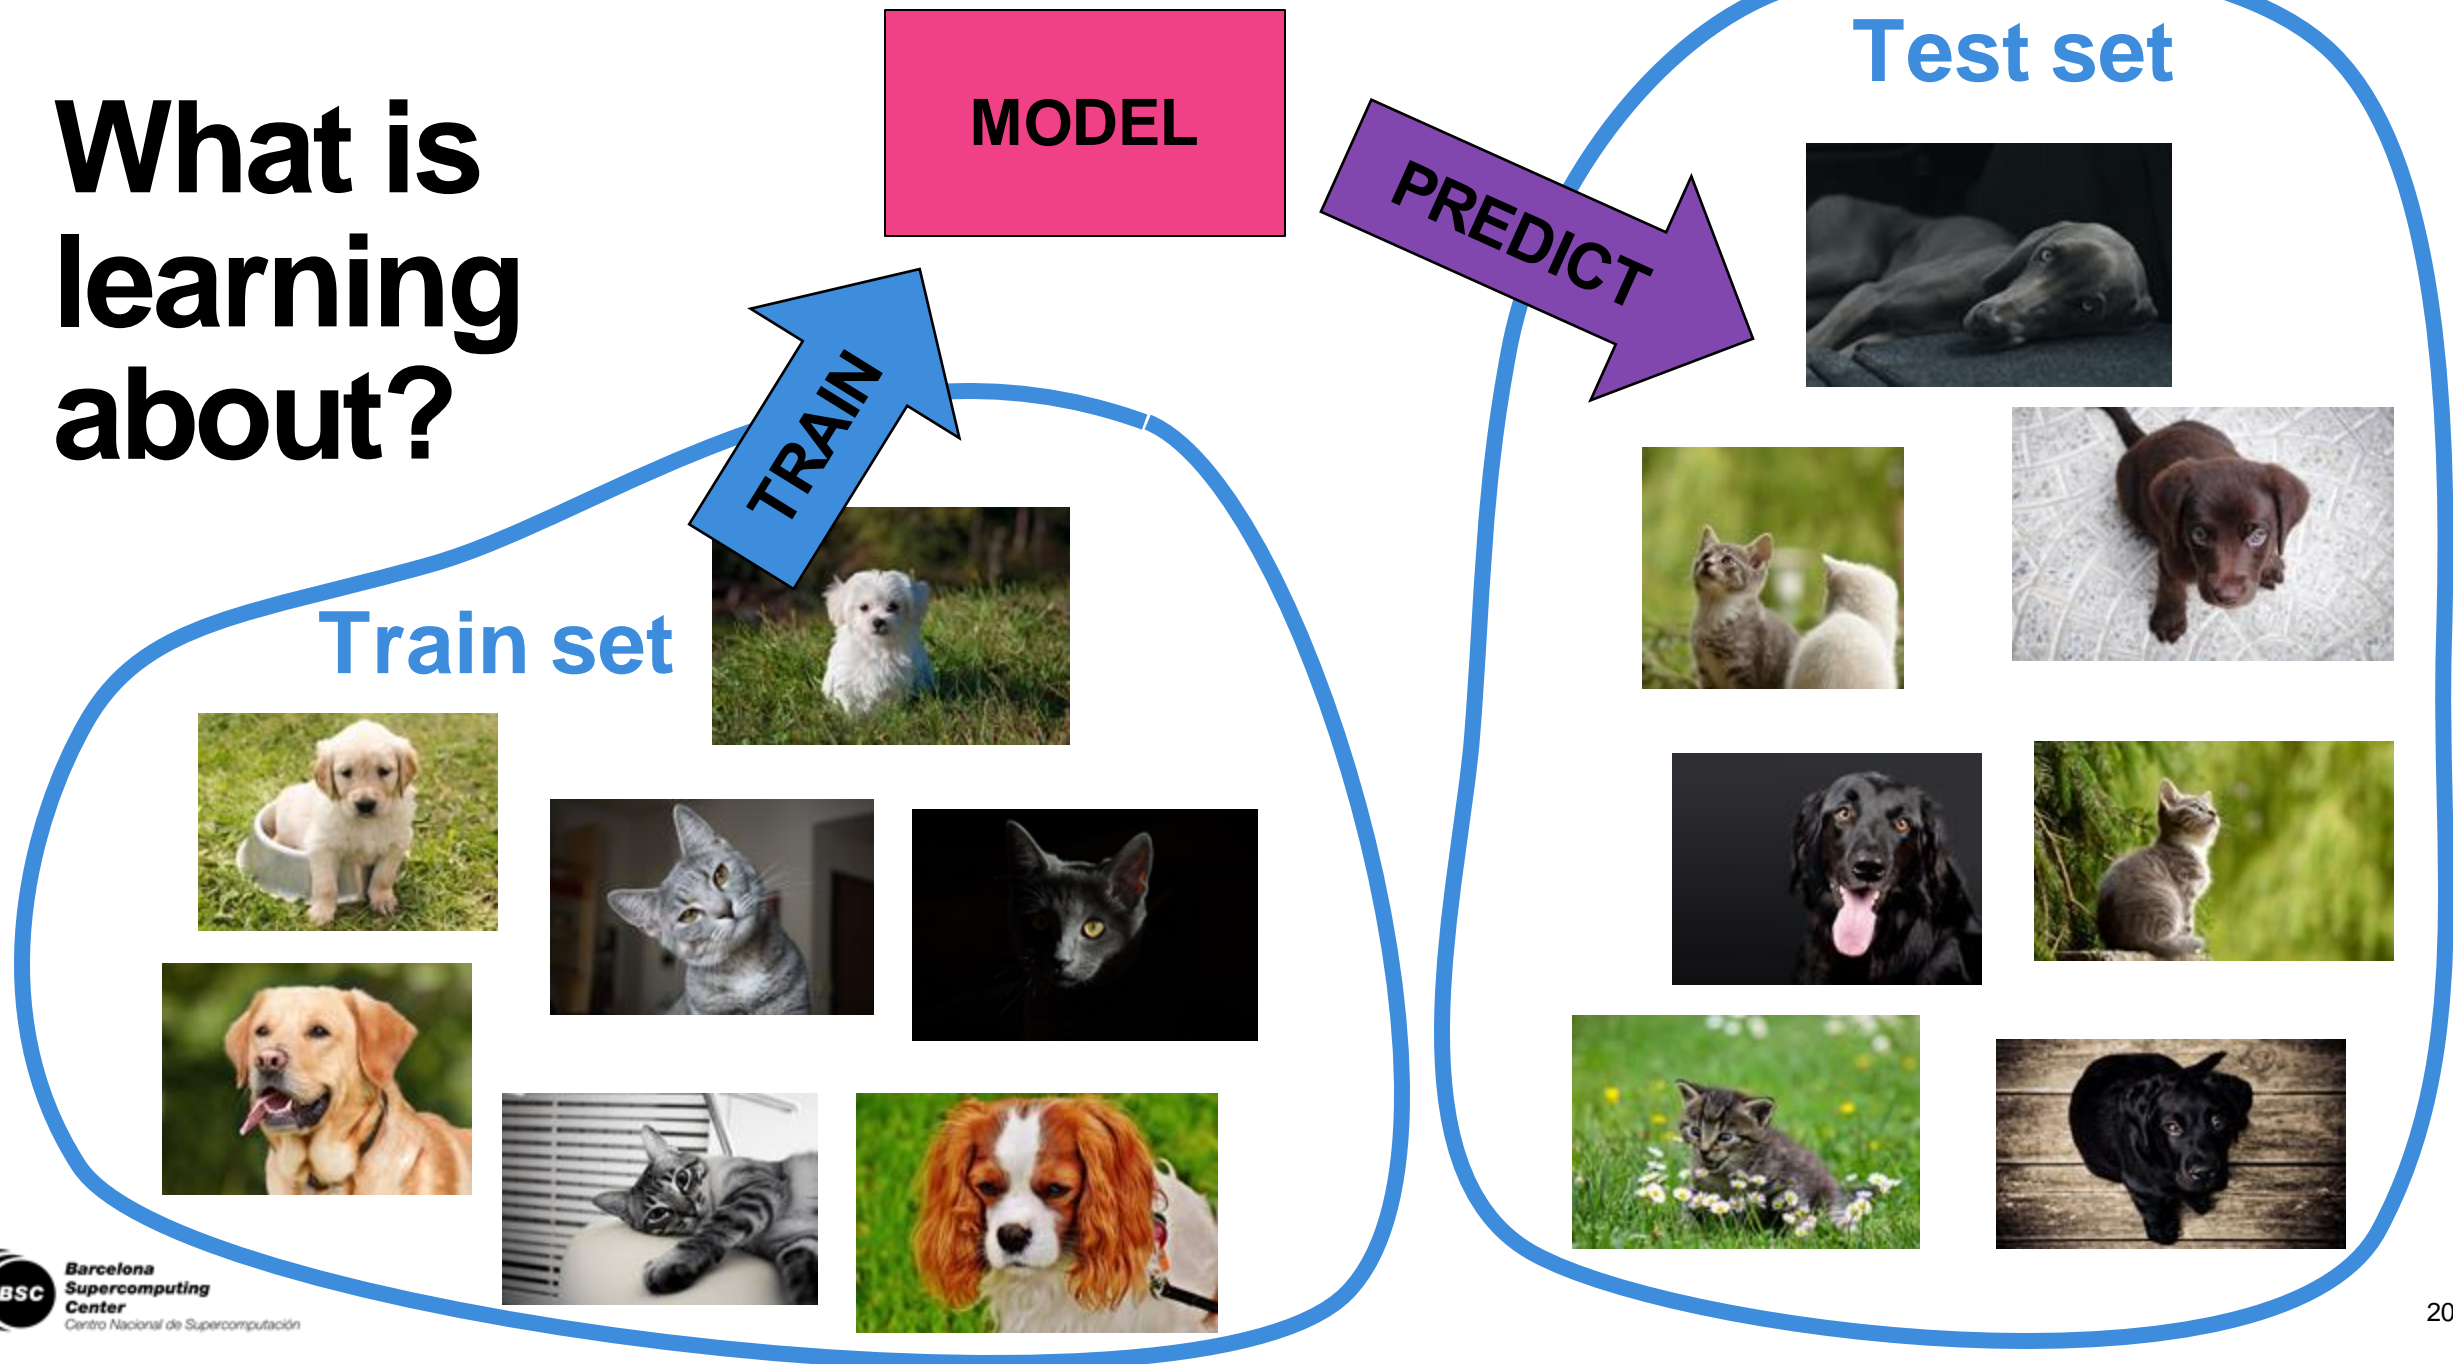
Train set



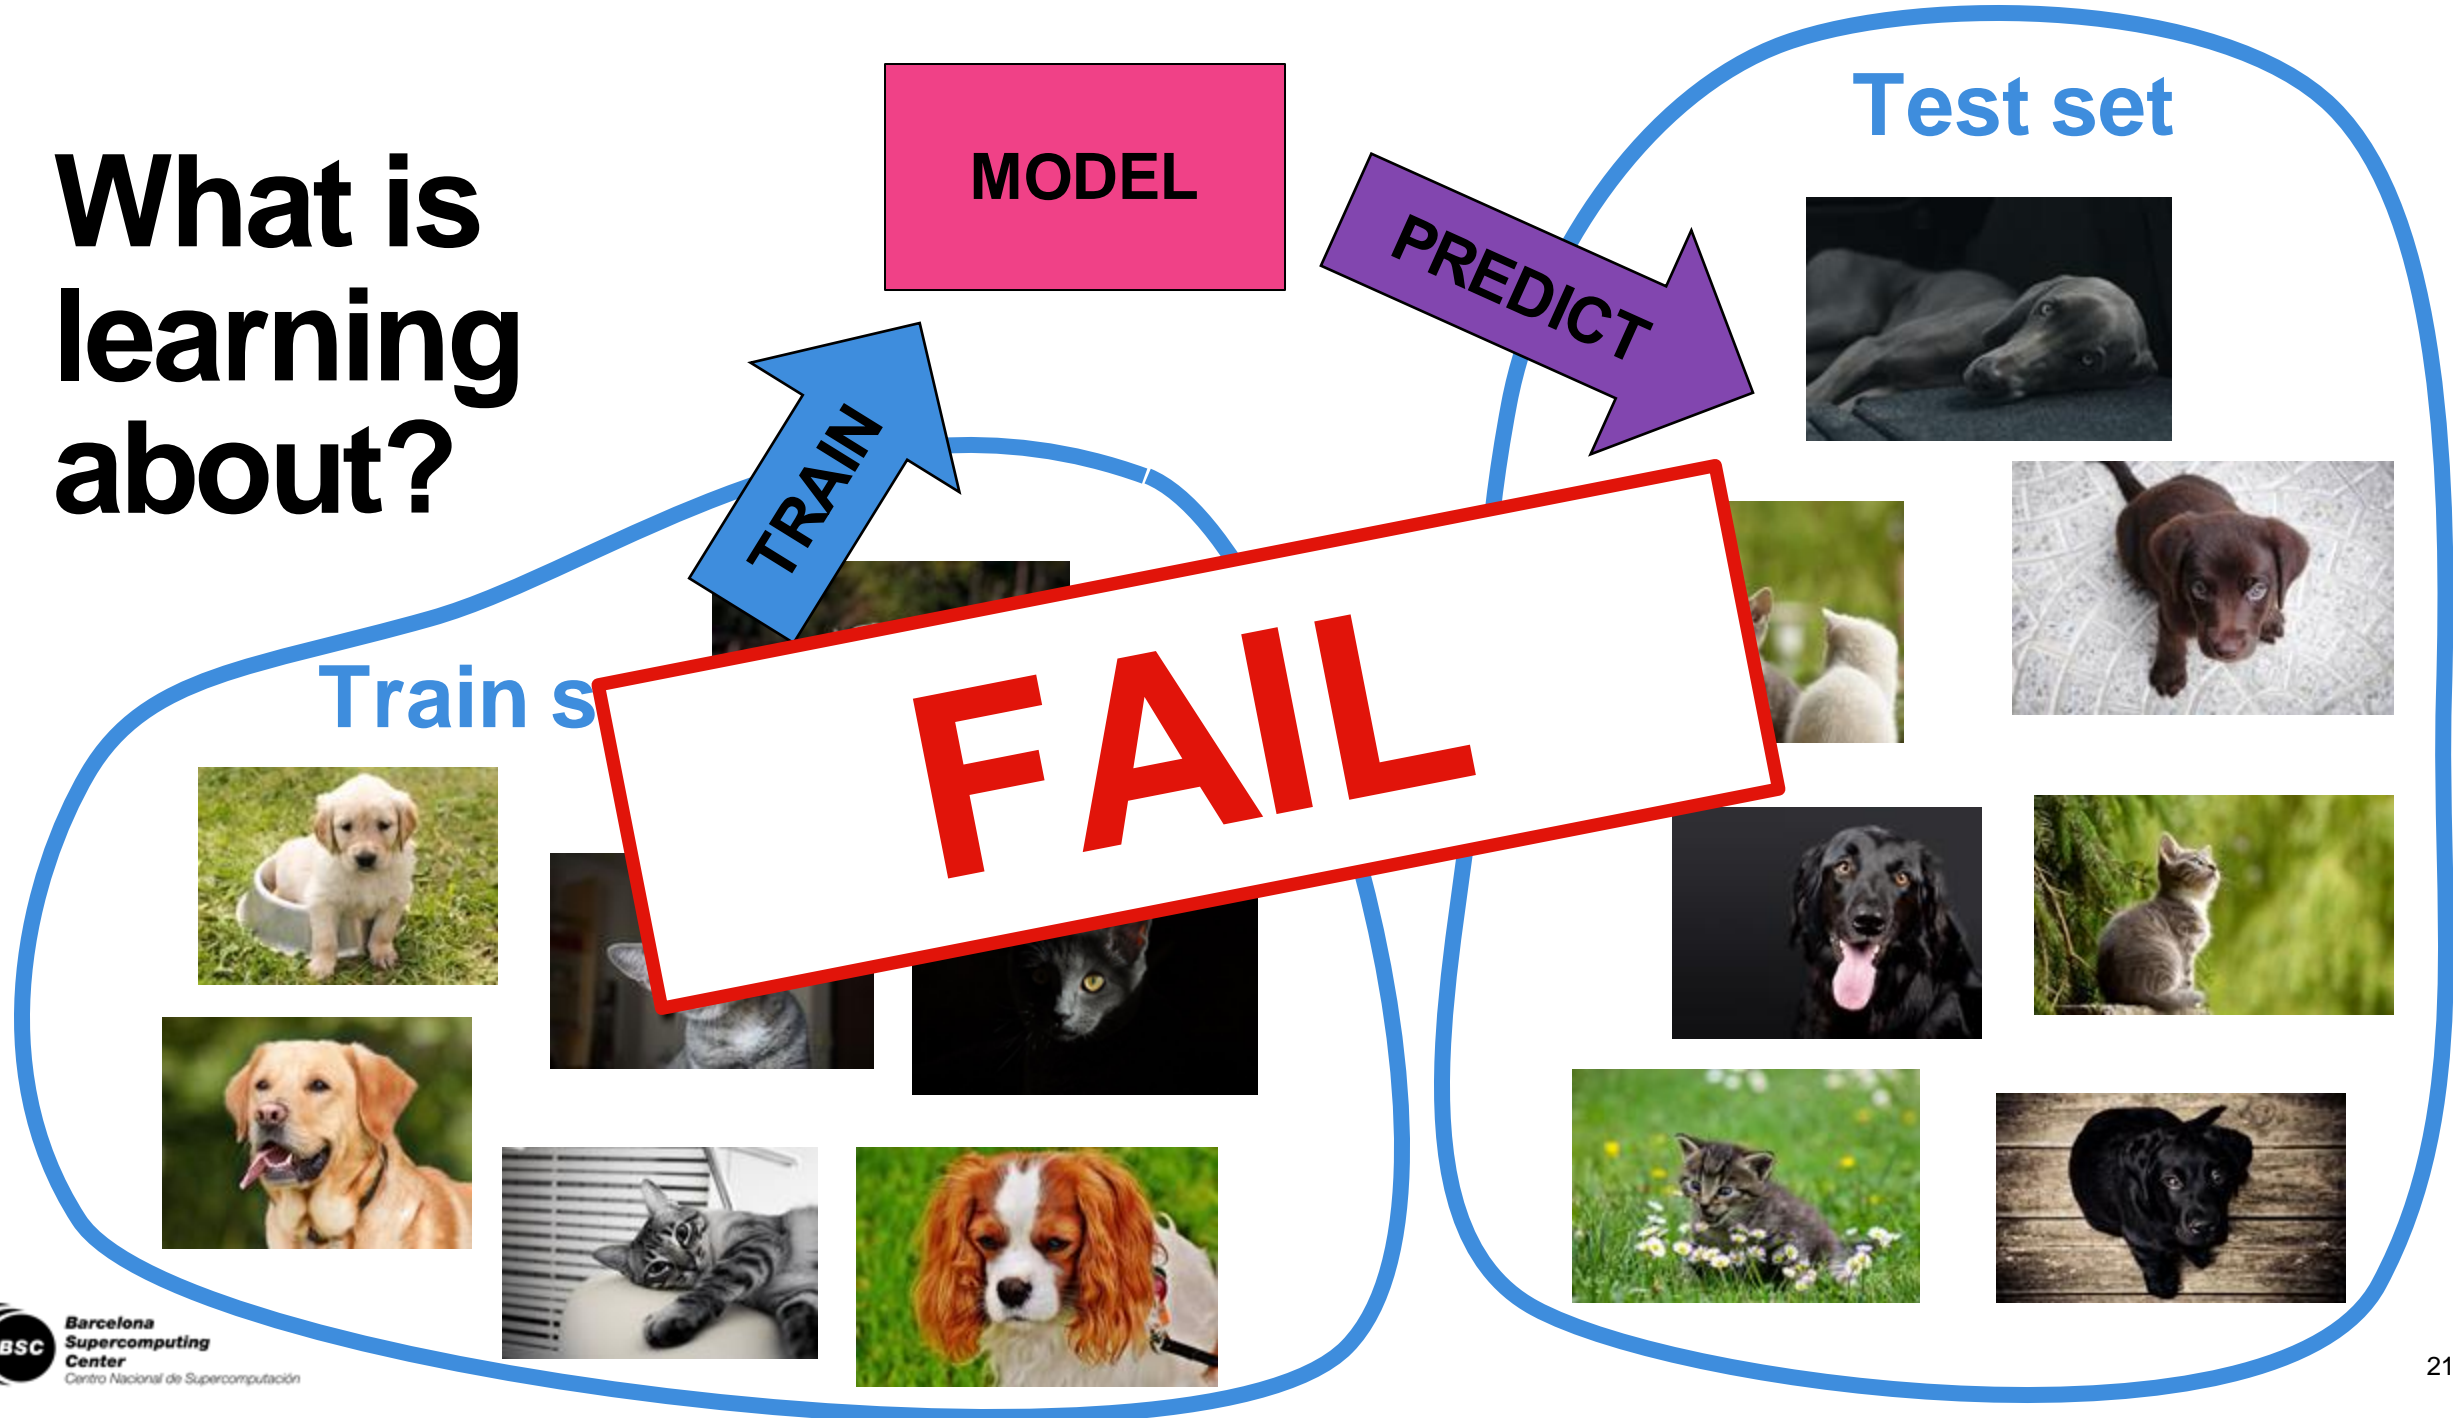
Test set



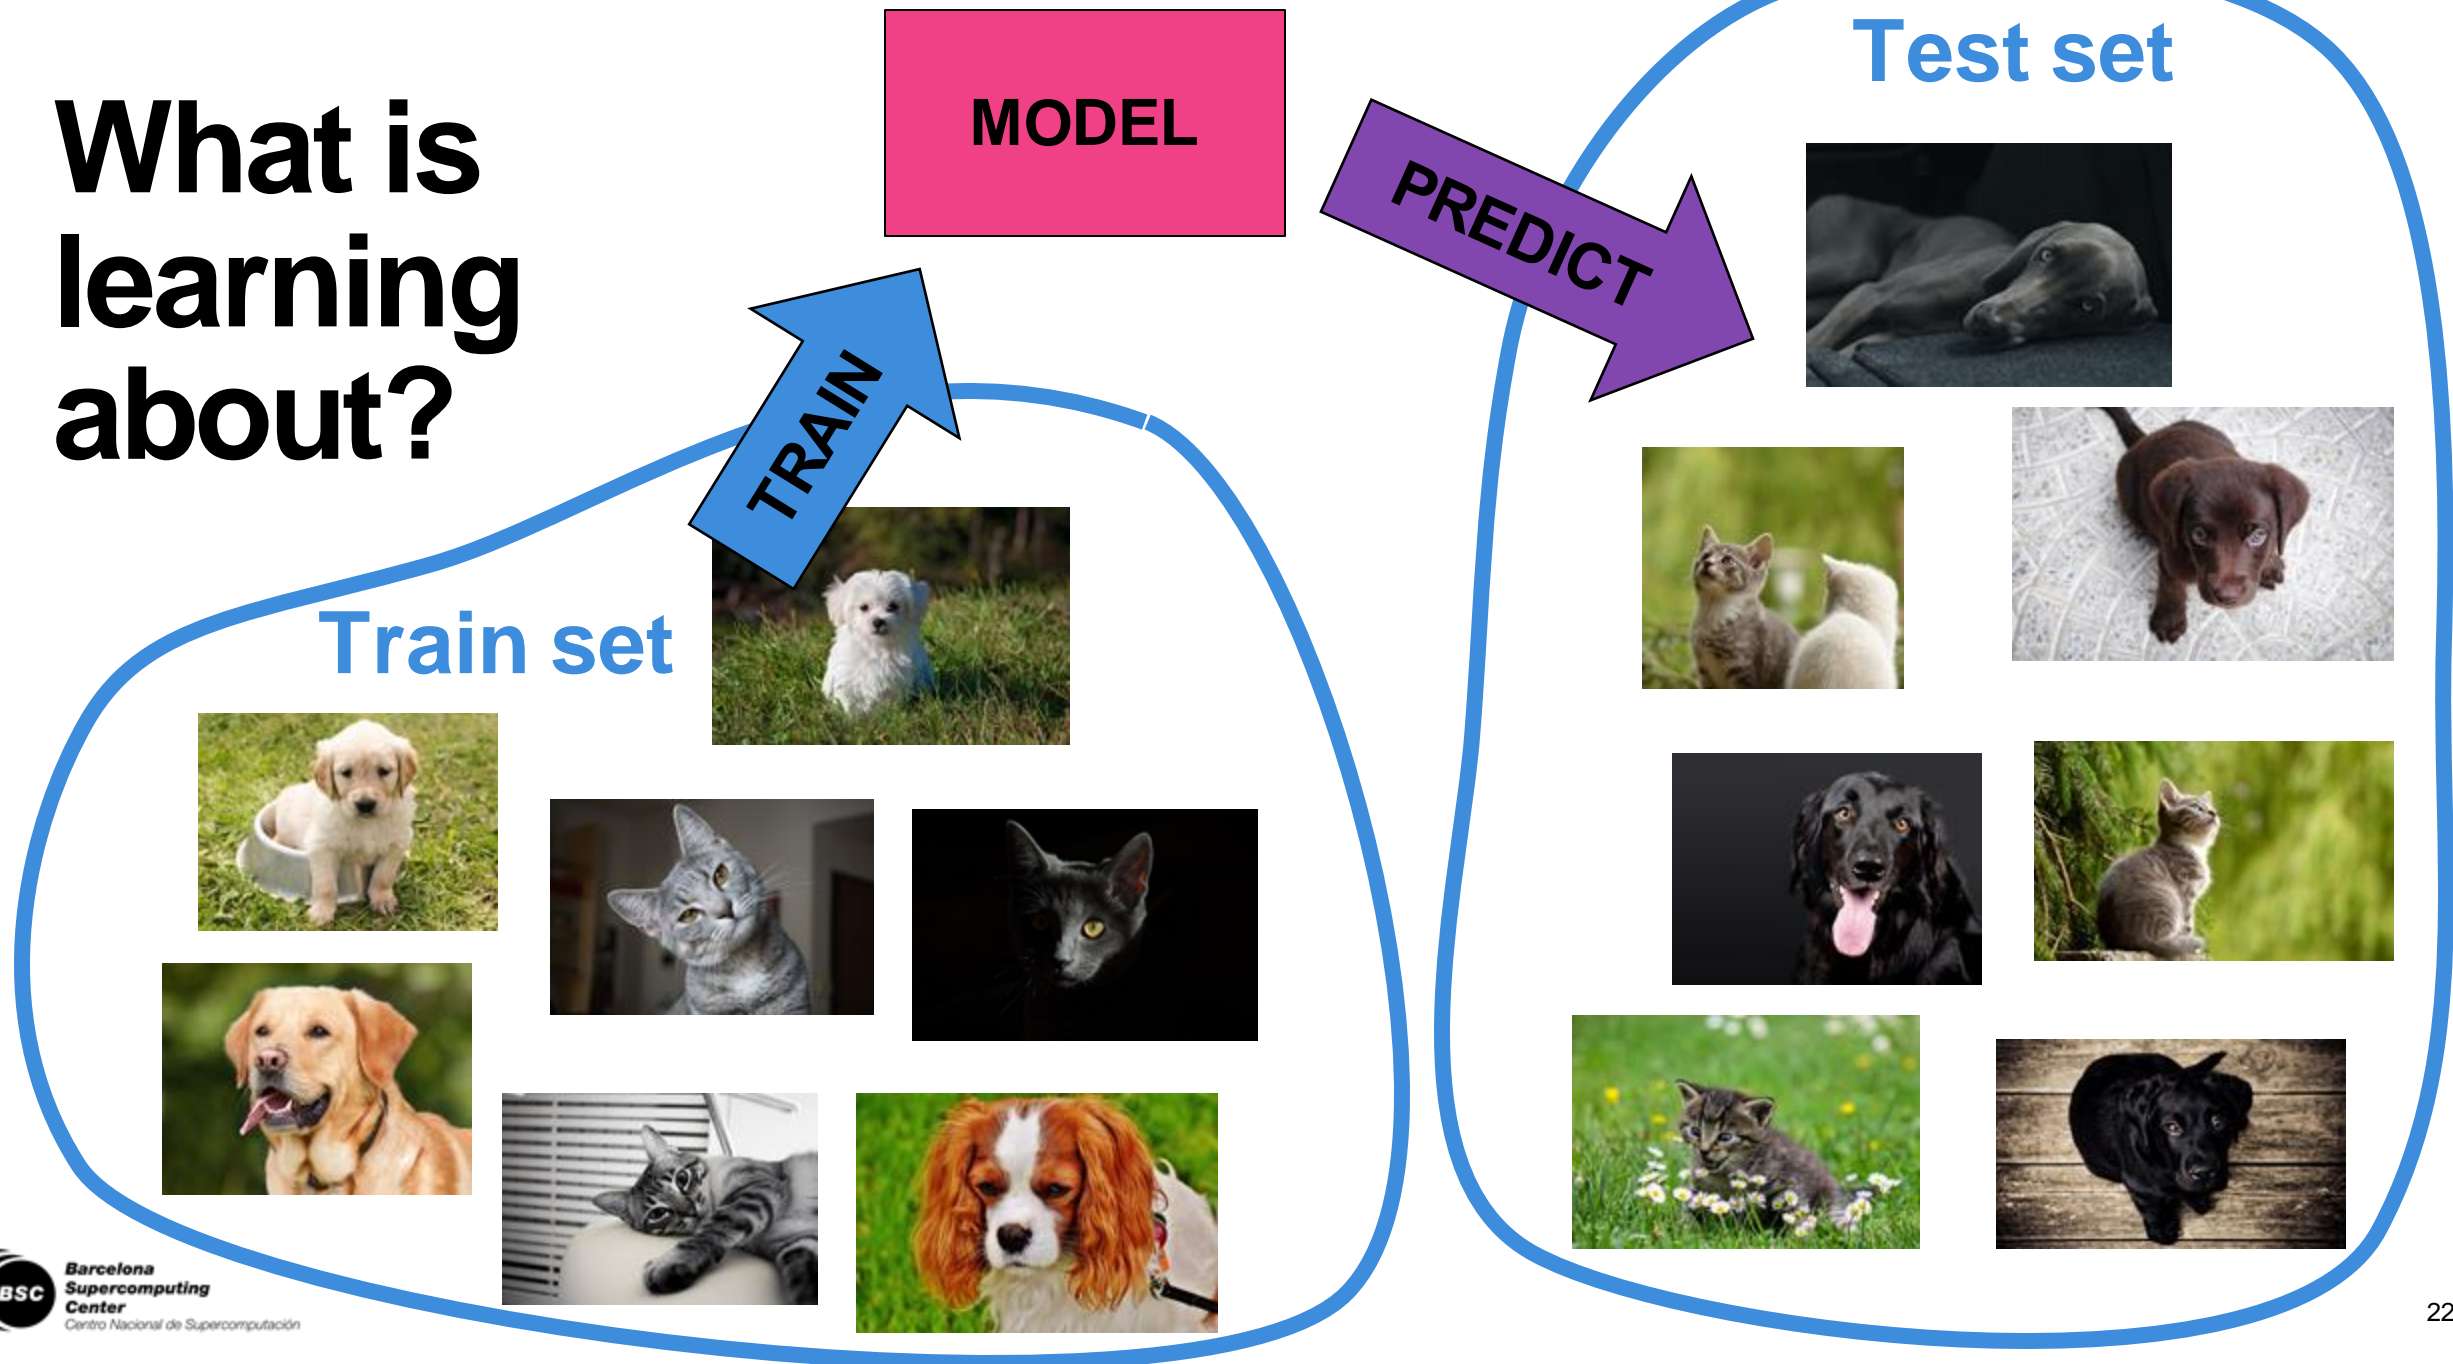
What is learning about?



What is learning about?



What is learning about?



What is learning about?

Train set



Test set



What is learning about?

GREEN
BACKGROUND

Train set



Test set



What is learning about?

GREEN
BACKGROUND

Train set




Test set



What is learning about?

GREEN
BACKGROUND



DOG


Test set

Train set



What is learning about?

GREEN
BACKGROUND



DOG

Test set

Train set



What is learning about?

GREEN BACKGROUND

DOG

Test set

Train set

Train and test sets have different conditional probability distributions




What is learning about?

Solution?

Train set



Test set




What is learning about?

Solution? Randomize

Train set



Test set



What is learning about?

Solution?
Randomize

Test set

Train set

Ensure that train and test sets have **similar** conditional probability distributions



What is learning about?

Solution?
Randomize

Test set

Ensure that train and test sets have **similar** conditional probability distributions

Train set



THEY WILL NEVER BE
EXACTLY EQUAL



What is learning about?

Train set



Test set



What is learning about?

MODEL

TRAIN

Train set



PREDICT

Test set



What is learning about?

MODEL

TRAIN

Train set



PREDICT

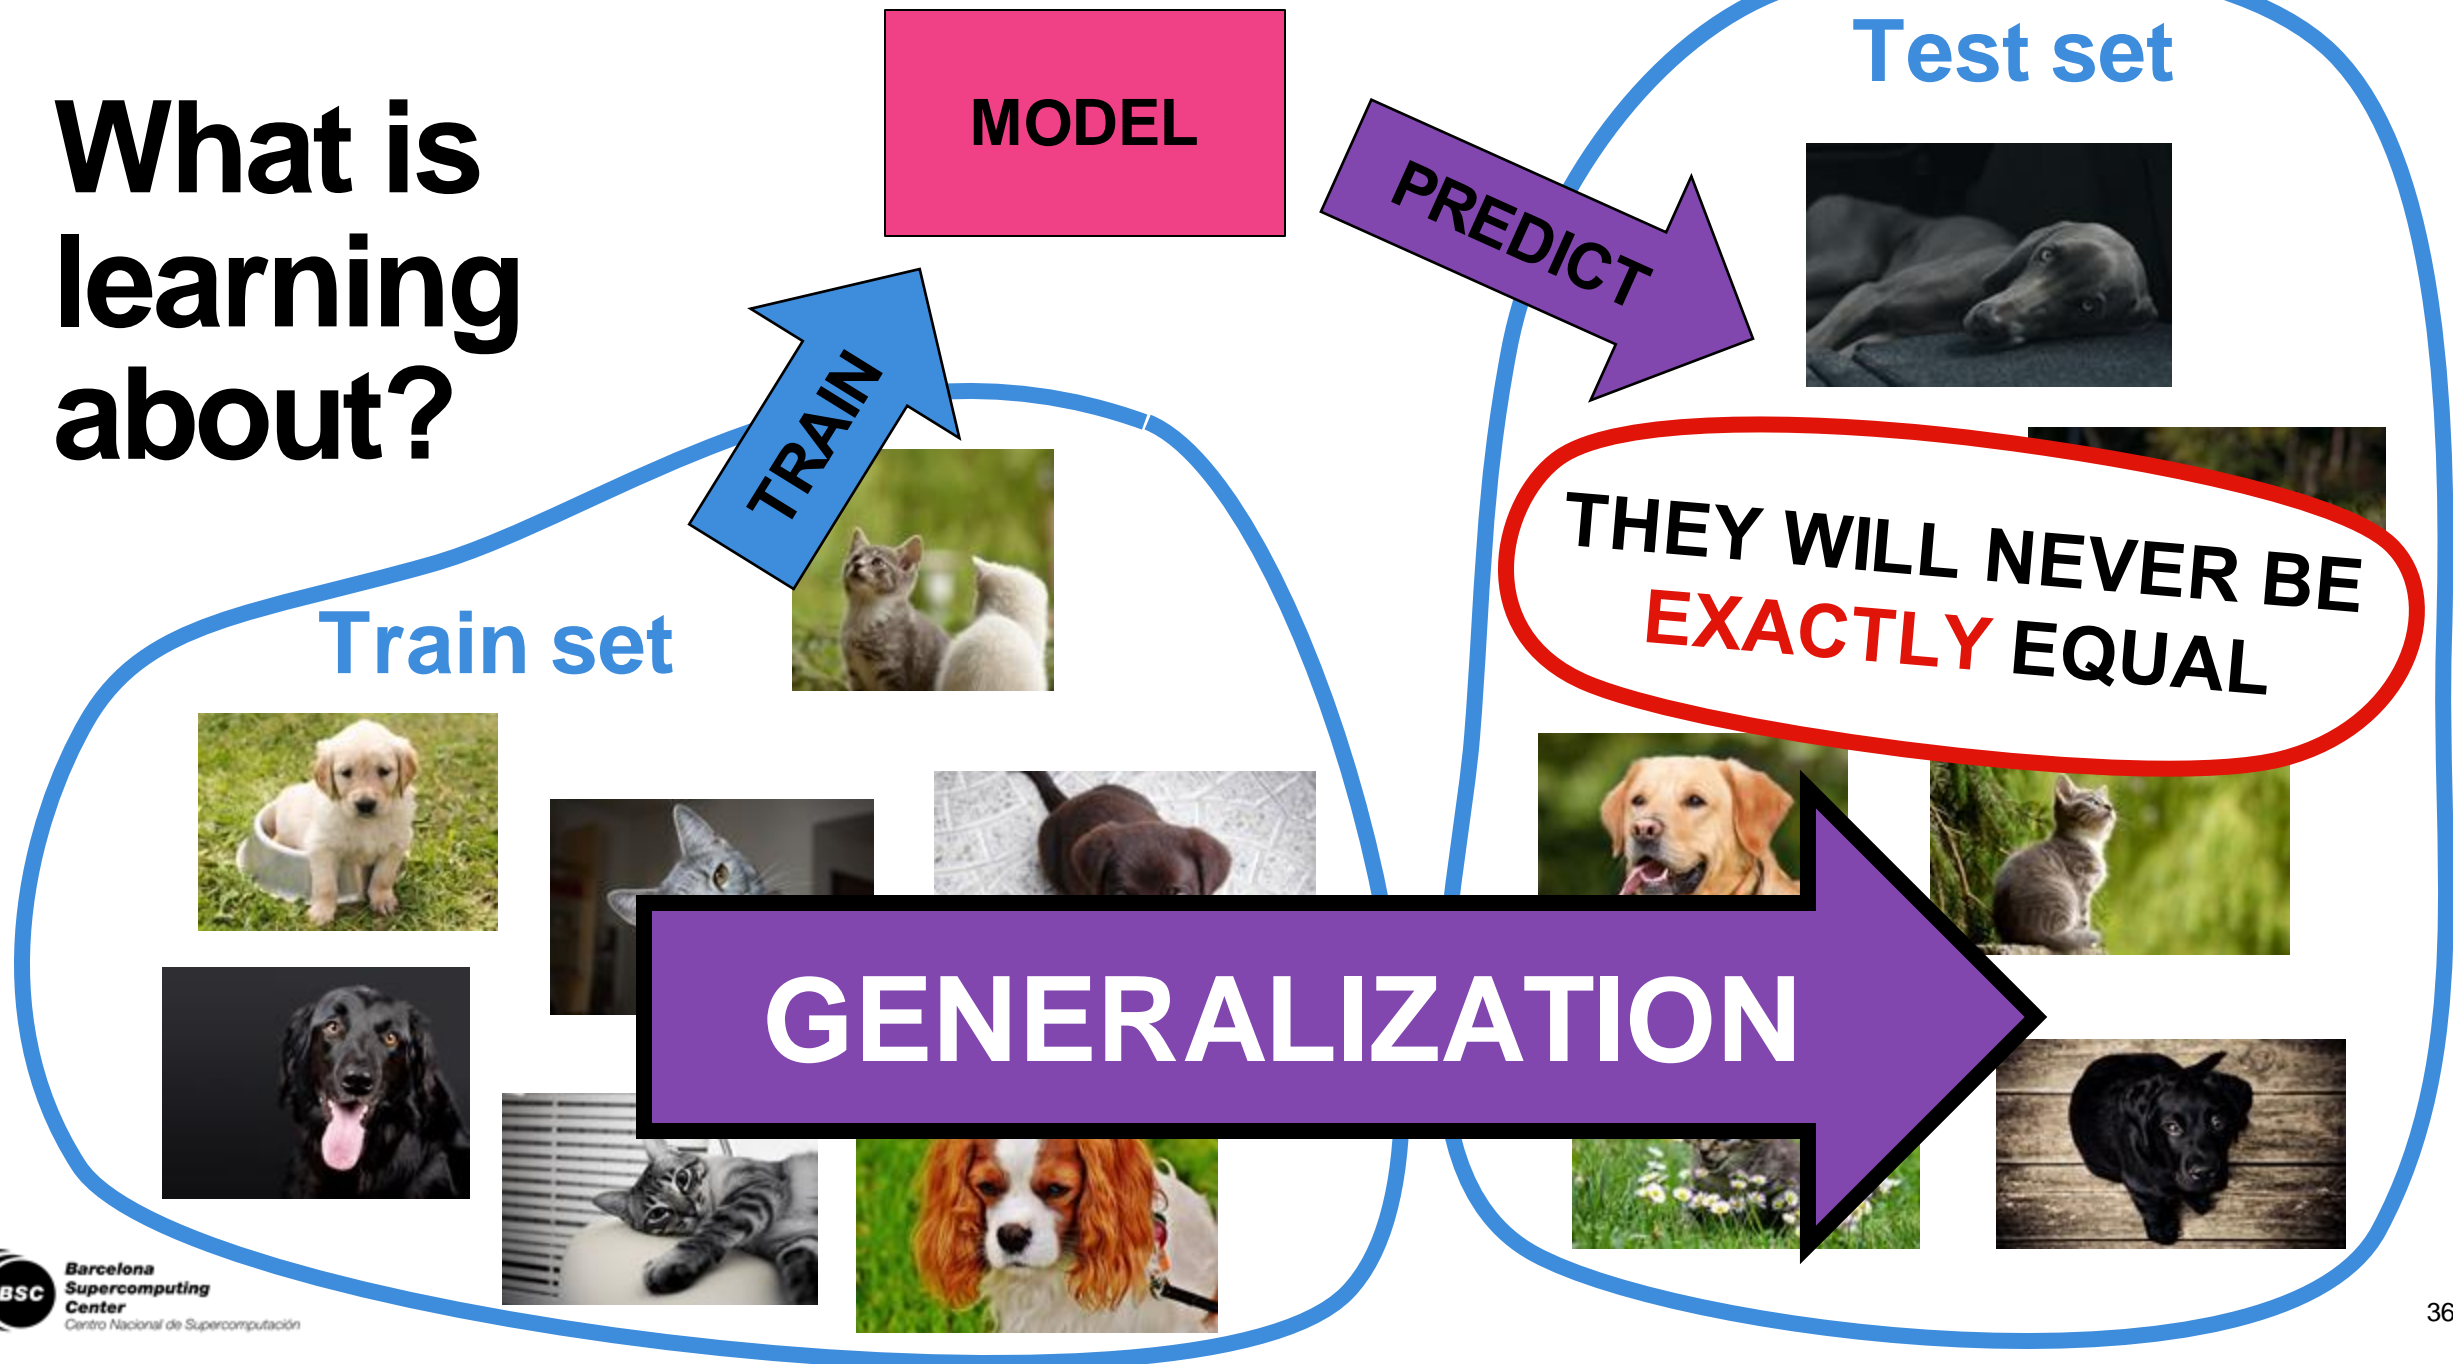
Test set



THEY WILL NEVER BE
EXACTLY EQUAL



What is learning about?




What is learning about?

Train set



GENERALIZATION

Test set



What is learning about?

Train a **Machine Learning Model** on a **train set** with the hope that what has been learnt will be useful to solve the **same task** on a **different test set**.

Train set



GENERALIZATION

Test set



What is learning about?

Train a **Machine Learning Model** on a **train set** with the hope that what has been learnt will be useful to solve the **same task** on a **different test set**.

Train set



GENERALIZATION

Test set



Train and test sets must have **similar** conditional probability distributions

THEY WILL NEVER BE EXACTLY EQUAL

What is **transfer** learning about?

Train set



Test set



What is **transfer** learning about?

Train set



GENERALIZATION

Test set



What is transfer learning about?

Train a **Machine Learning Model** on a **train set** with the hope that what has been learnt will be useful to solve a **different task** on a **different test set**.

Train set



GENERALIZATION

Test set



What is transfer learning about?

Train a **Machine Learning Model** on a **train set** with the hope that what has been learnt will be useful to solve a **different task** on a **different test set**.

Train set



GENERALIZATION


Train and test sets are drawn from a **not so similar** underlying probability distribution.

Test set



Formalizing transfer learning

Pan, Sinno Jialin, and Qiang Yang.
"A survey on transfer learning."
*IEEE Transactions on knowledge
and data engineering* 22.10 (2010):
1345-1359.



Formalizing **transfer** learning


Domain:

Task:

Formalizing **transfer** learning

Domain: $\mathcal{D} = \{\mathcal{X}, P(X)\}$

- A feature space \mathcal{X}



\neq

“The Elgar Concert
Hall at the University
of Birmingham for
our third conference”

→ Bag of words

→ Content vector

- A marginal probability distribution $P(X)$, where $X = \{x_1, \dots, x_n\} \in \mathcal{X}$



\neq




Task:

Formalizing **transfer** learning

Domain: $\mathcal{D} = \{\mathcal{X}, P(X)\}$

- A feature space \mathcal{X}



\neq

"The Elgar Concert
Hall at the University
of Birmingham for
our third conference"

→ Bag of words

→ Content vector

- A marginal probability distribution $P(X)$, where $X = \{x_1, \dots, x_n\} \in \mathcal{X}$



\neq




Task: $\mathcal{T} = \{\mathcal{Y}, f(\cdot)\}$

- A label space \mathcal{Y}

CAT, DOG \neq LION, WOLF

- An objective predictive function $f(\cdot) \Leftrightarrow P(y|x)$



Formalizing **transfer** learning

Source

Target

Domain: $\mathcal{D} = \{\mathcal{X}, P(X)\}$

Task: $\mathcal{T} = \{y, f(\cdot)\}$


Formalizing **transfer** learning

Domain: $\mathcal{D} = \{\mathcal{X}, P(X)\}$


- A feature space \mathcal{X}
 - **The Same (different)**
- A marginal probability distribution $P(X)$
 - **Different**
 - **Similar**

Task: $\mathcal{T} = \{y, f(\cdot)\}$

Source



Target



Formalizing **transfer** learning


Domain: $\mathcal{D} = \{\mathcal{X}, P(X)\}$

- A feature space \mathcal{X}
 - **The Same (different)**
- A marginal probability distribution $P(X)$
 - **Different**
 - **Similar**

Task: $\mathcal{T} = \{y, f(\cdot)\}$

- A label space y
 - **Different**
 - **The same**
- An objective predictive function
 - **Different (but similar?)**


Source



{CAT, DOG}
{FELINE, CANINE}

$f_S(\cdot)$

Target



{LION, WOLF}
{FELINE, CANINE}

$f_T(\cdot)$

What is **transfer** learning about?

Train set



Test set



What is **transfer** learning about?

Source
domain



Target domain



What is transfer learning about?

Source domain



Source Task $y=\{\text{CAT, DOG}\}, f_s(\cdot)$

Target domain



Target Task $y=\{\text{LION, WOLF}\}, f_T(\cdot)$

Formalizing **transfer** learning


Domain: $\mathcal{D} = \{\mathcal{X}, P(X)\}$

- A feature space \mathcal{X}
 - **The Same (different)**
- A marginal probability distribution $P(X)$
 - **Different**
 - **Similar**

Task: $\mathcal{T} = \{y, f(\cdot)\}$

- A label space y
 - **Different**
 - **The same**
- An objective predictive function
 - **Different (but similar?)**


Source



{CAT, DOG}
{FELINE, CANINE}

$f_S(\cdot)$

Target



{LION, WOLF}
{FELINE, CANINE}

$f_T(\cdot)$

Formalizing **transfer** learning

Domain: $\mathcal{D} = \{\mathcal{X}, P(X)\}$


- A feature space \mathcal{X}
 - **The Same (different)**
- A marginal probability distribution $P(X)$
 - **Different**
 - **Similar**

Task: $\mathcal{T} = \{y, f(\cdot)\}$

- A label space y
 - **Different**
 - **The same**


- An objective predictive function
 - **Different (but similar?)**

Source



{CAT, DOG}
{FELINE, CANINE}

Target



{LION, WOLF}
{FELINE, CANINE}


$f_S(\cdot)$

$f_T(\cdot)$

transfer learning


Source
 $f_s(\cdot)$

Target
 $f_T(\cdot)$




- Are they similar?
- Can we just use $\hat{f}_S(\cdot)$ to approximate $f_T(\cdot)$?
- Can we reuse part of it?

Representation learning




Deep Neural Networks are **representation** learning techniques

- **Support Vector Machine** (SVM) is just a **classifier** (a very good one).
- SVM find the best boundary separating the data instances into different classes in a **given** feature space.




Deep Neural Networks are **representation** learning techniques

- **Support Vector Machine** (SVM) is just a **classifier** (a very good one).
- SVM find the best boundary separating the data instances into different classes in a **given** feature space.




Deep Neural Networks are **representation** learning techniques


- SVMs using the **kernel trick** can overcome the linear limitation through an **implicit** mapping to a higher dimensional feature space



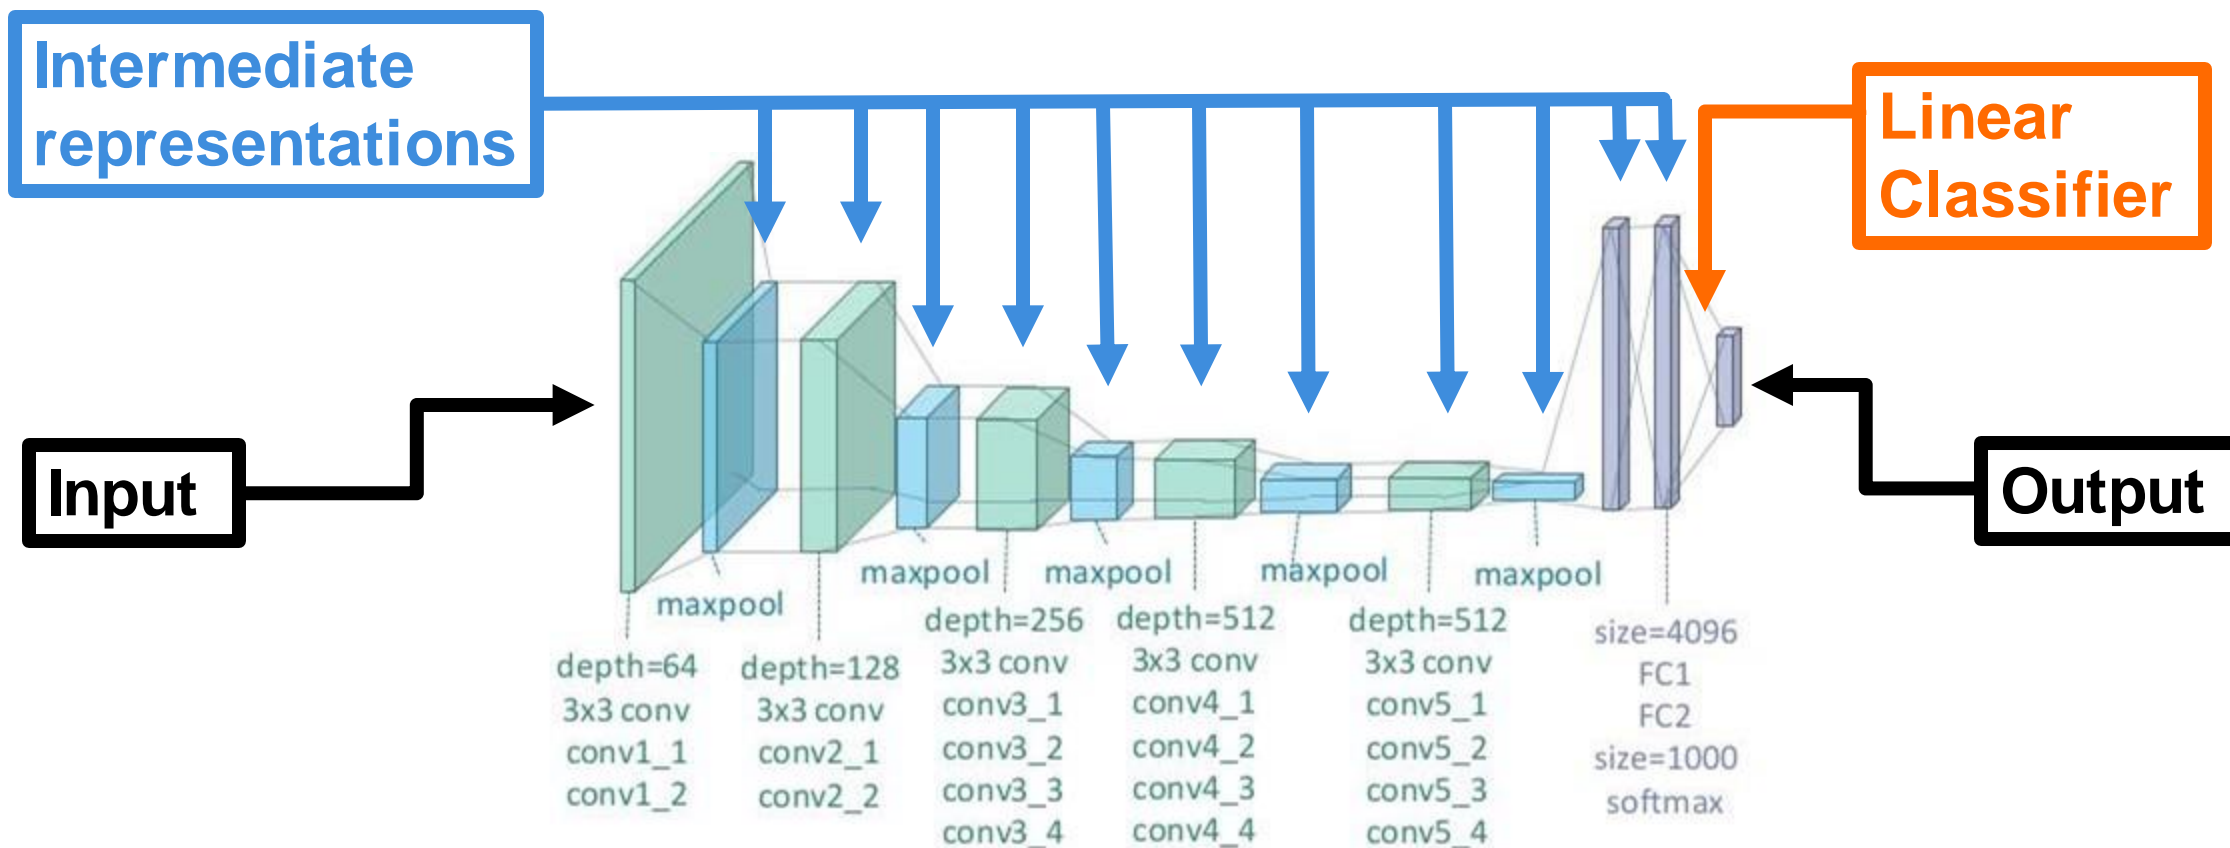
Deep Neural Networks are **representation** learning techniques




Deep Neural Networks are **representation** learning techniques



Deep Neural Networks are **representation** learning techniques



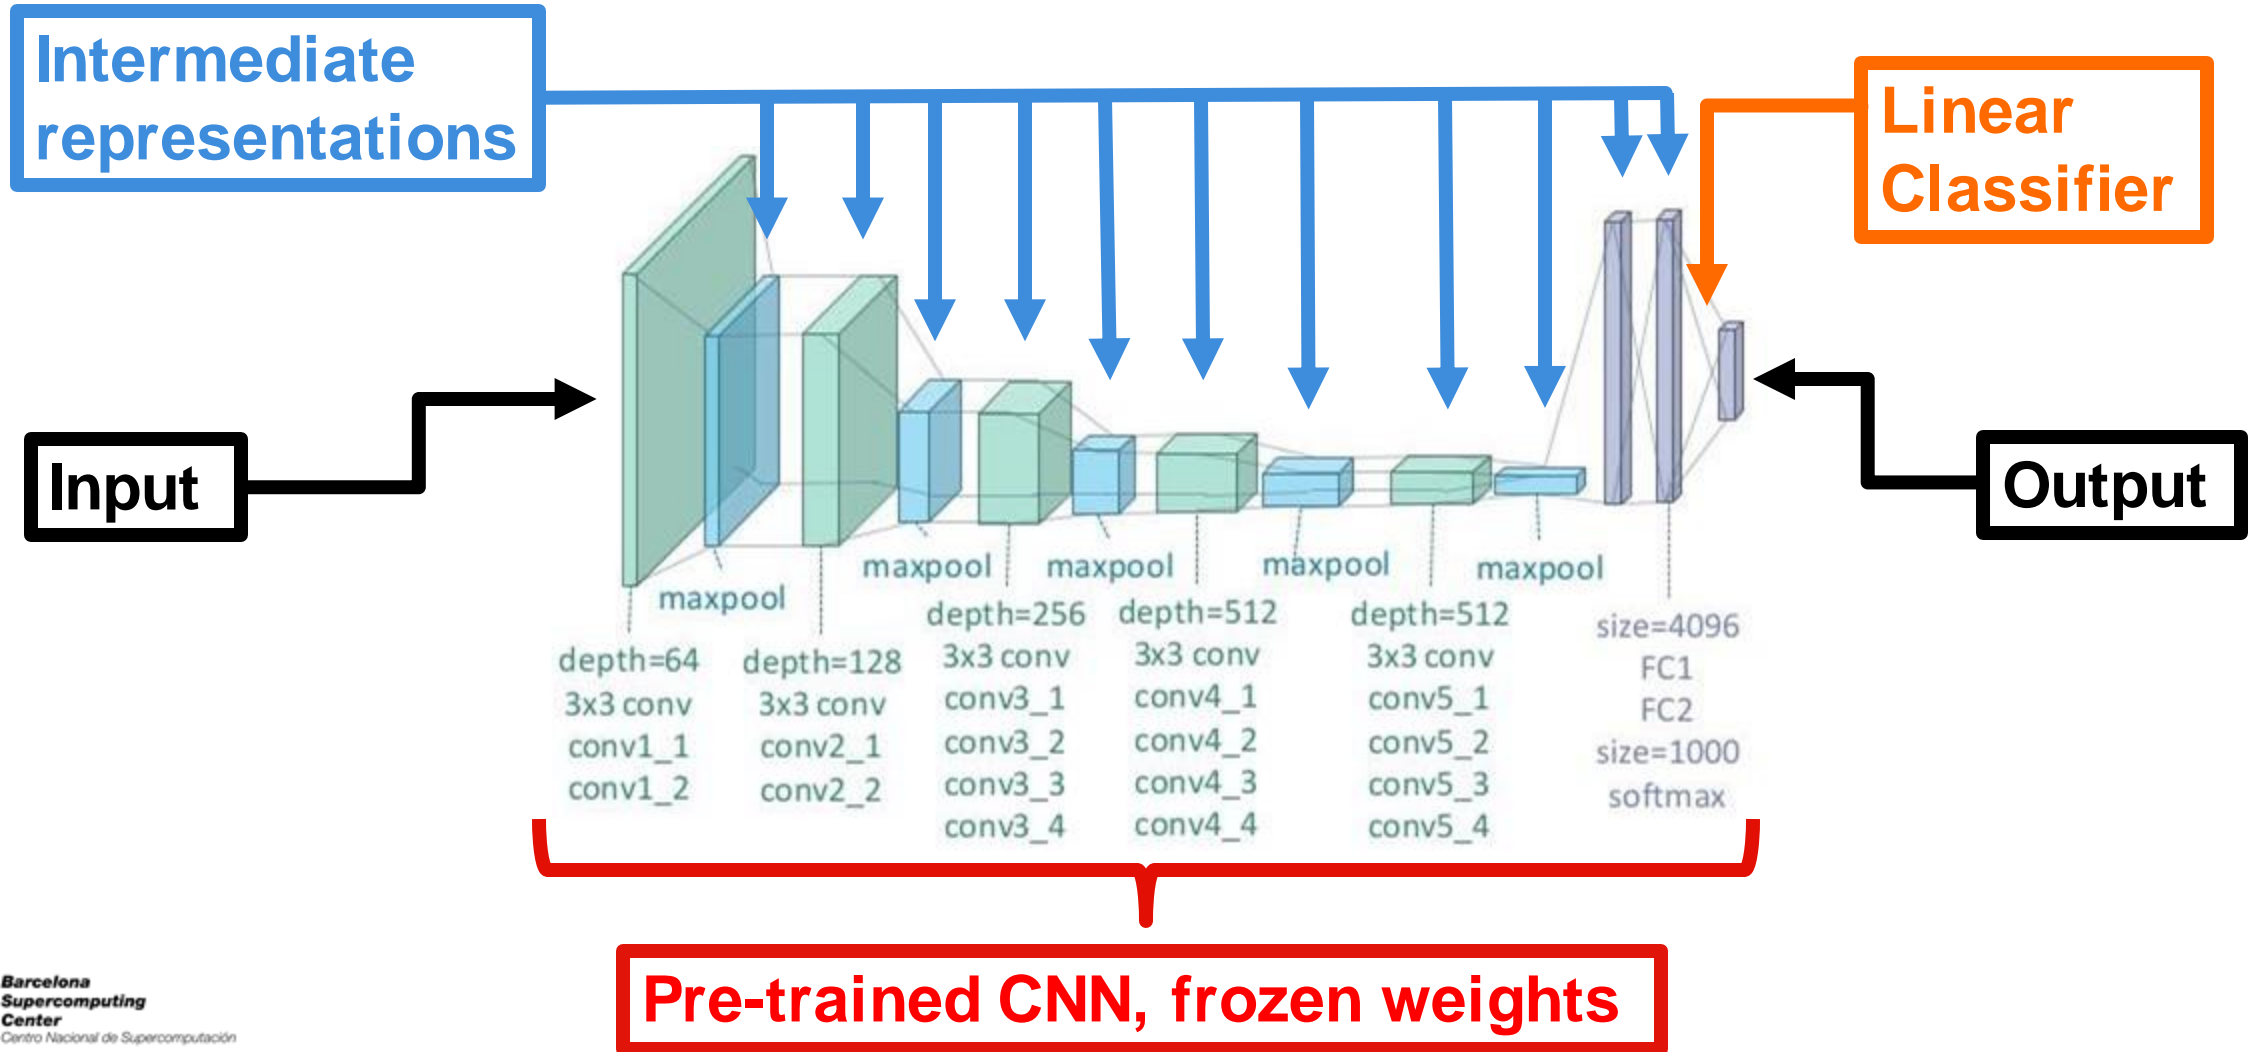
**Reusing
DNNs
knowledge**



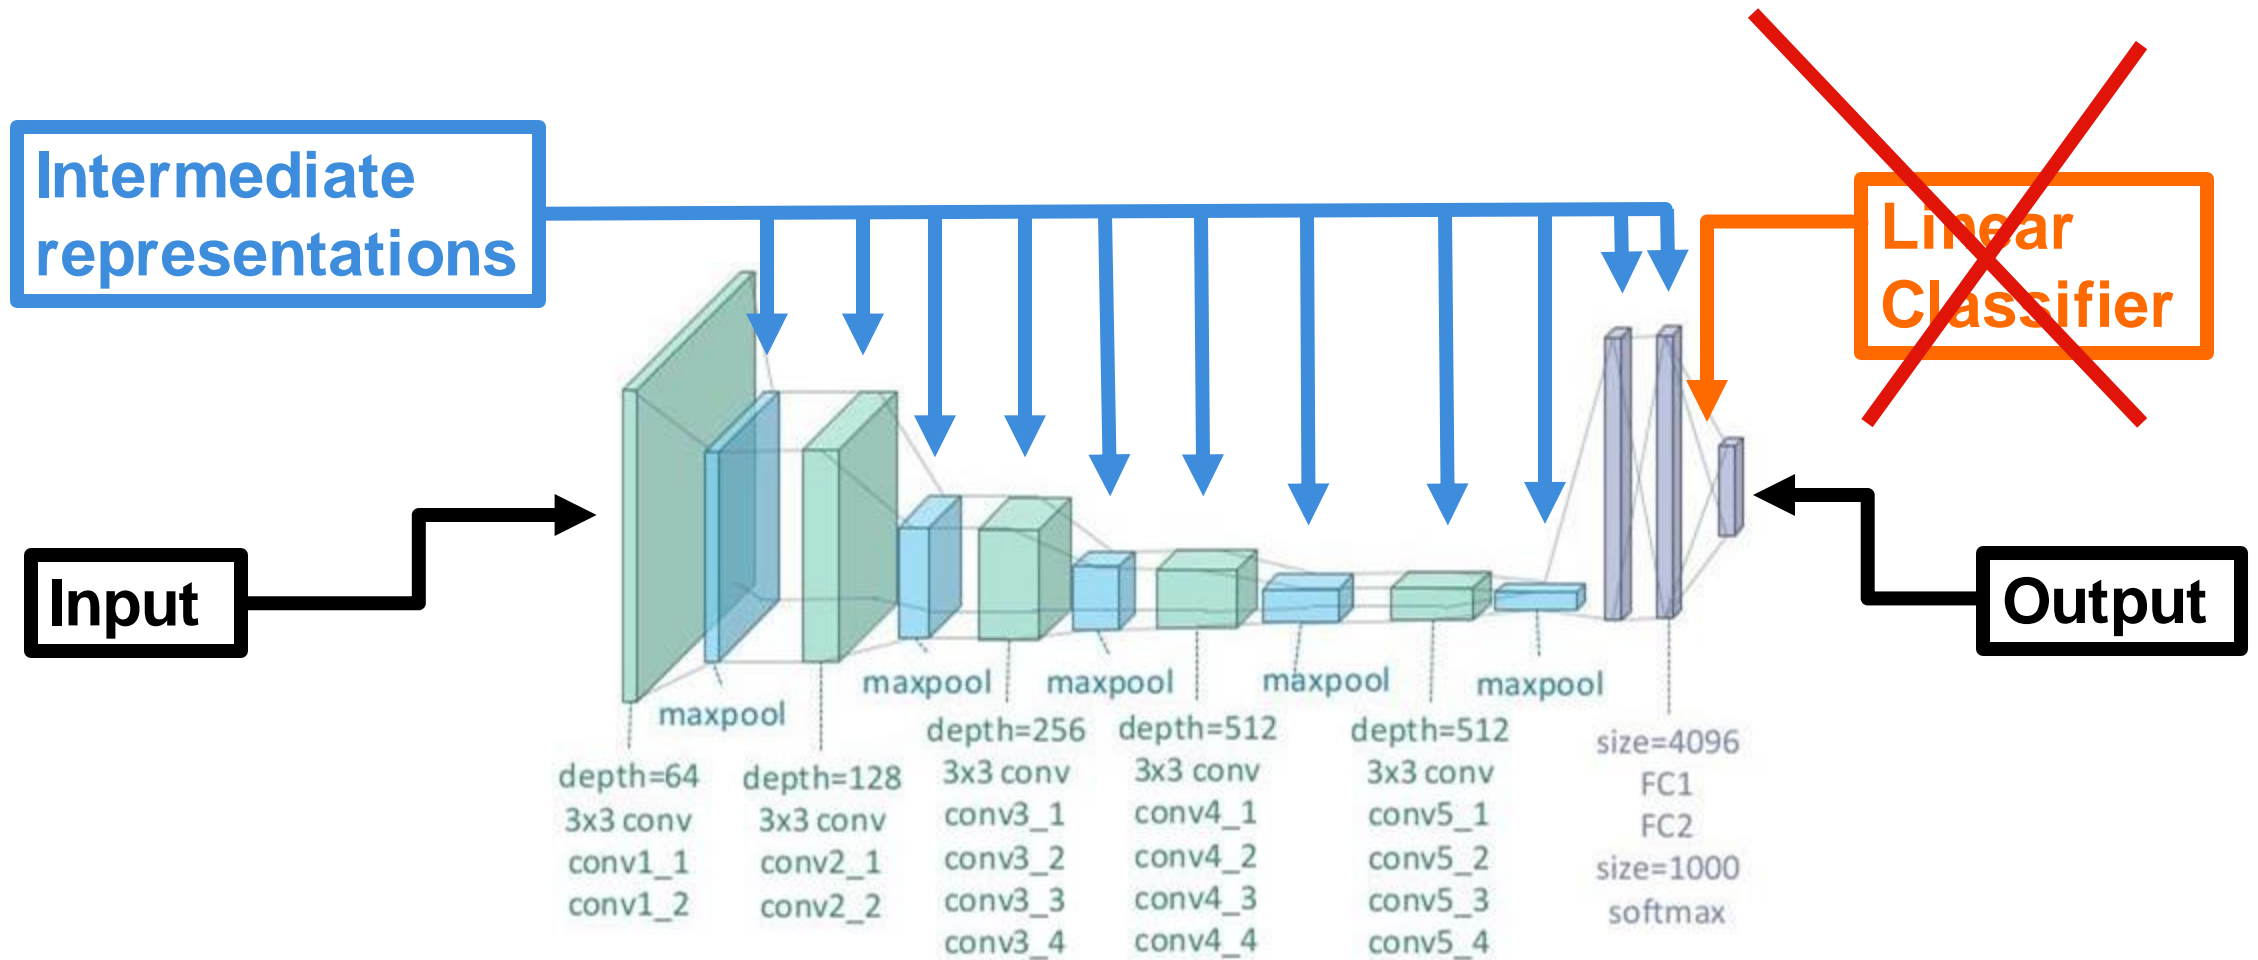
**SAVE THE
EARTH**

**REUSE
DNNs**

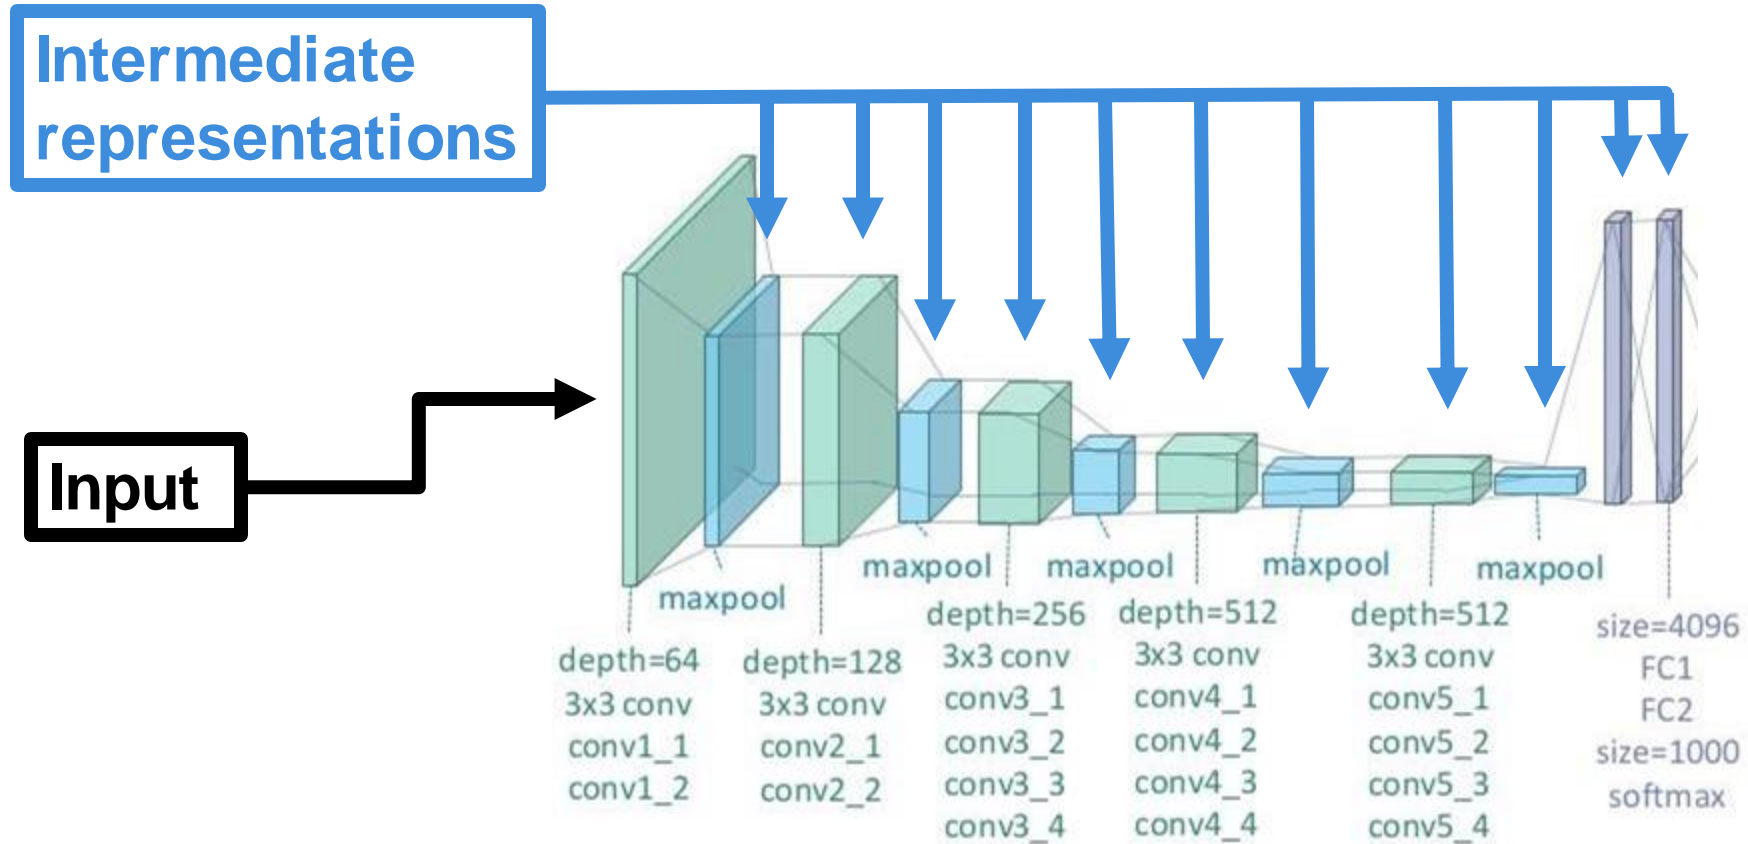
Feature extraction



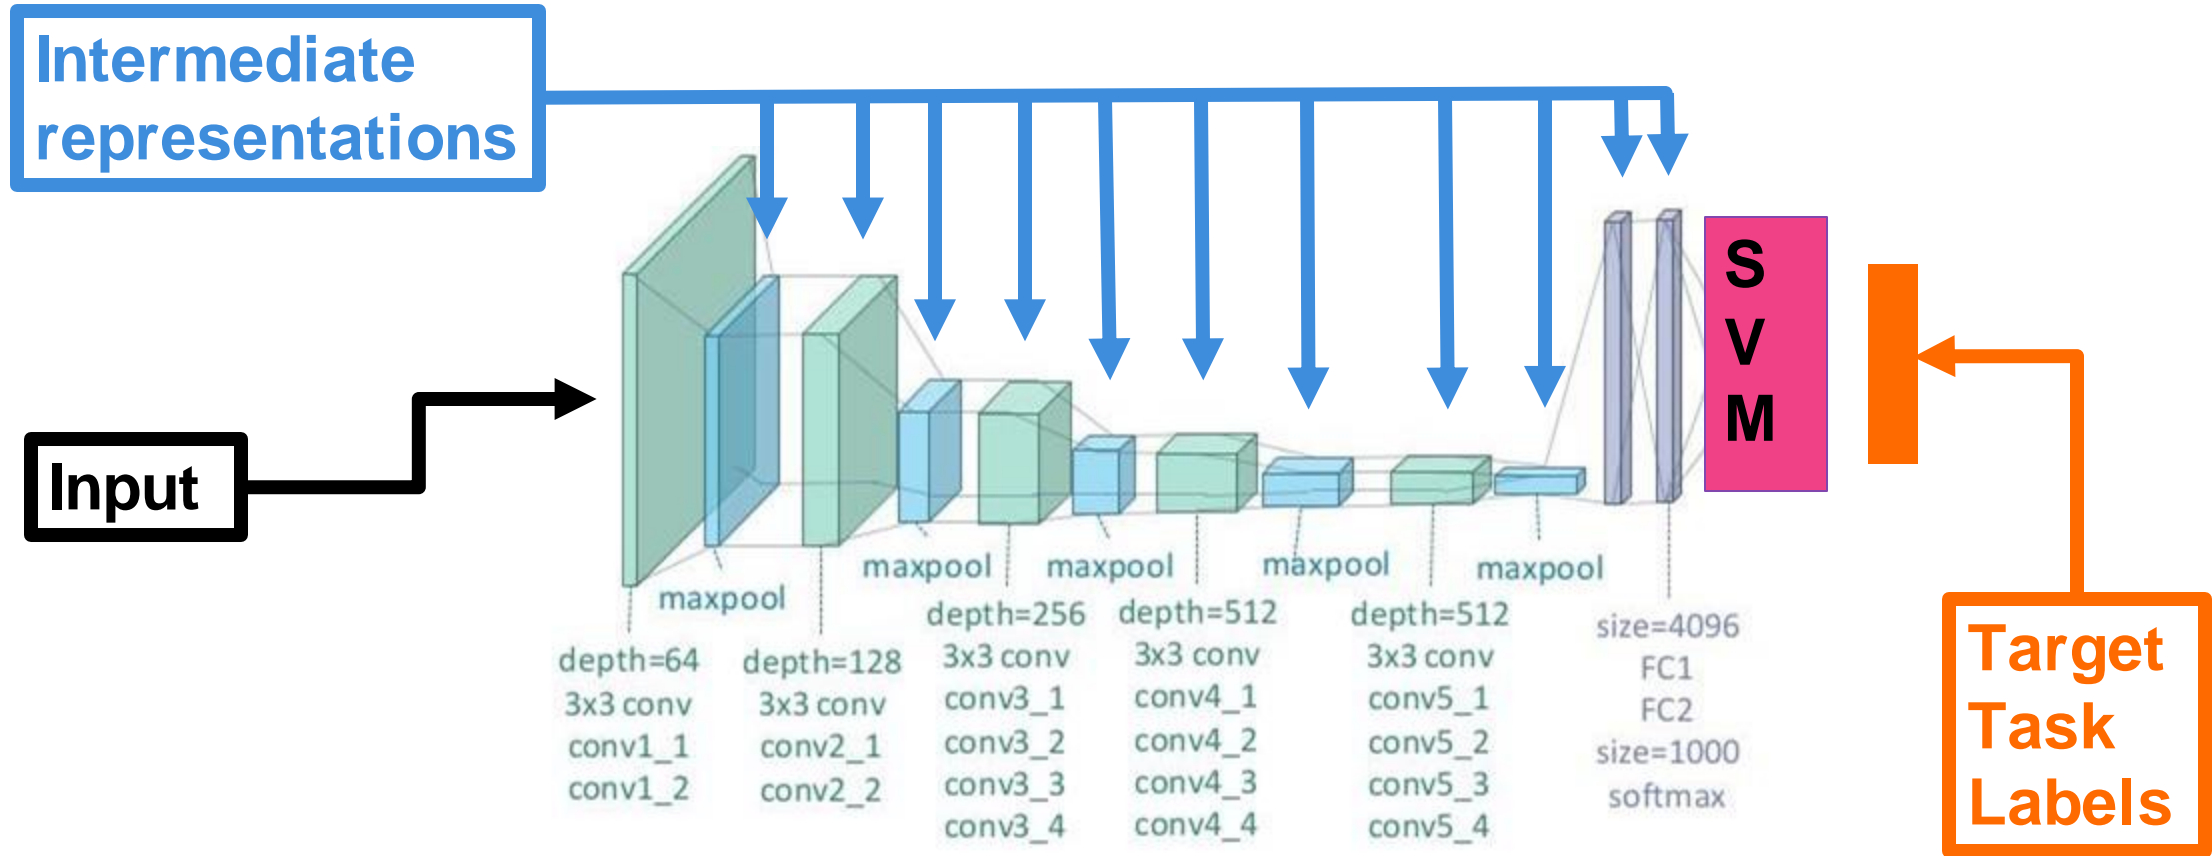
Feature extraction



Feature extraction




Feature extraction




Simple solutions

- DNN last layer features + SVM
(Feature extraction)




Simple solutions

- DNN last layer features + SVM
(Feature extraction)
 - **Similar task and domain**

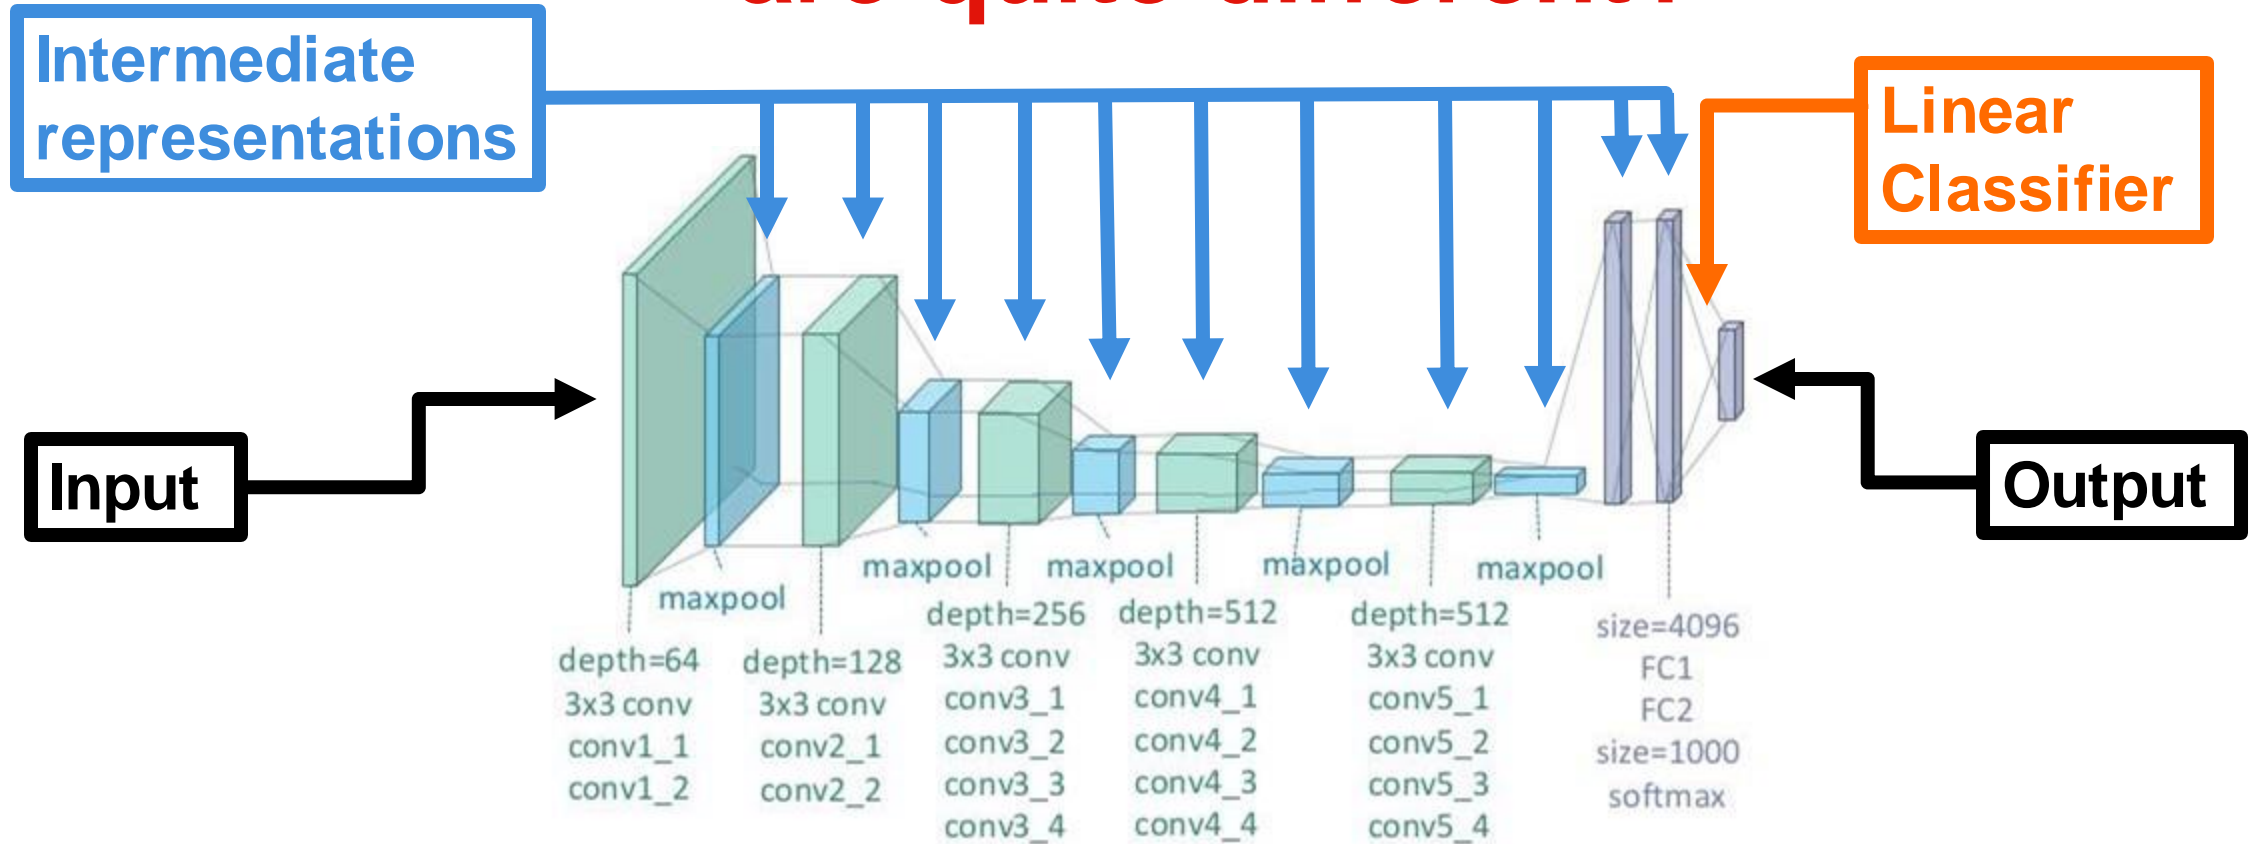


Fine tuning



Fine tuning

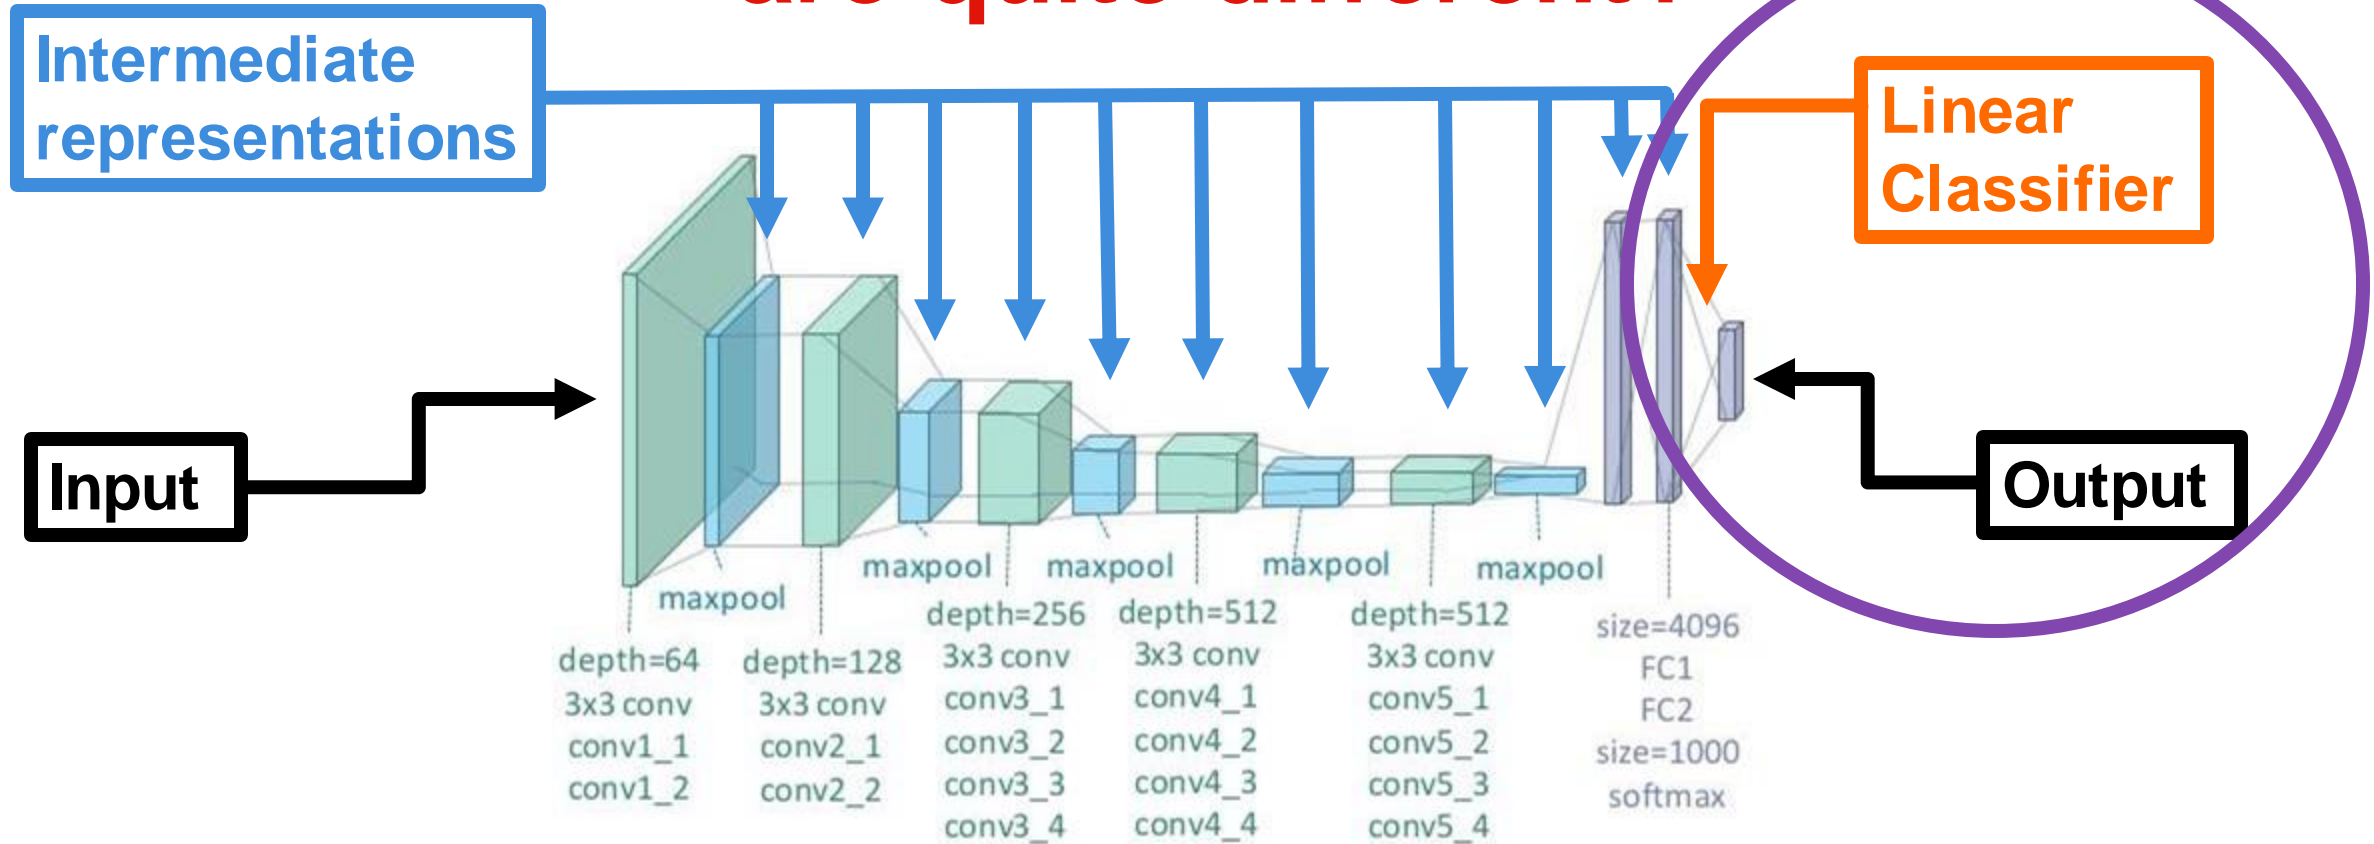
What if the tasks
are quite different?



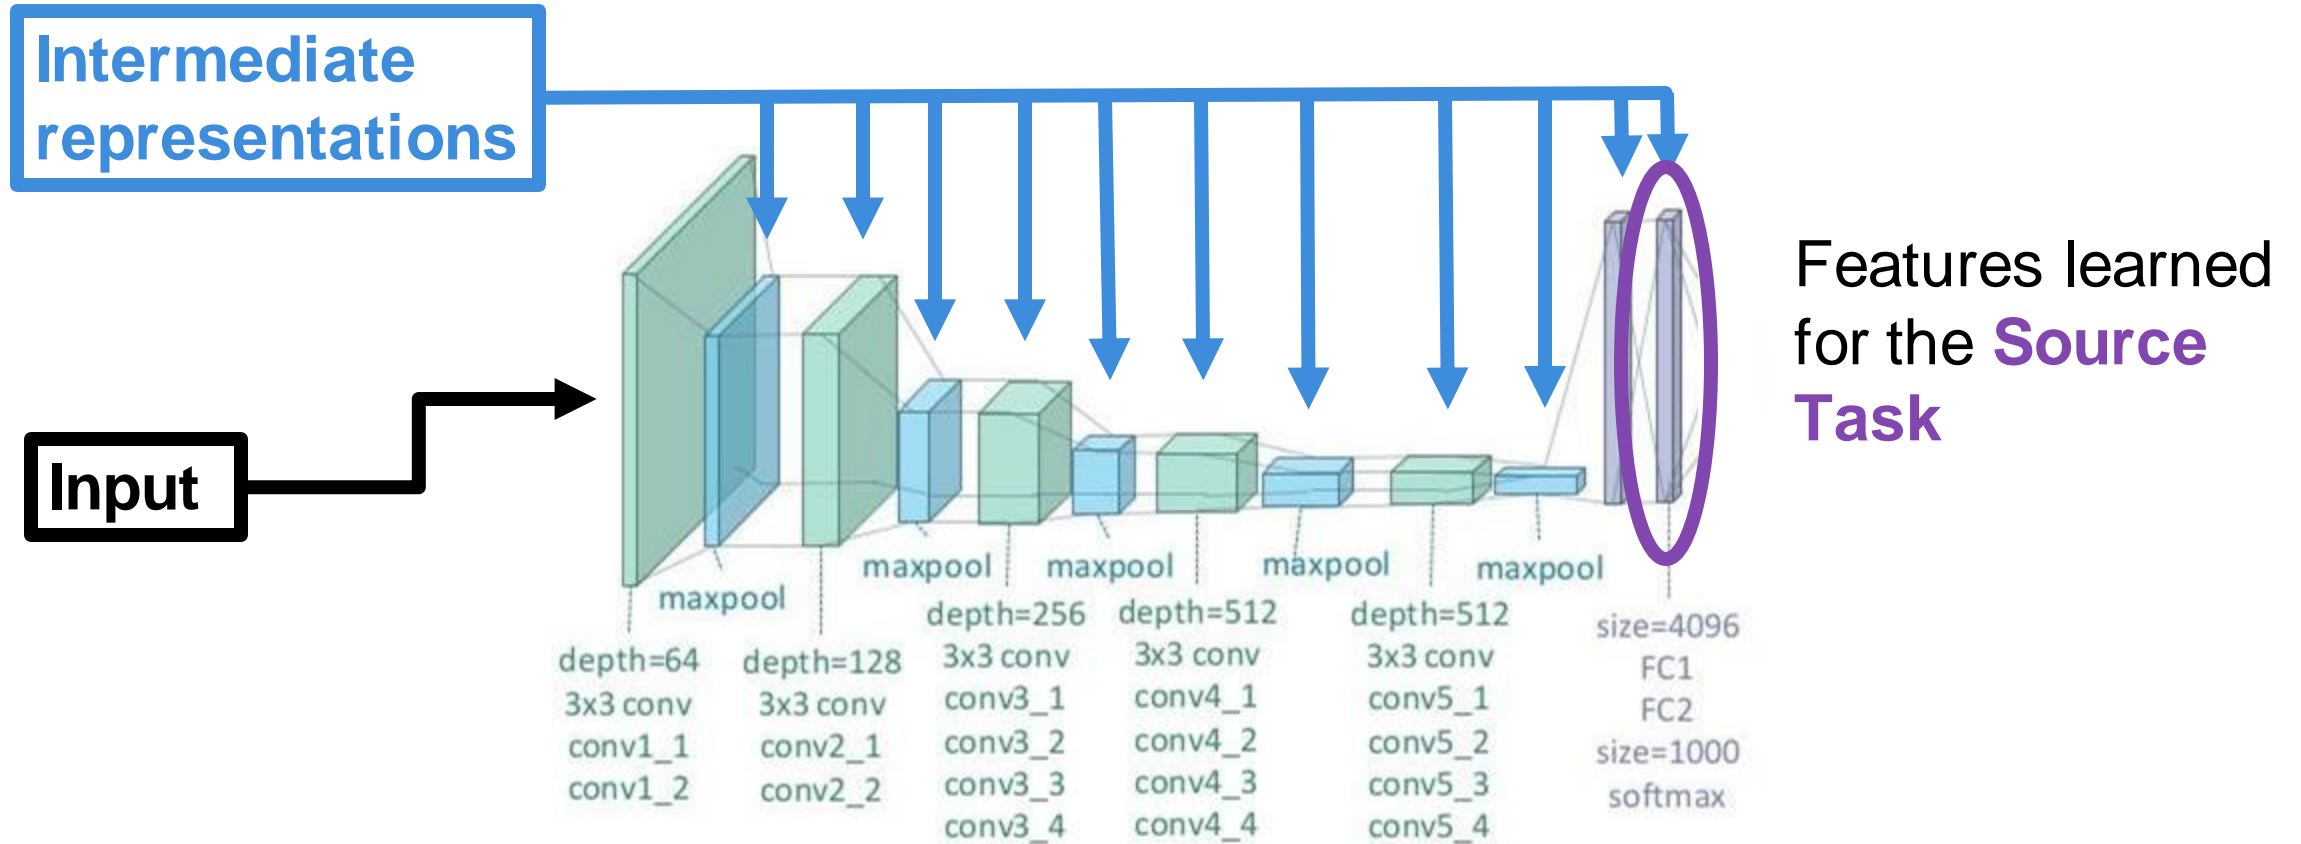
Fine tuning

What if the tasks are quite different?

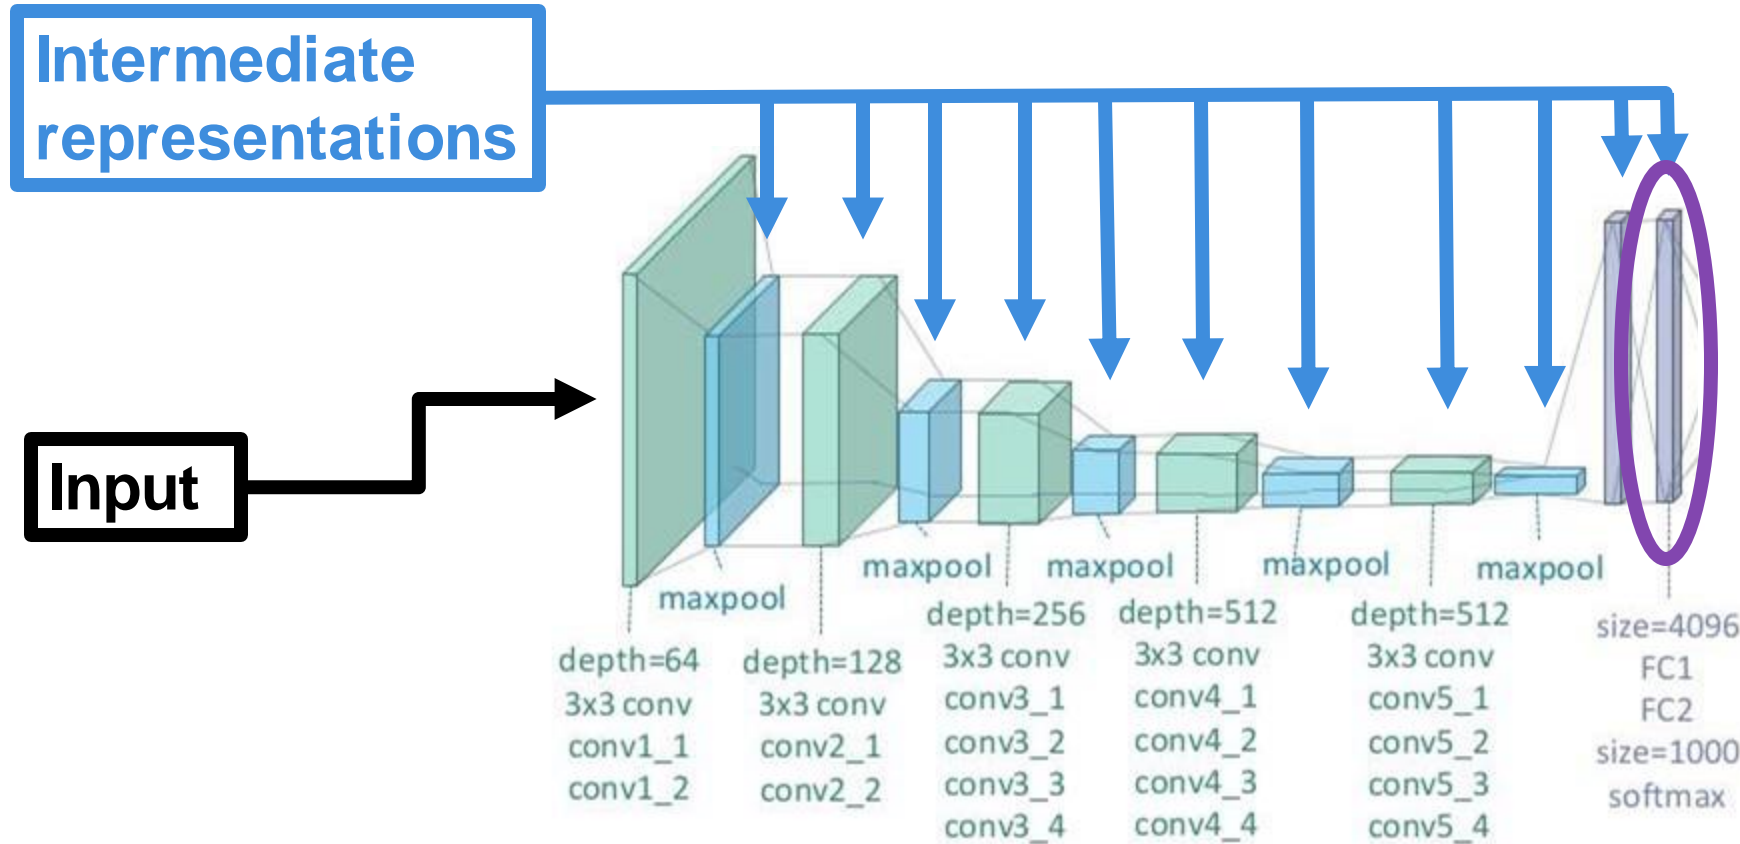
Source Task



Fine tuning



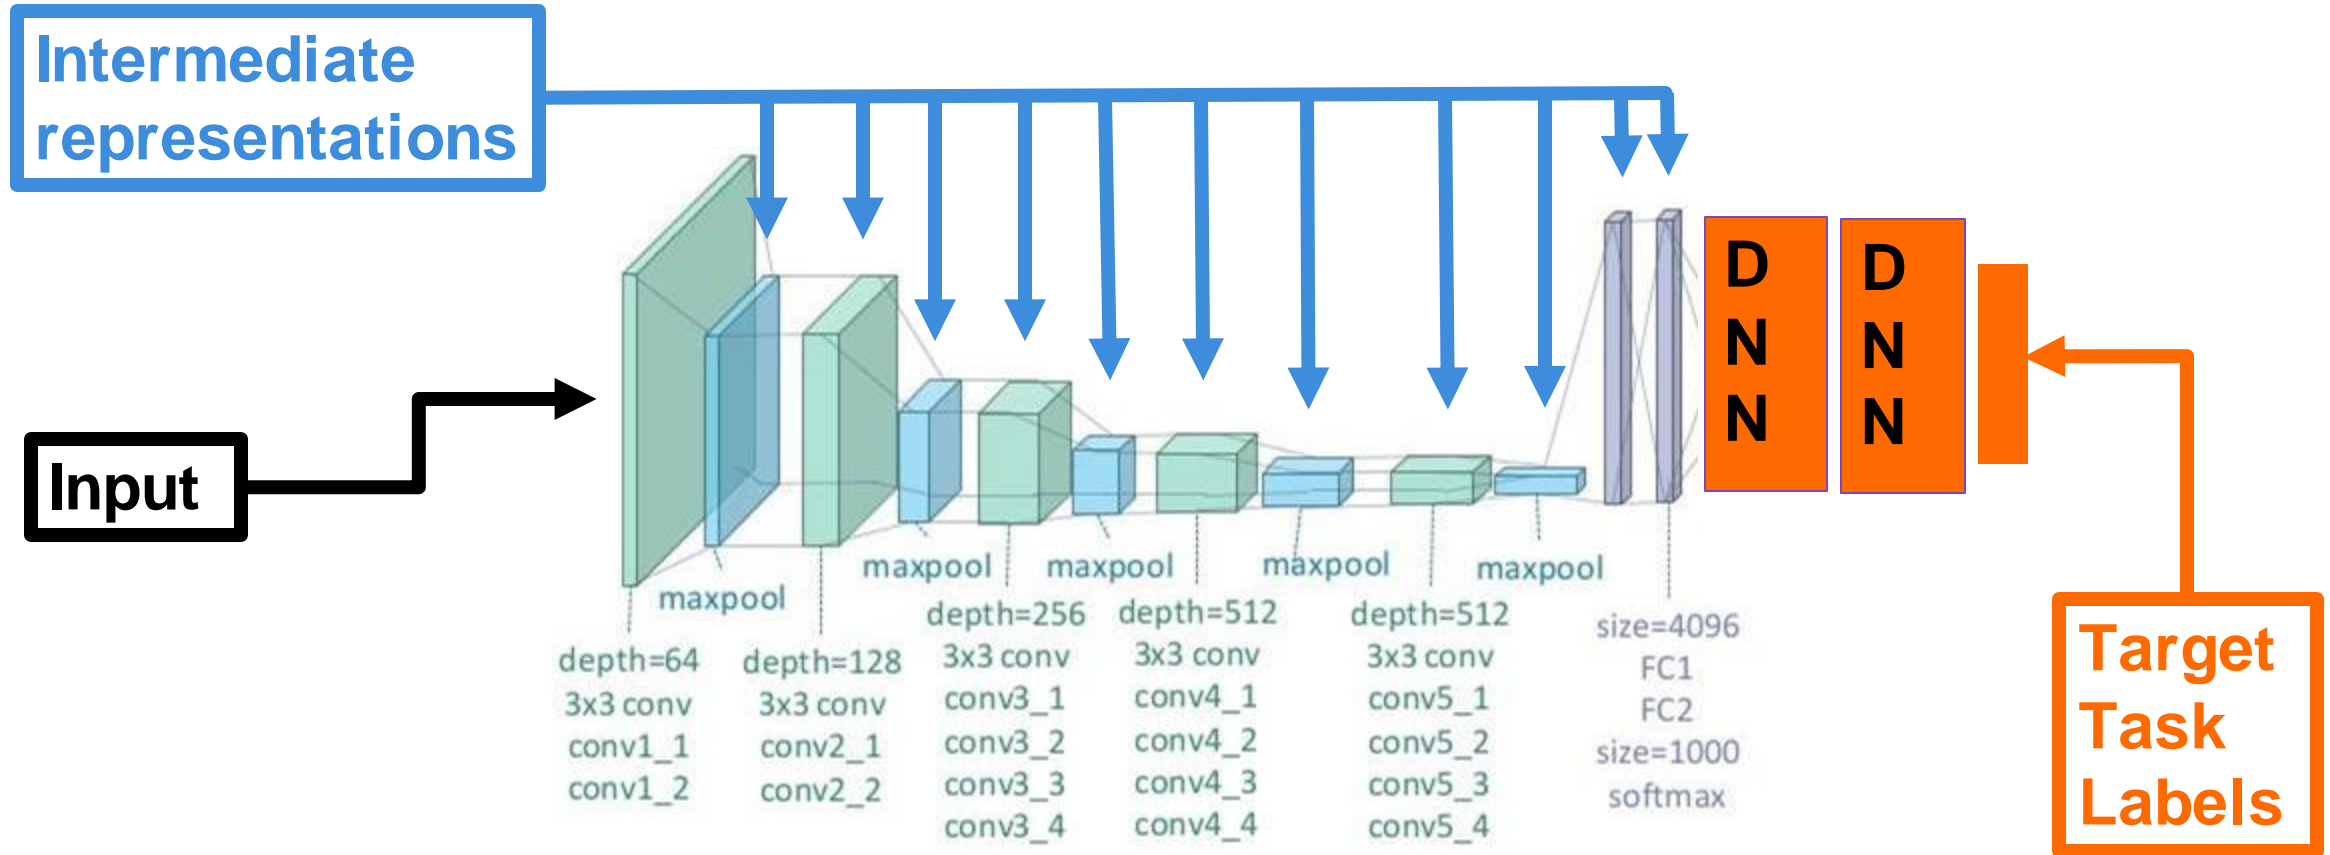
Fine tuning



Features learned for the **Source Task**

Can we make them better?

Fine tuning

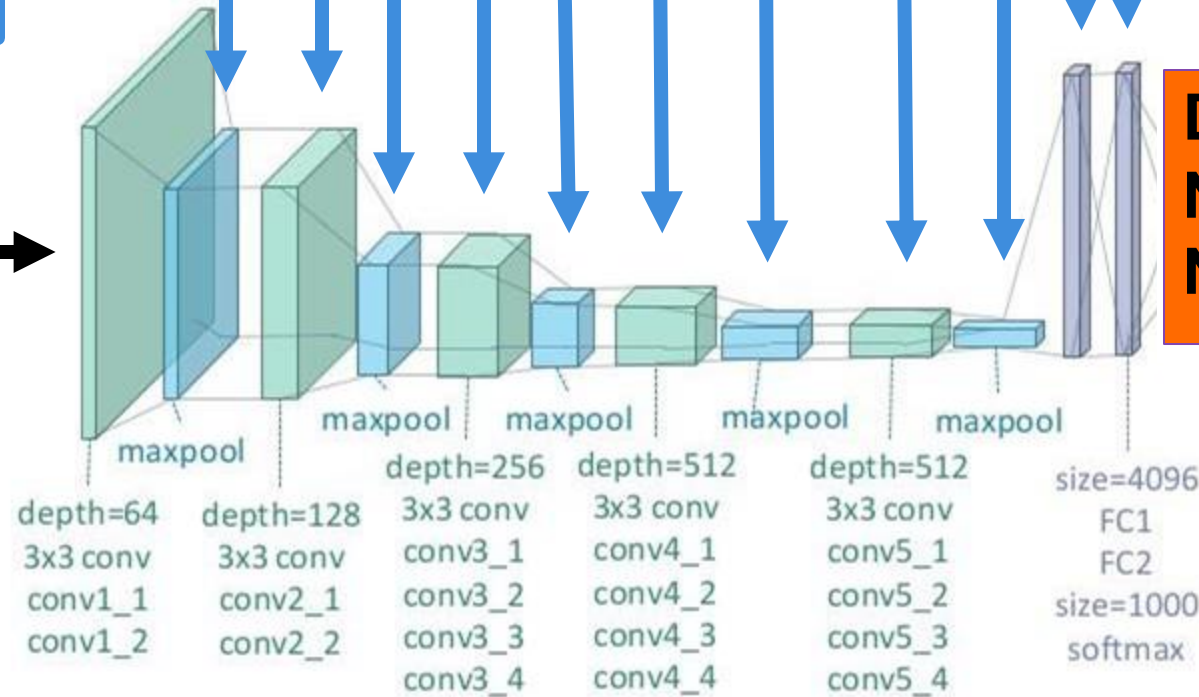


Fine tuning

Error back-propagation

Intermediate representations

Input



D
N
N

D
N
N

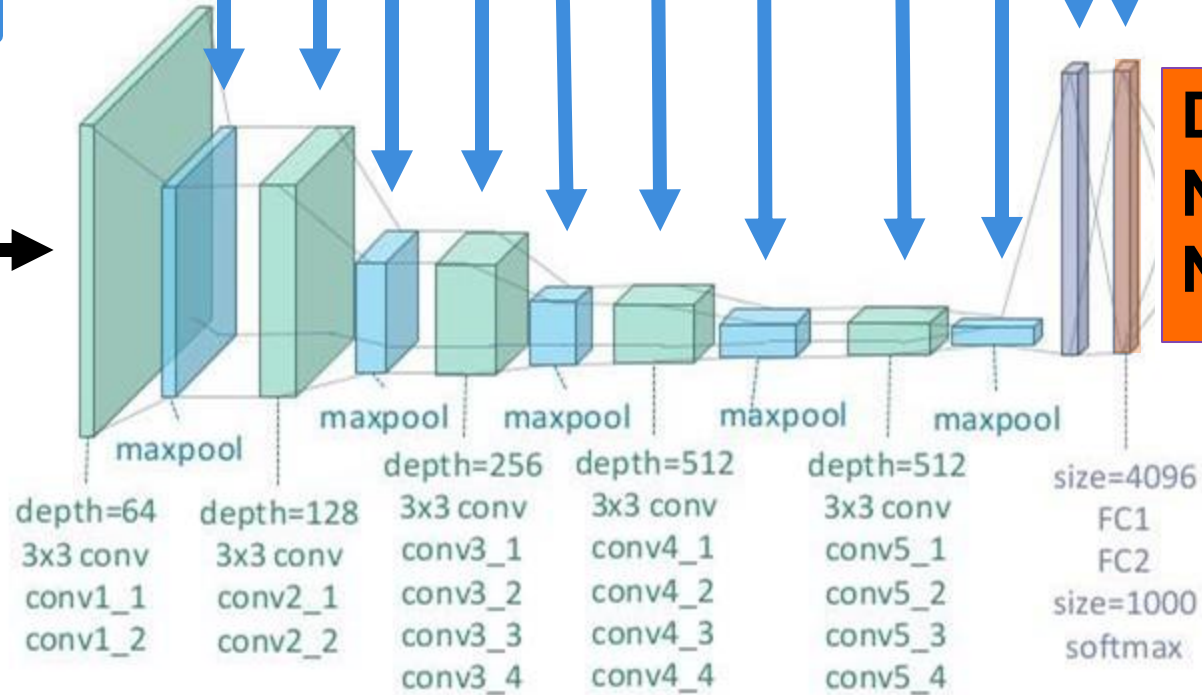
Target
Task
Labels

Fine tuning

Error back-propagation

Intermediate representations

Input



D
N
N

D
N
N

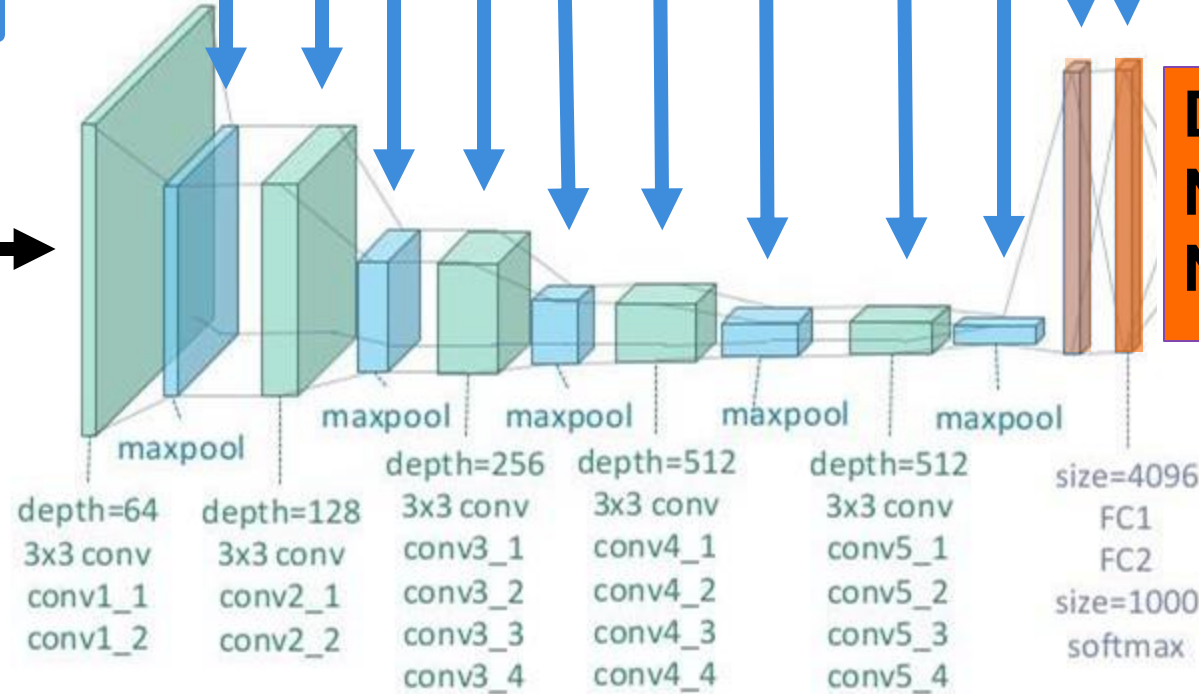
Target Task Labels

Fine tuning

Error back-propagation

Intermediate representations

Input



D
N
N

D
N
N

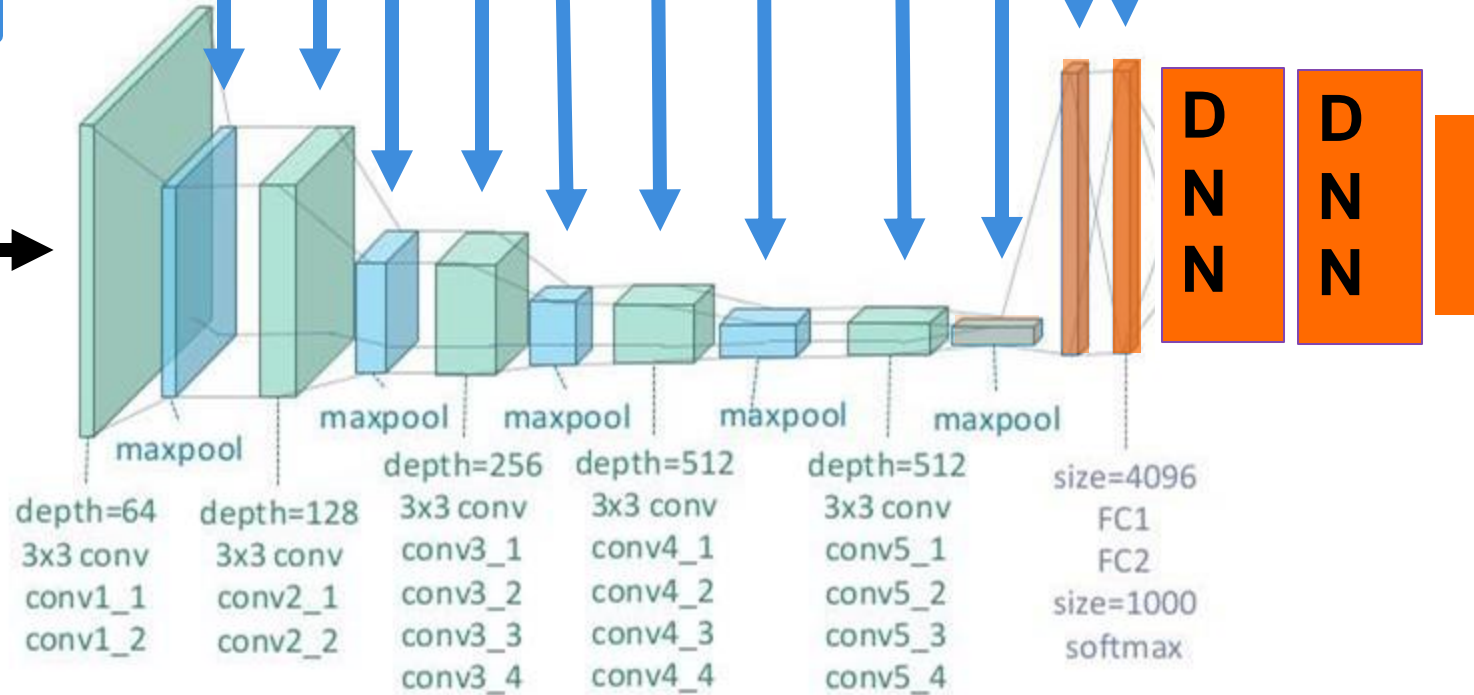
Target Task Labels

Fine tuning

Error back-propagation

Intermediate representations

Input



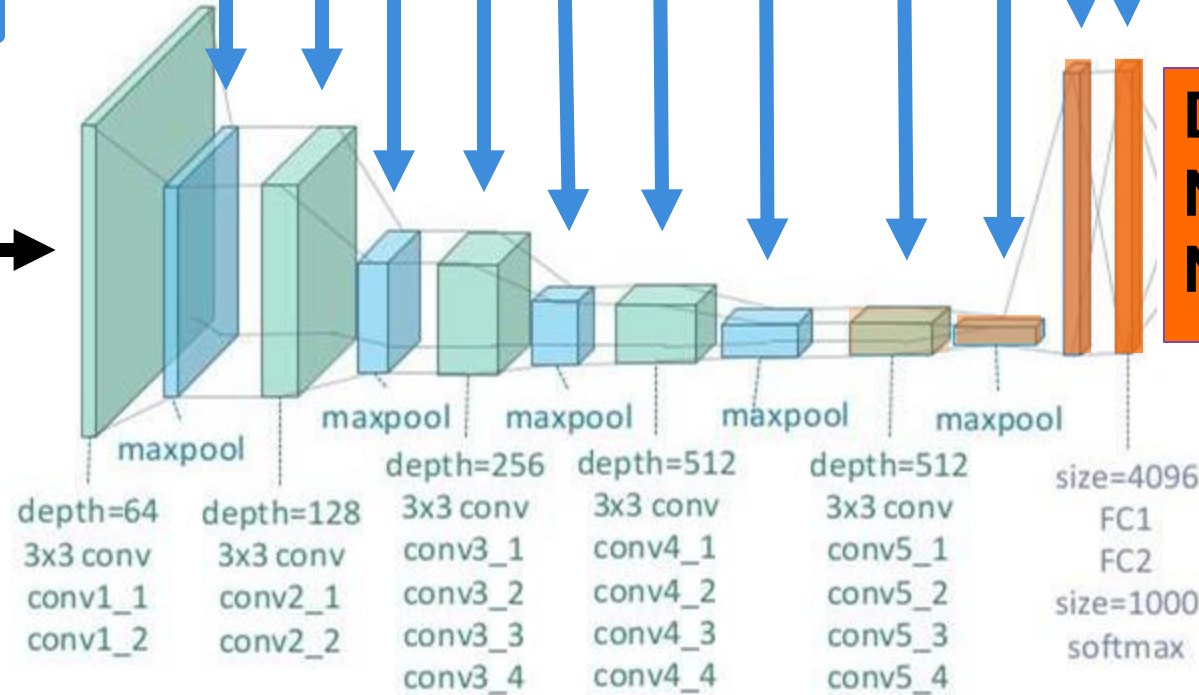
Target Task Labels

Fine tuning

Error back-propagation

Intermediate representations

Input

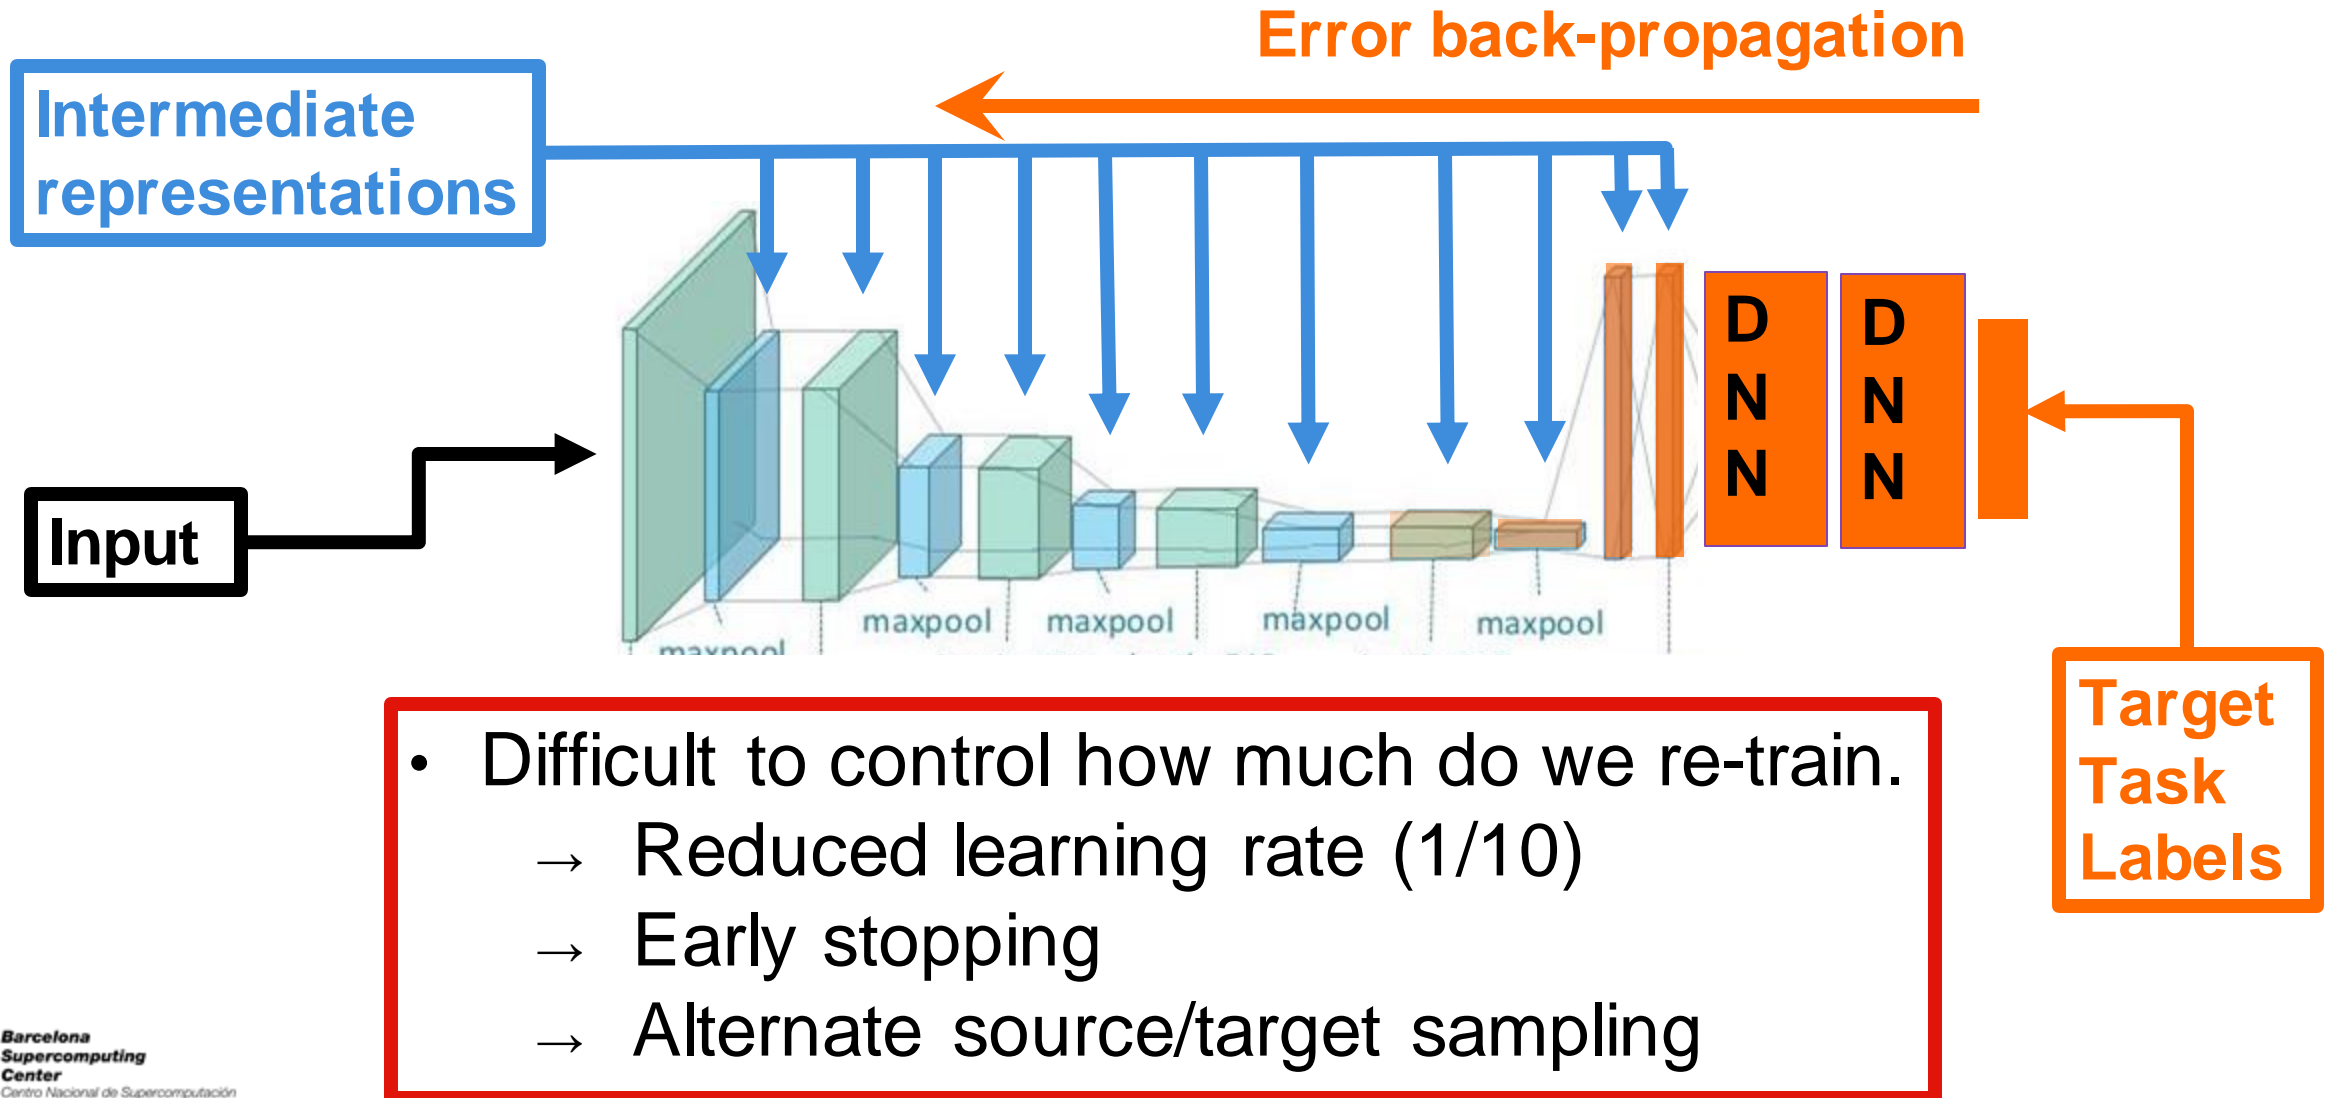


D
N
N

D
N
N



Target Task Labels

Fine tuning




Simple solutions

- DNN last layer features + SVM
(Feature extraction)
 - **Similar task and domain**
- Add one or several NN layers +
Fine-tuning pre-trained layers

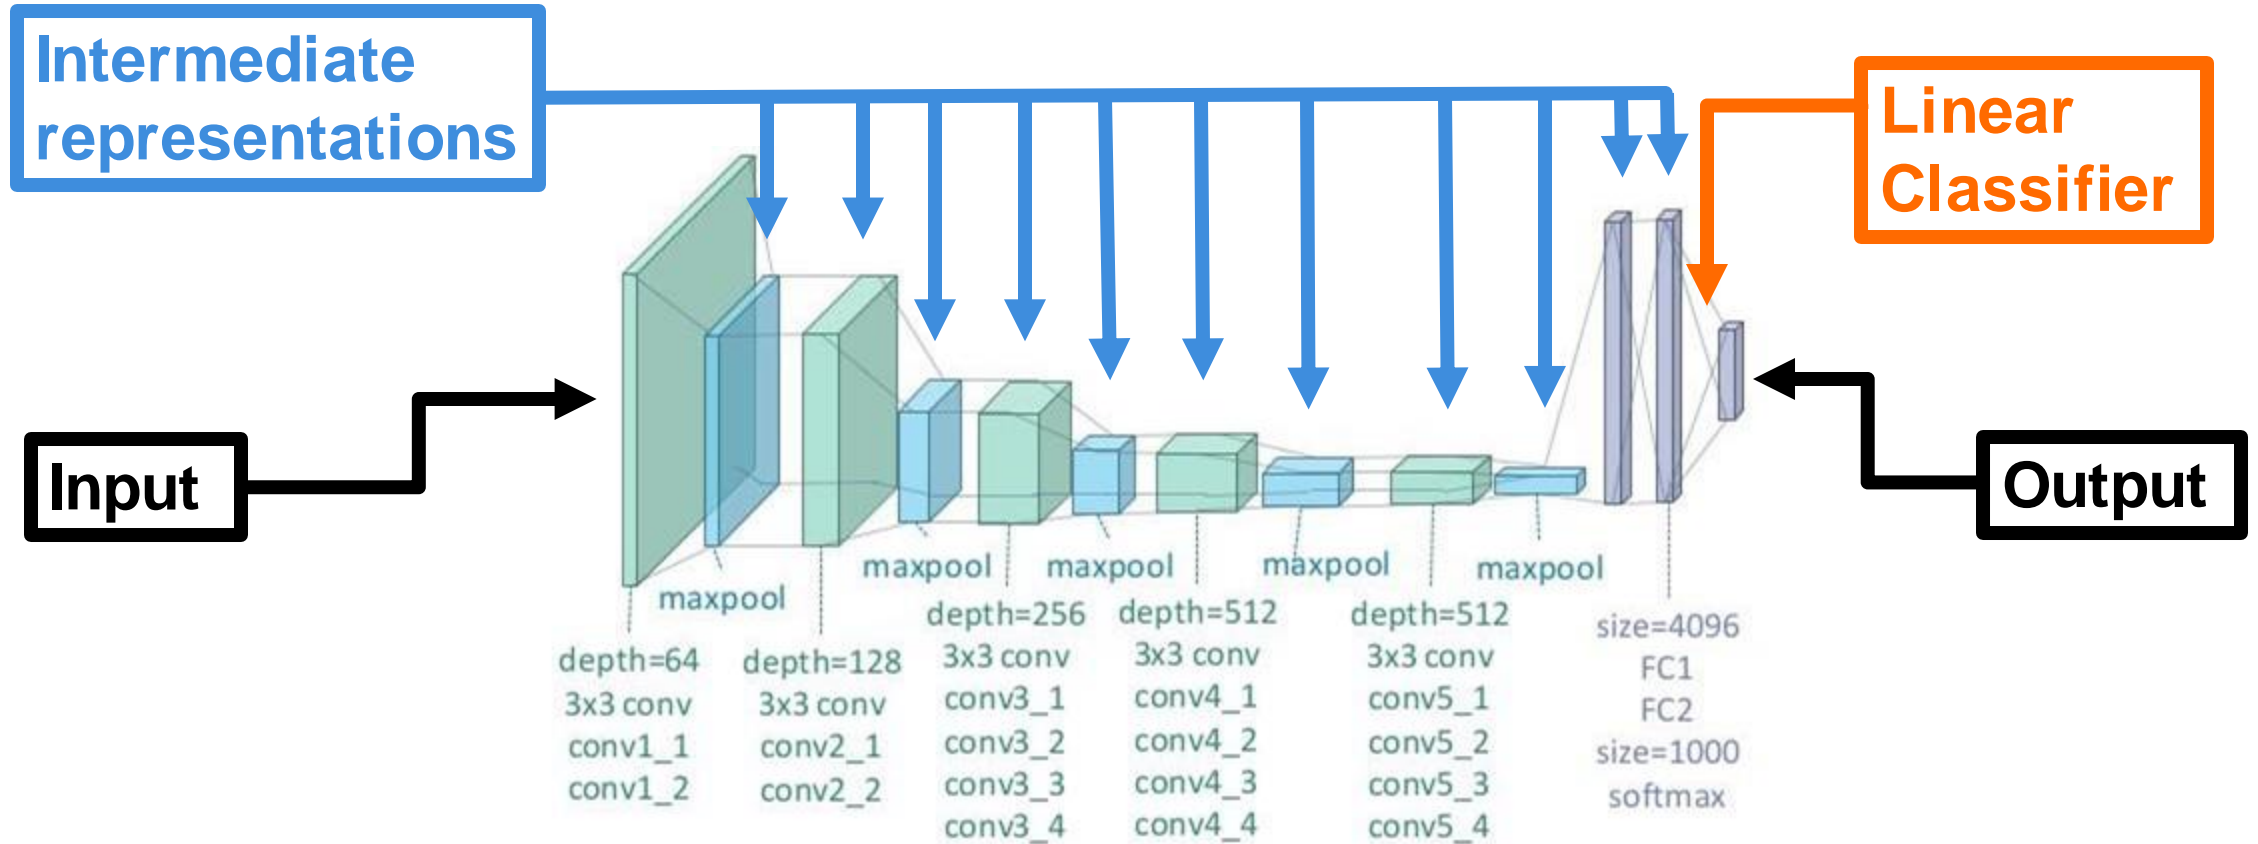


Simple solutions


- DNN last layer features + SVM
(Feature extraction)
 - **Similar task and domain**
- Add one or several NN layers +
Fine-tuning pre-trained layers
 - **Enough data**




Beyond the last layer



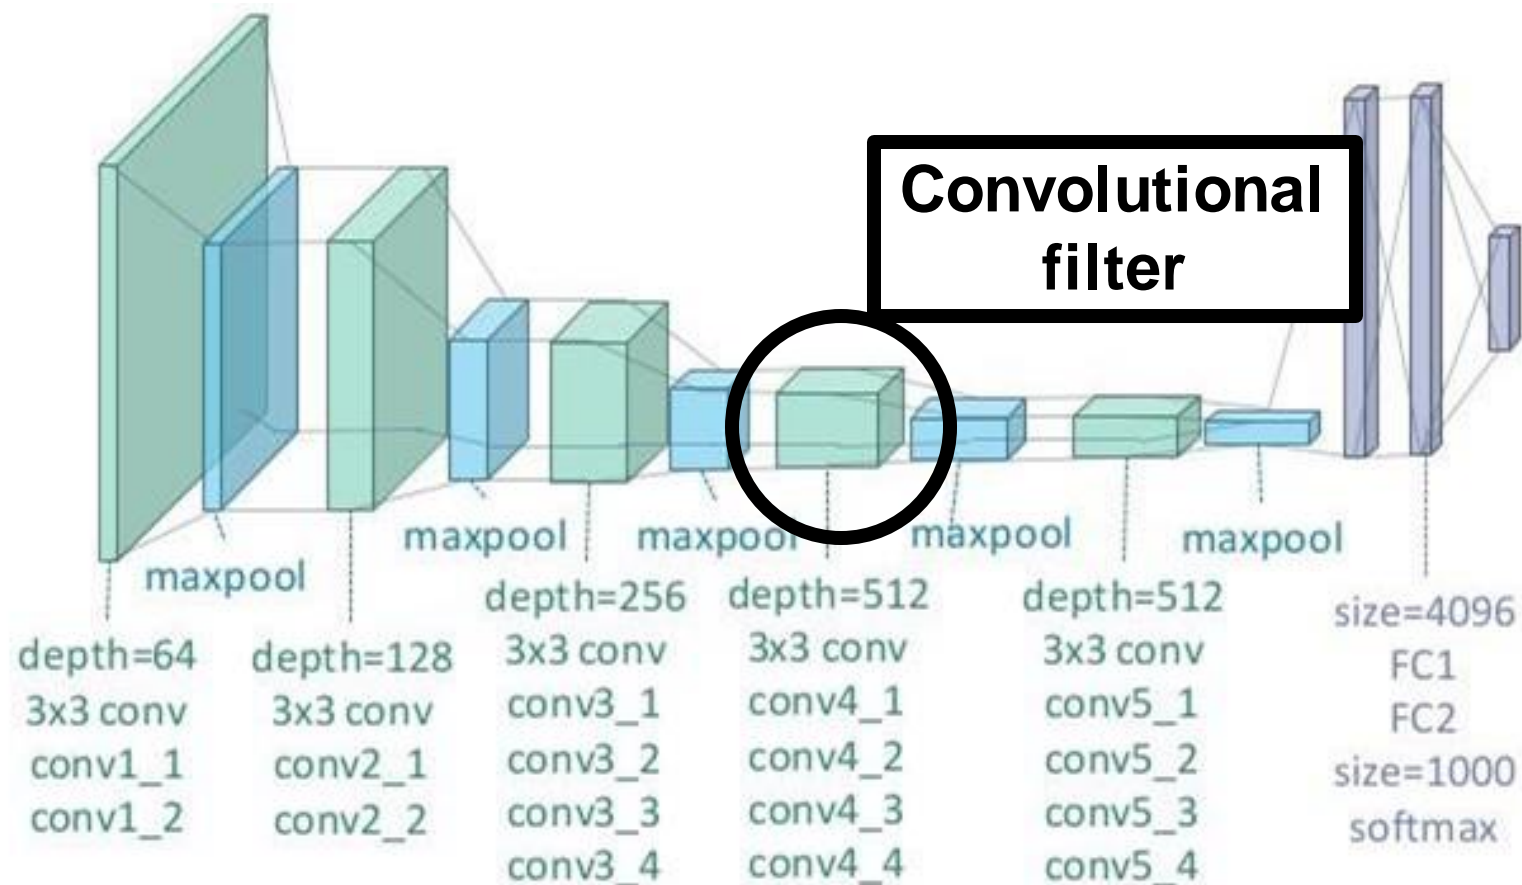
Knowledge **inside** DNN




Knowledge **inside** DNN



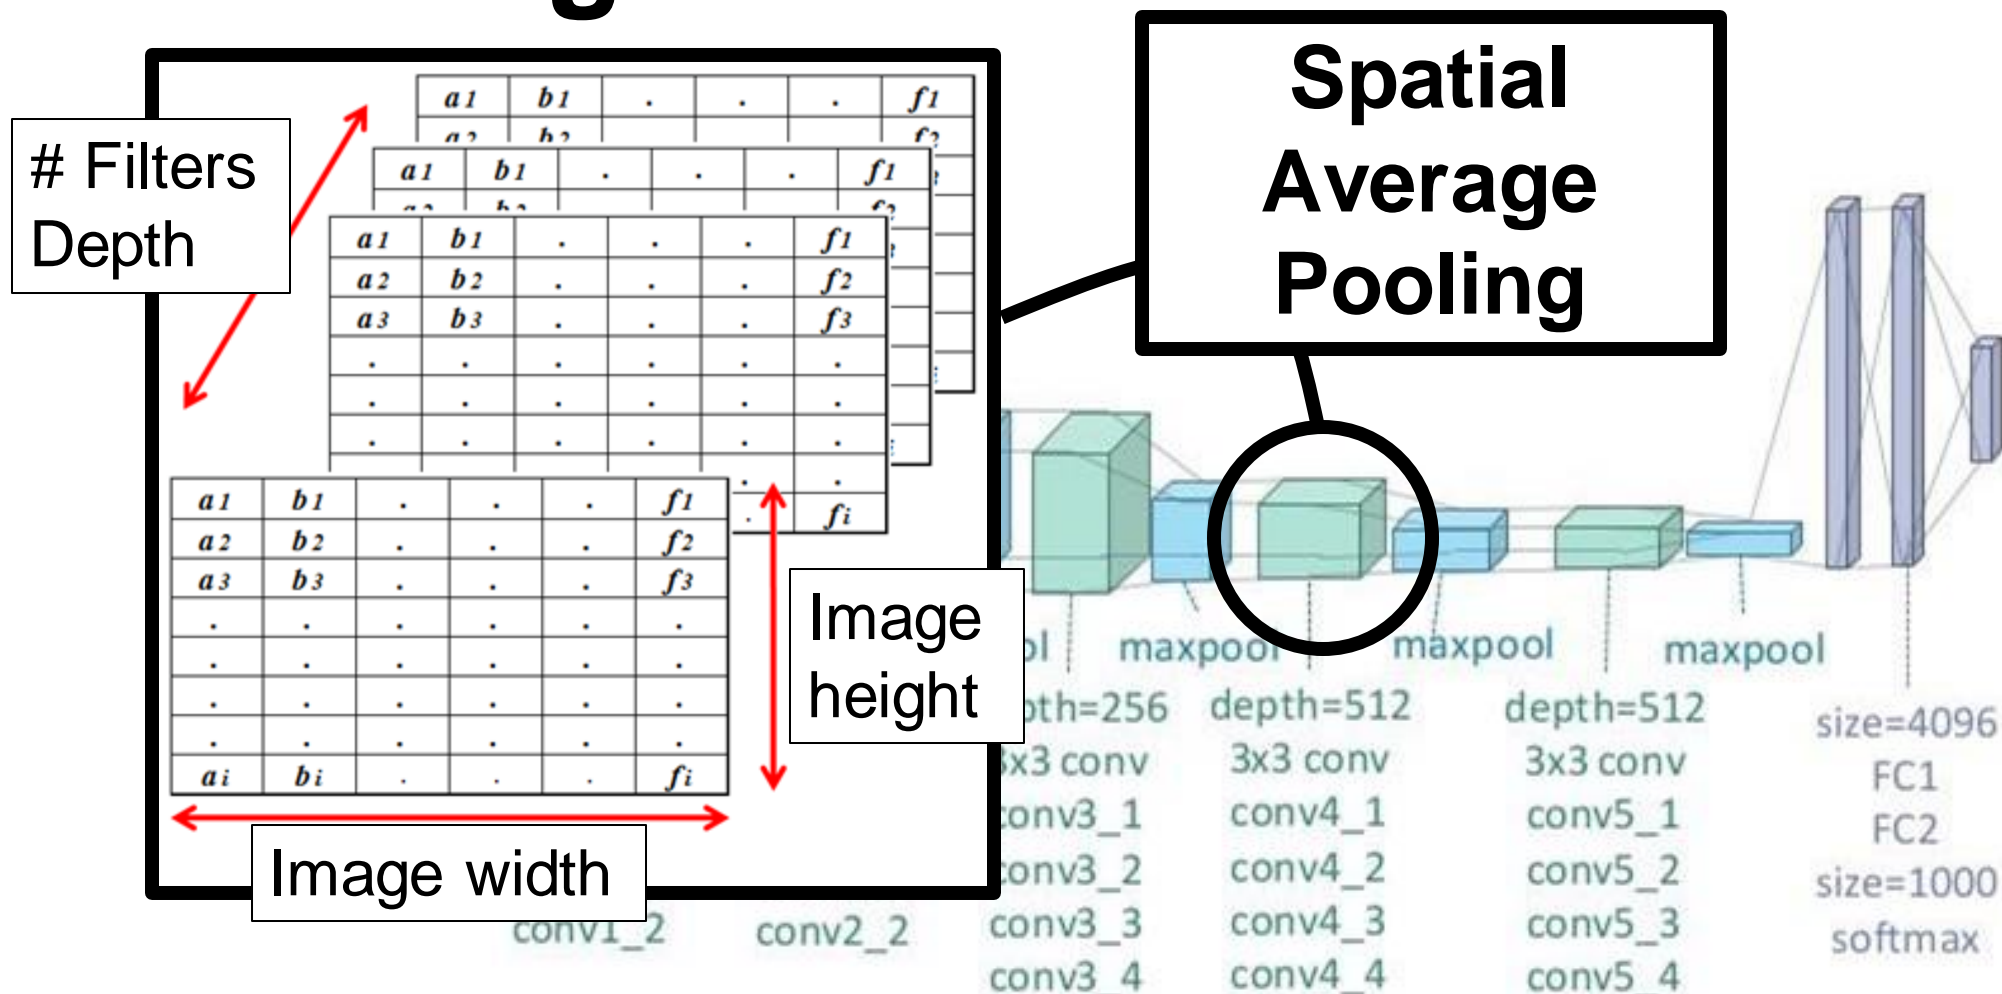
Knowledge **inside** DNN



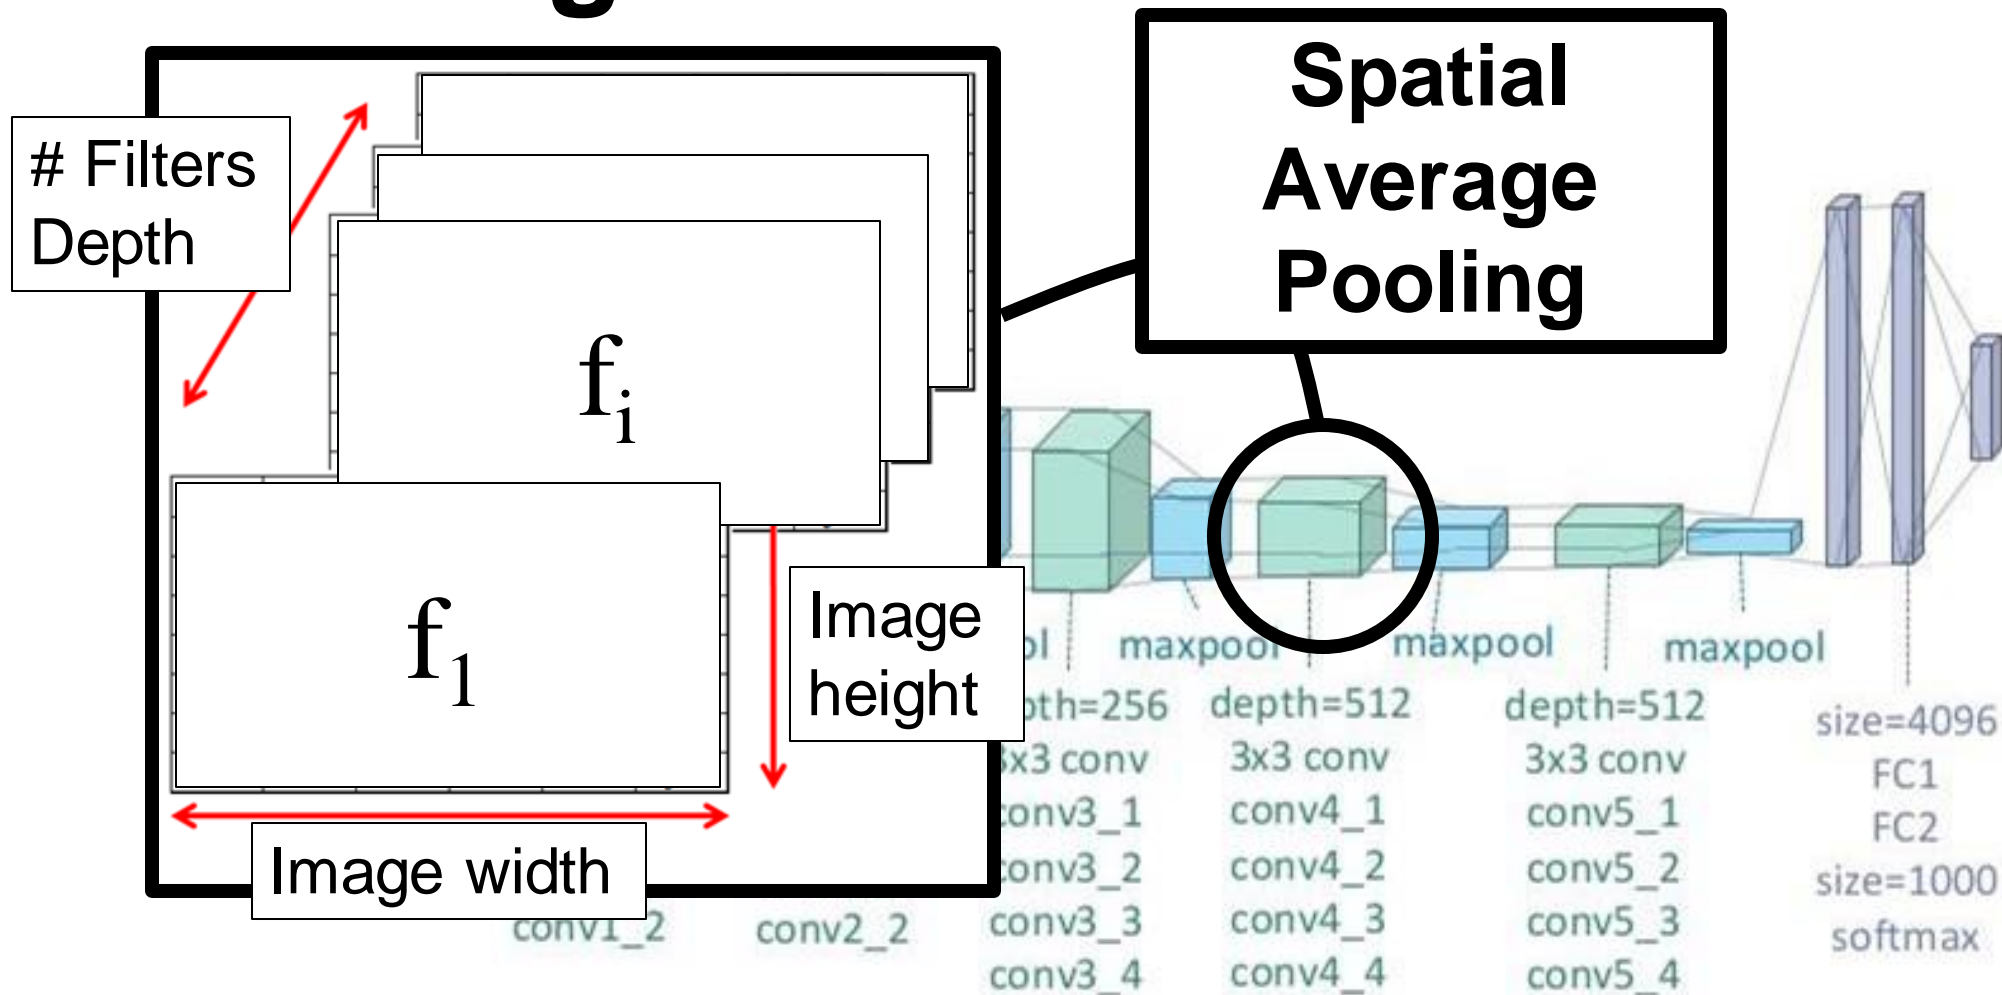
Knowledge **inside** DNN




Knowledge **inside** DNN




Knowledge **inside** DNN




Knowledge **inside** DNN




Knowledge **inside** DNN




Knowledge **inside** DNN



Knowledge **inside** DNN

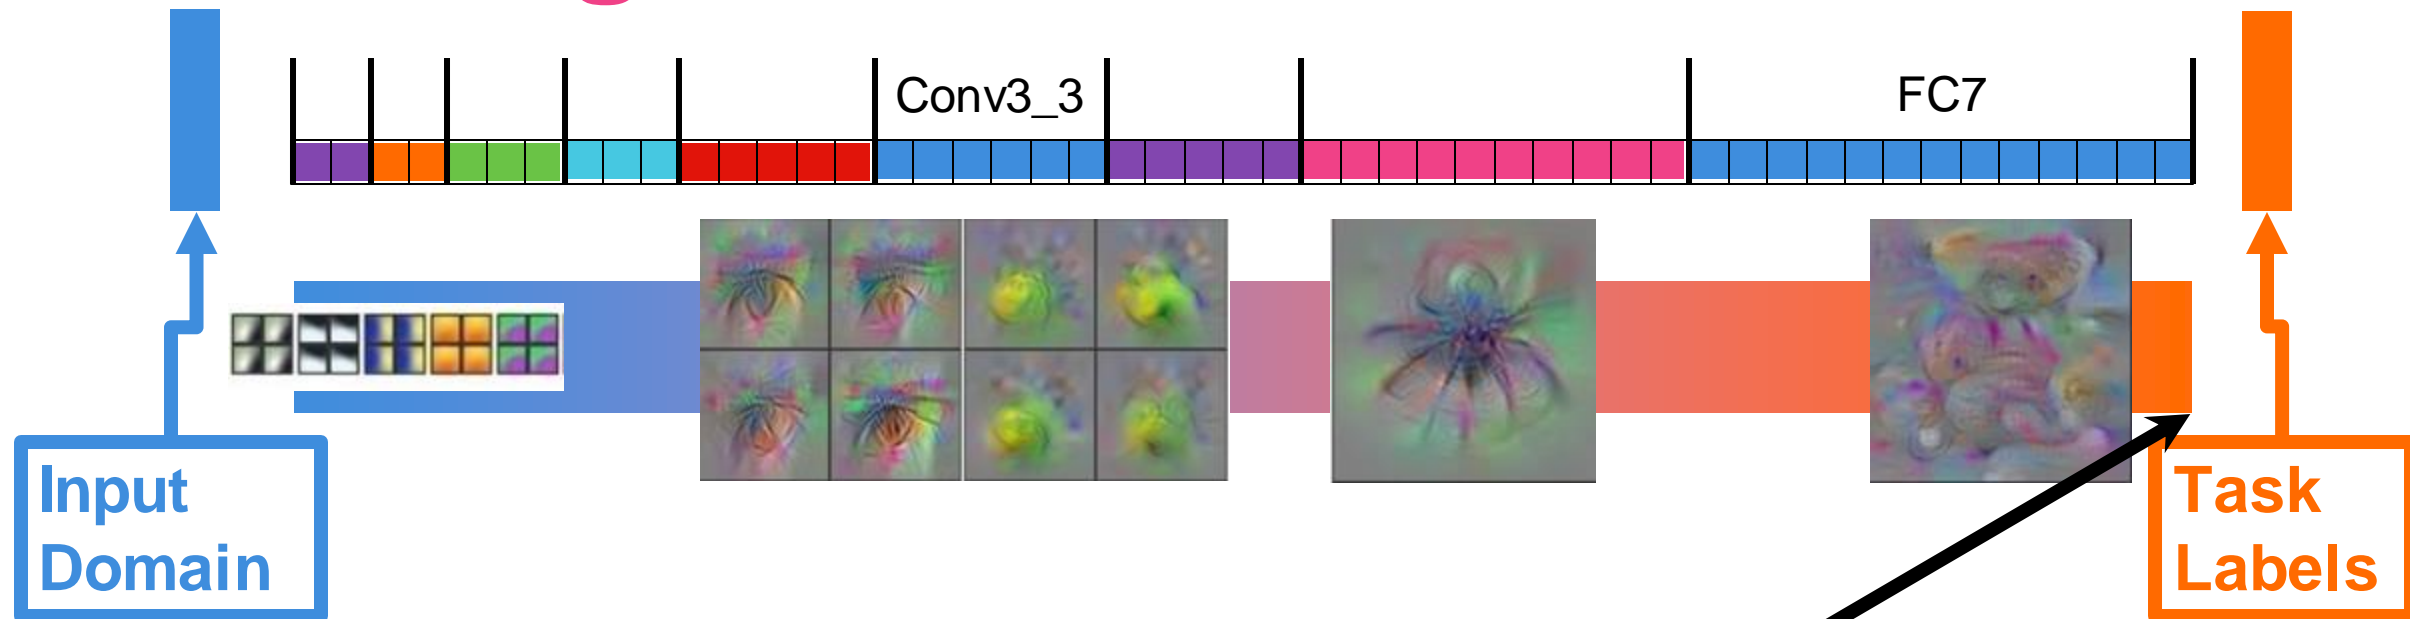


Knowledge **inside** DNN

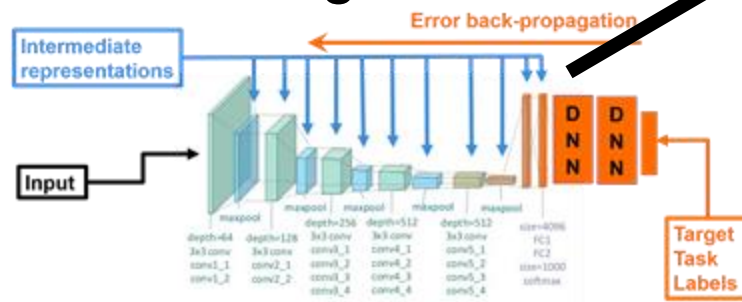


Visualizations from: Yosinski, Jason, et al. "Understanding neural networks through deep visualization." *arXiv preprint arXiv:1506.06579* (2015).


Fine tuning inside DNN?




Fine-tuning



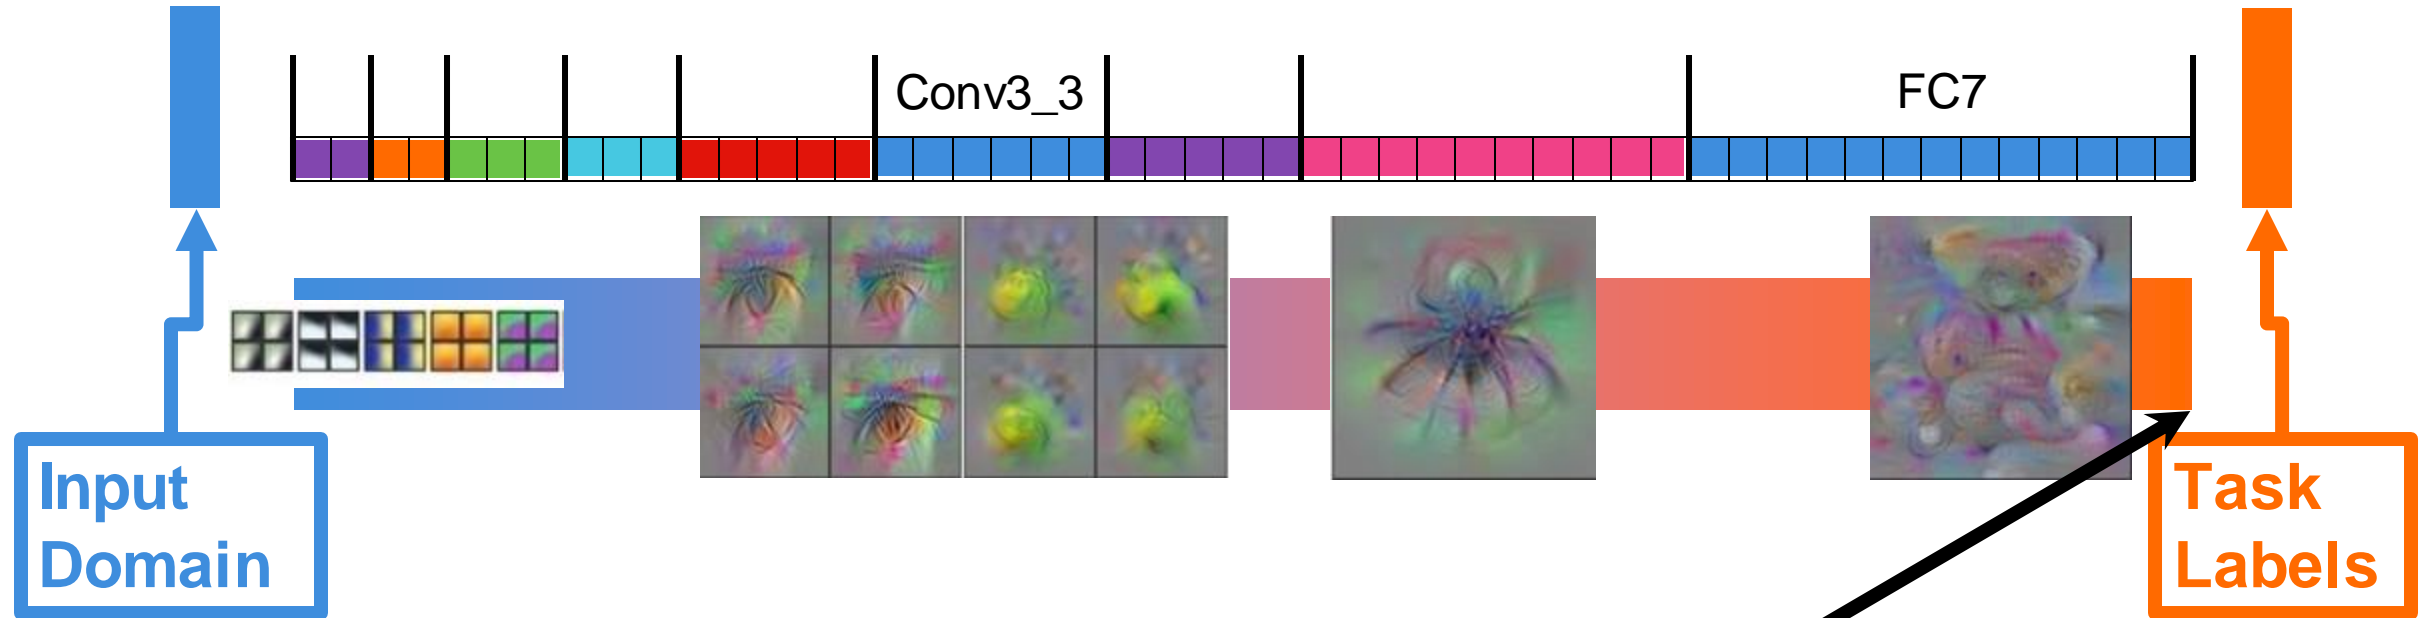
Fine tuning inside DNN?




Fine tuning inside DNN?



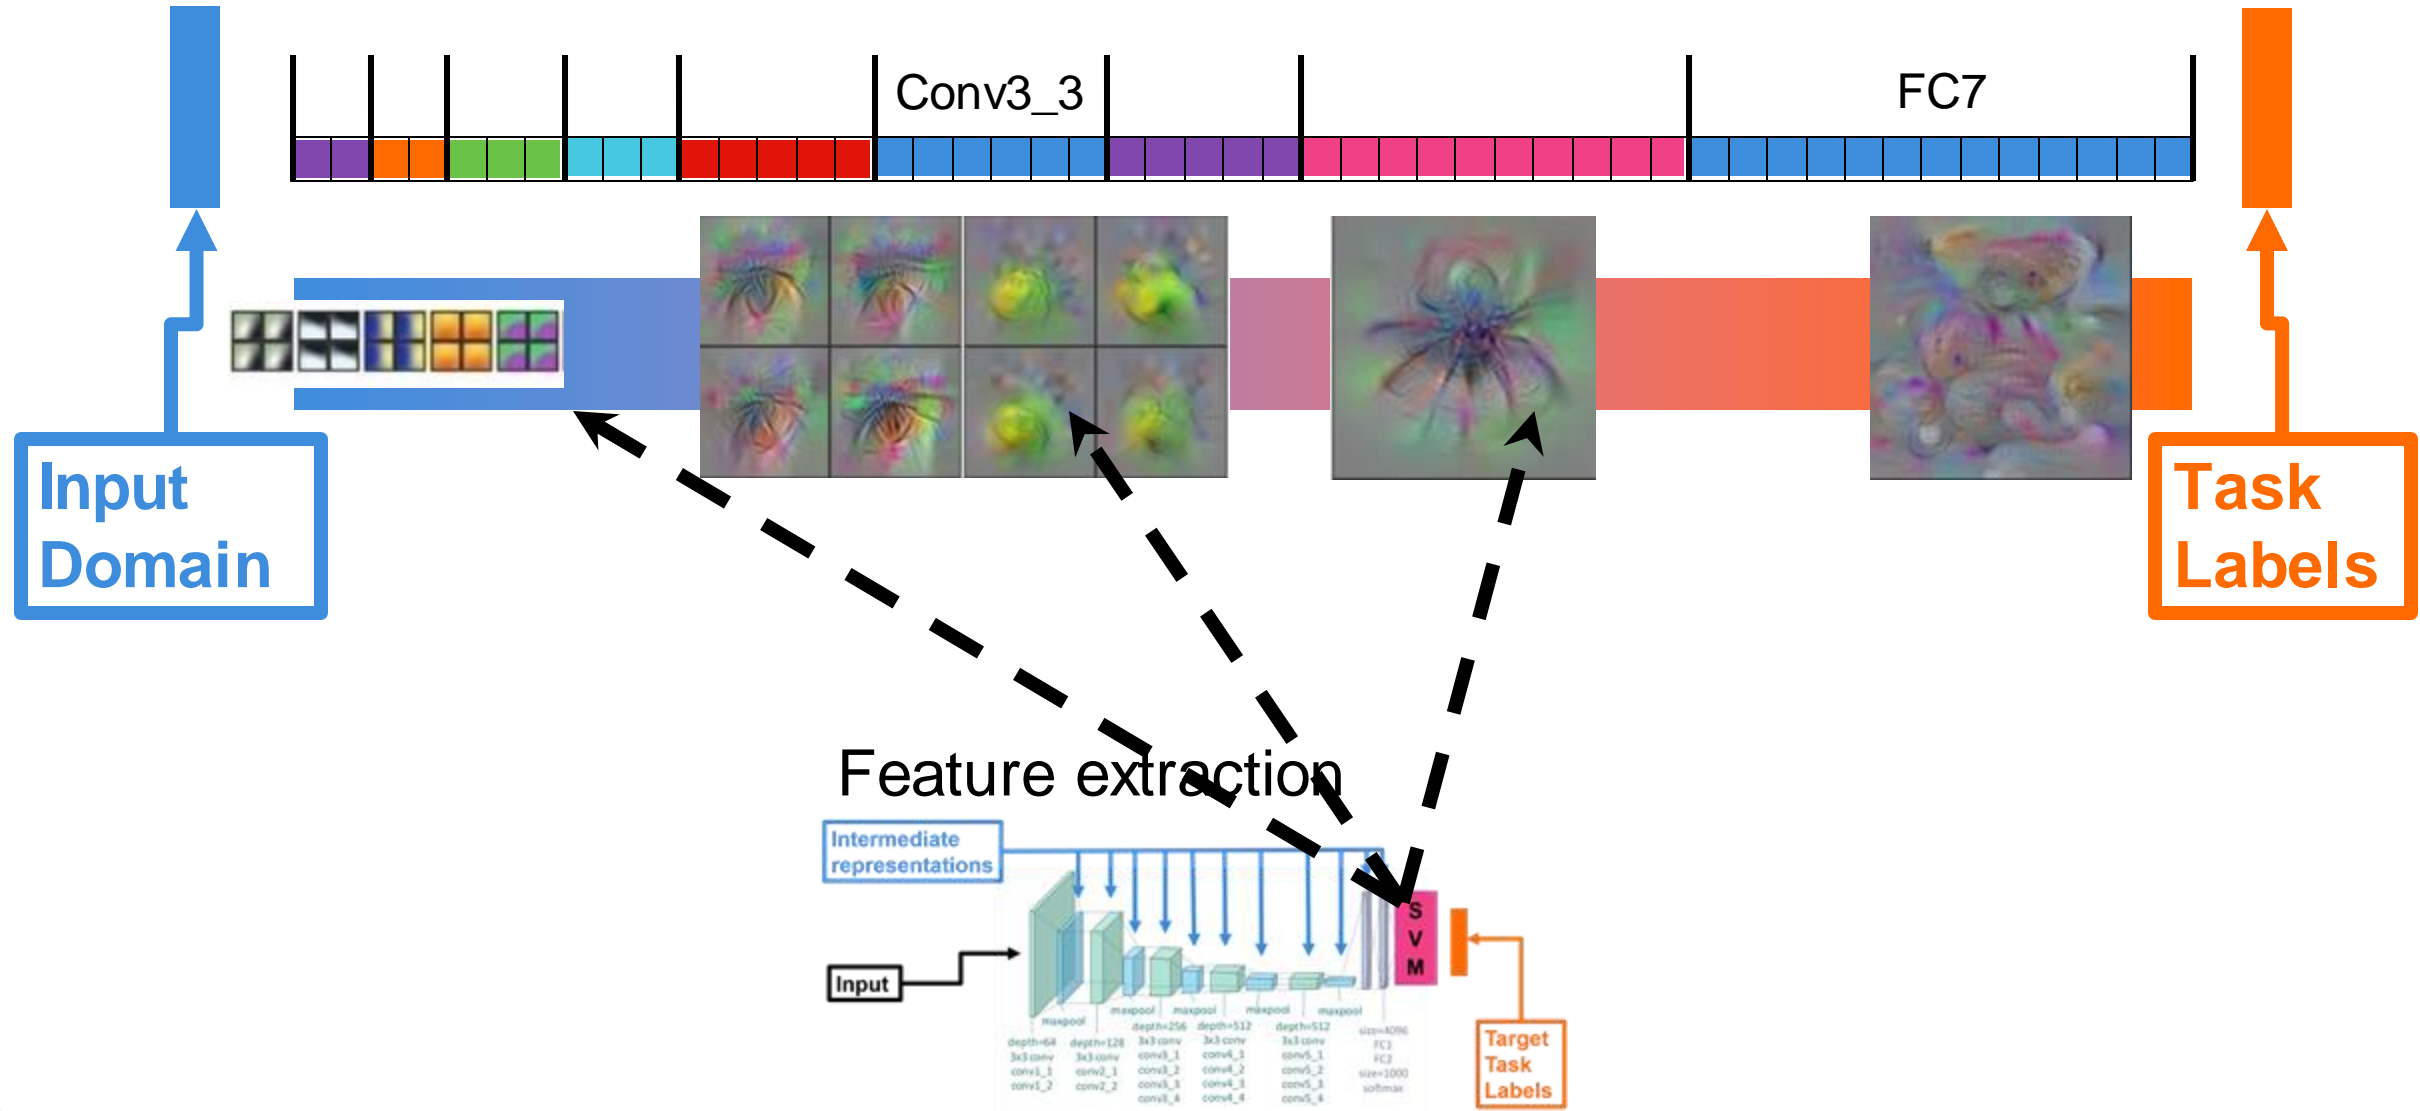
Feature extraction inside DNN?



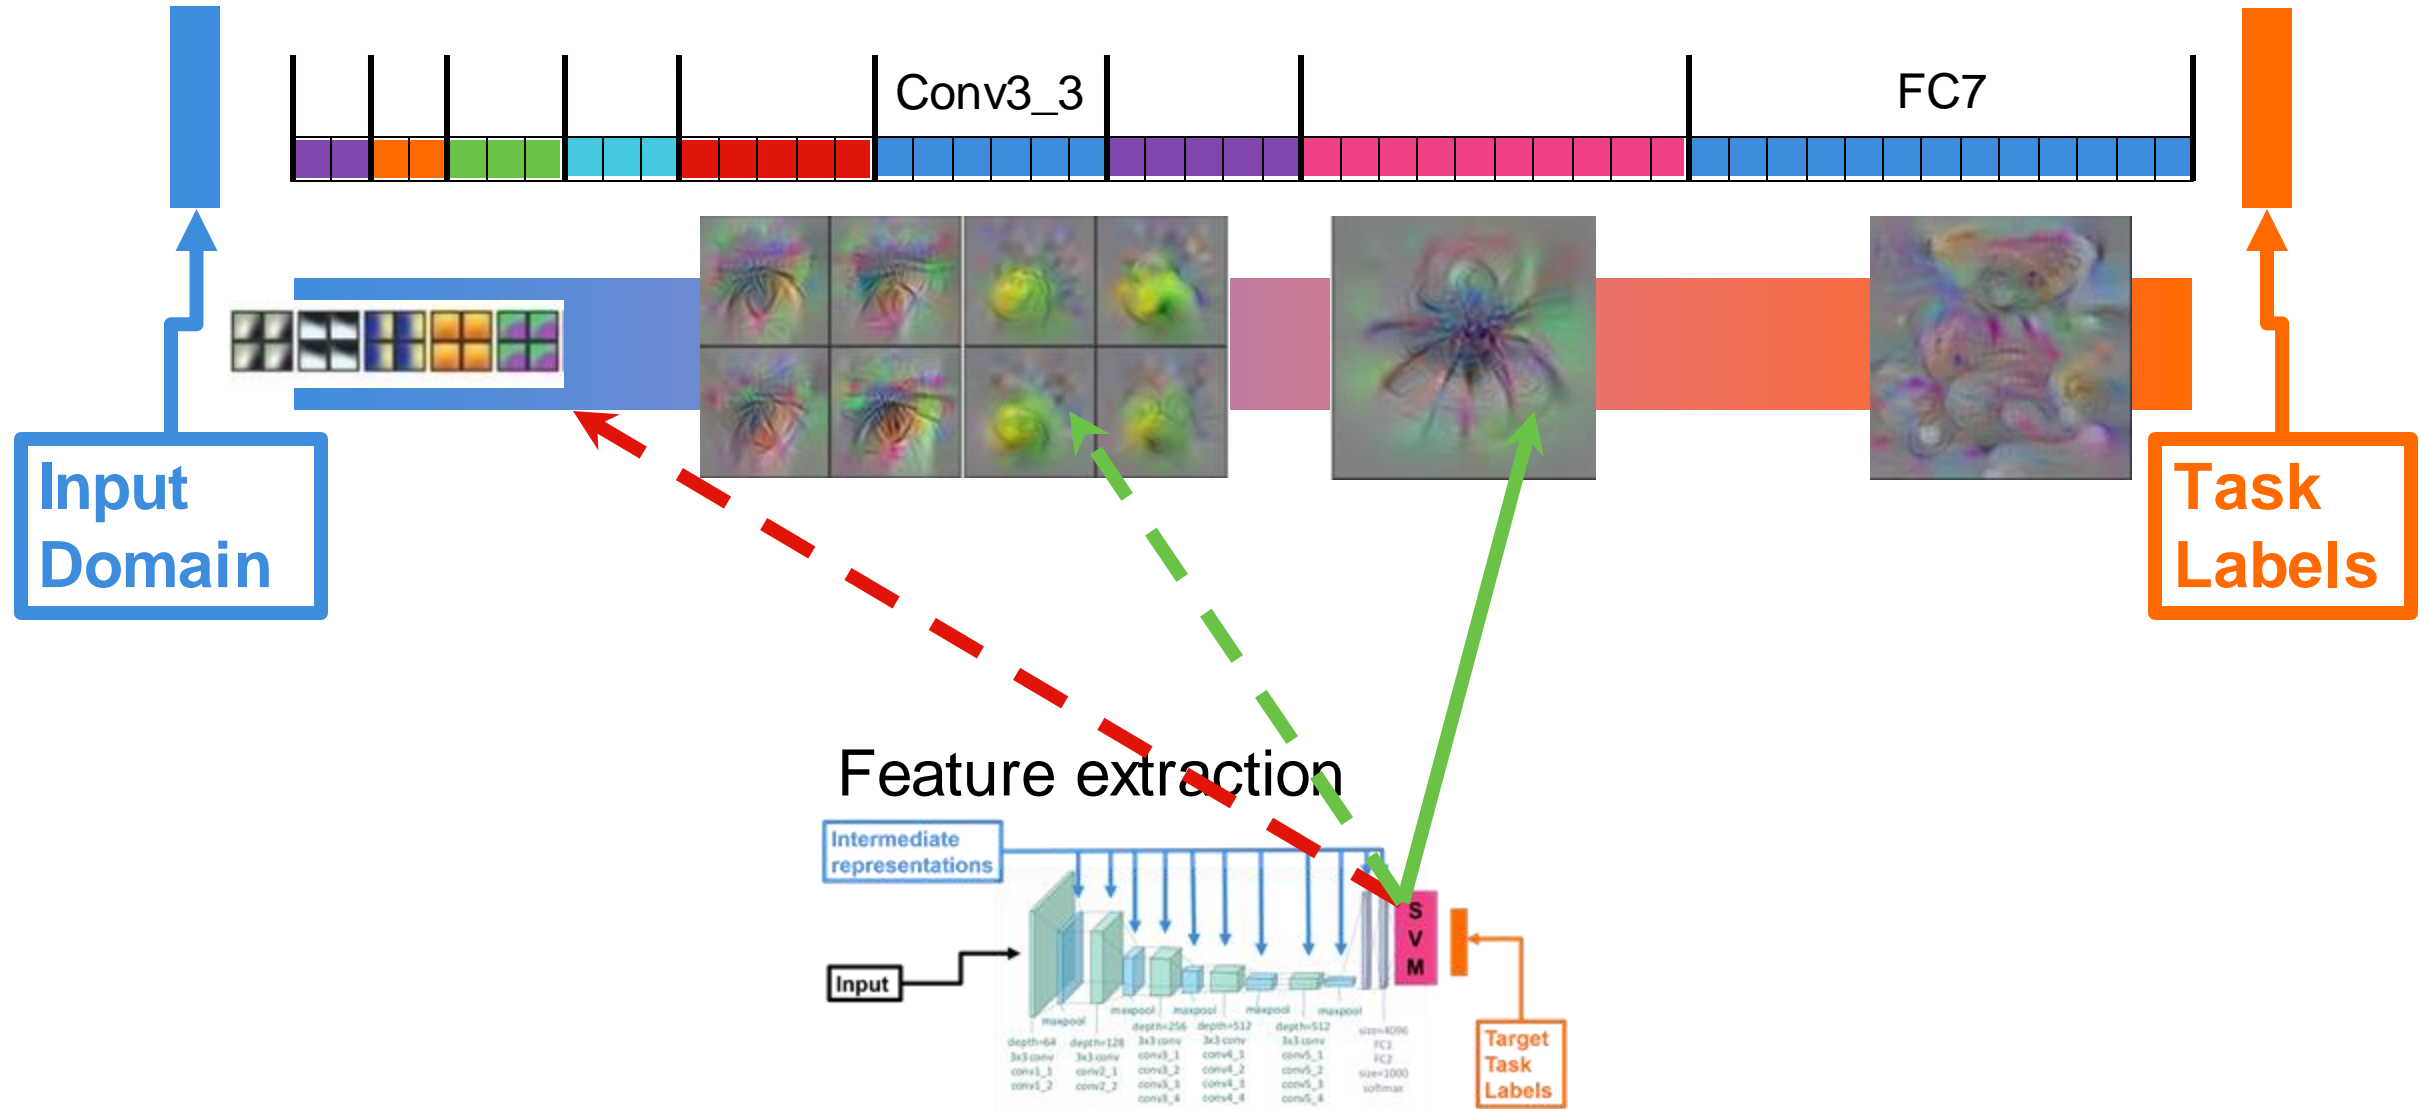
Feature extraction



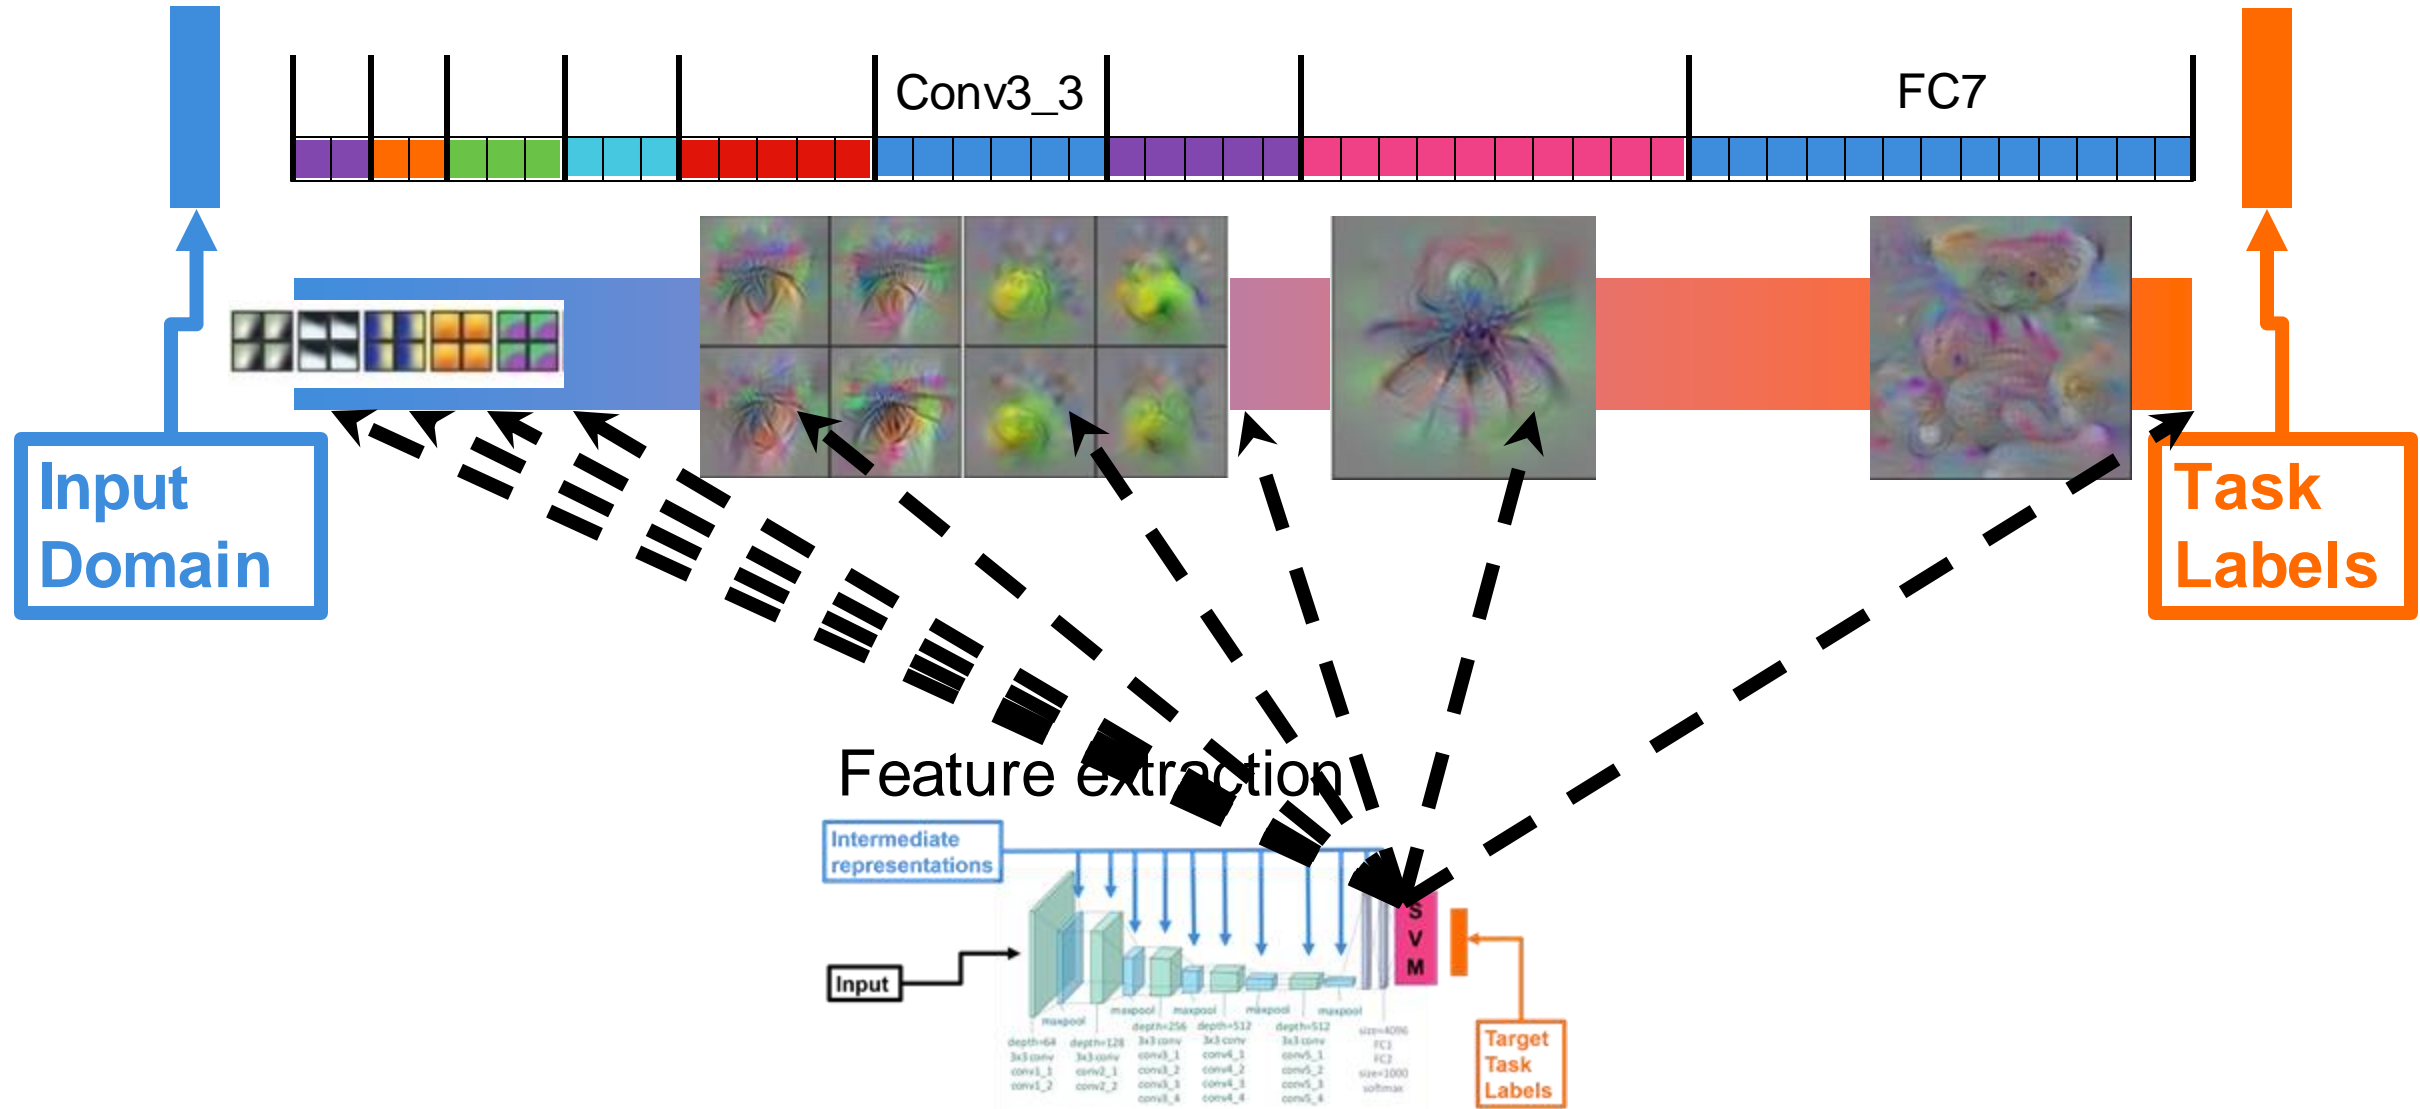
Feature extraction inside DNN?



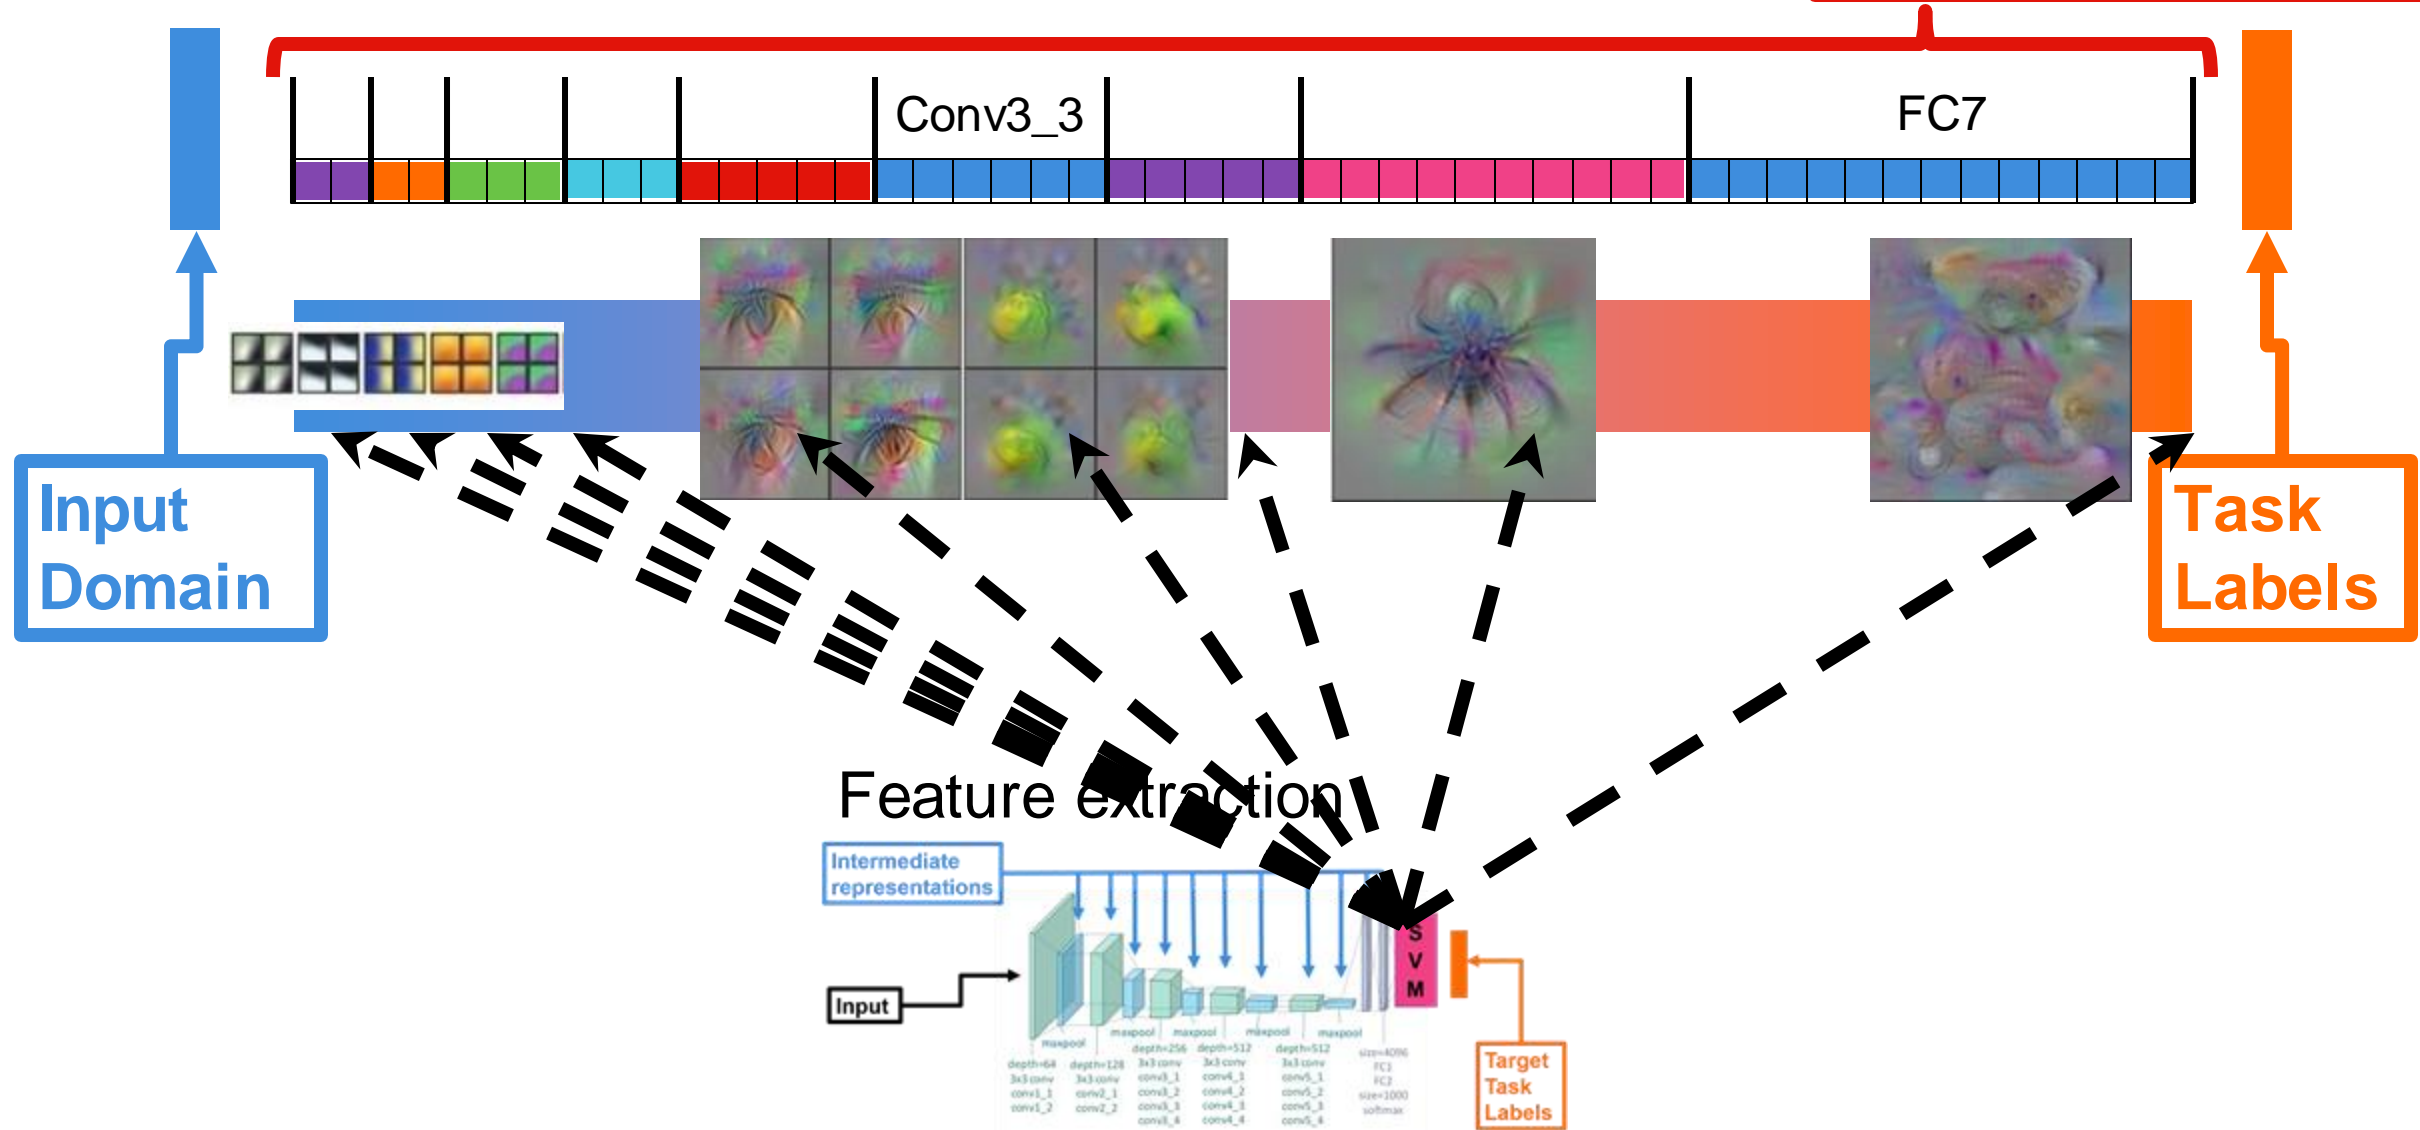
Feature extraction inside DNN?



Feature extraction inside DNN?

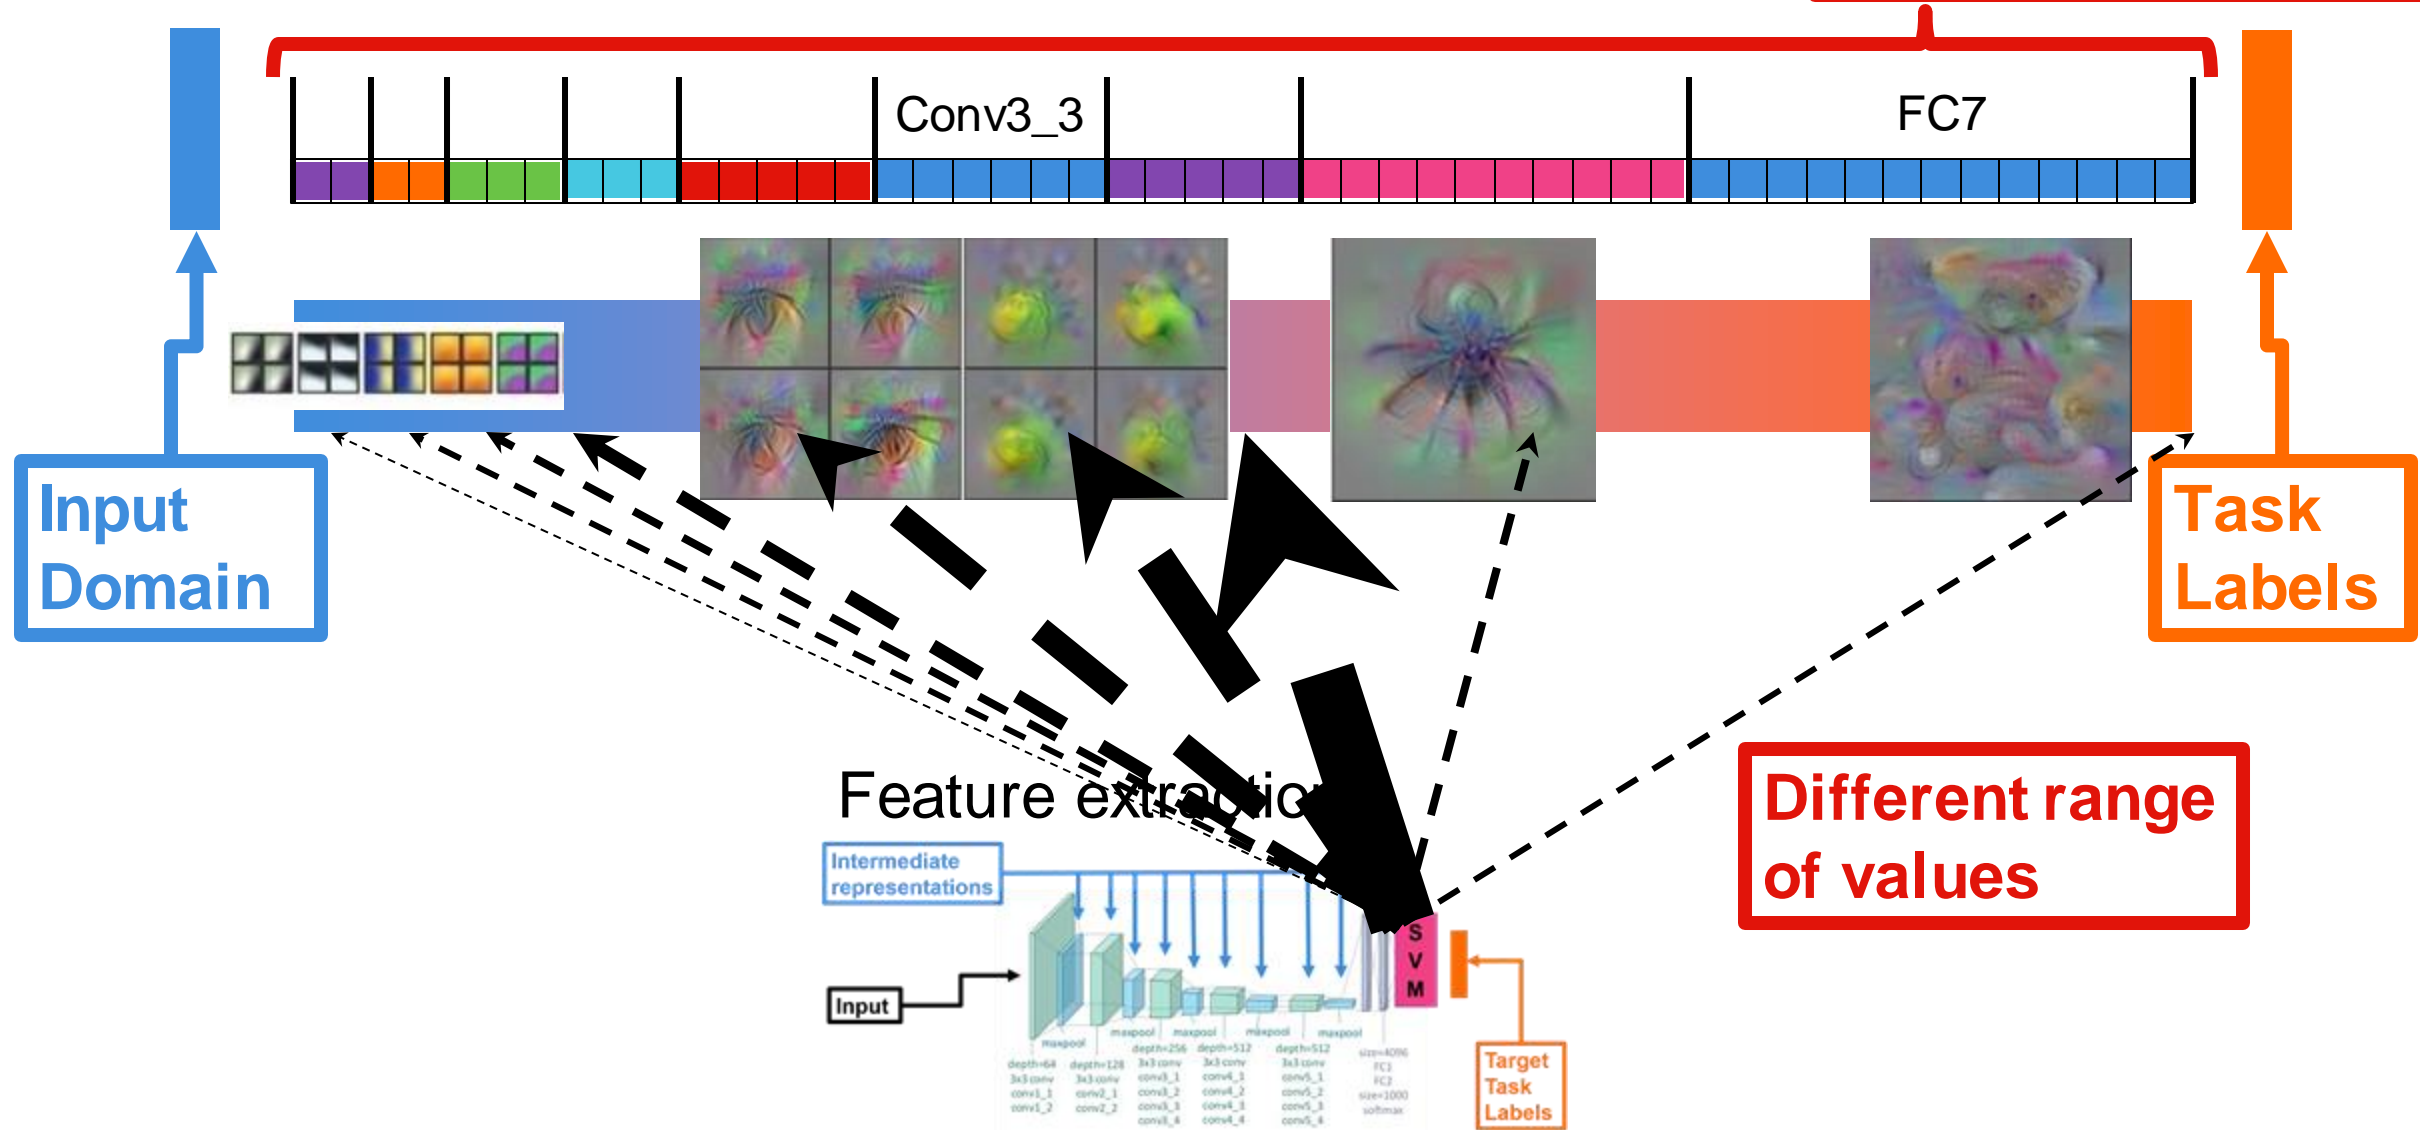


Feature extraction inside DNN? VGG16 dim: 12,416


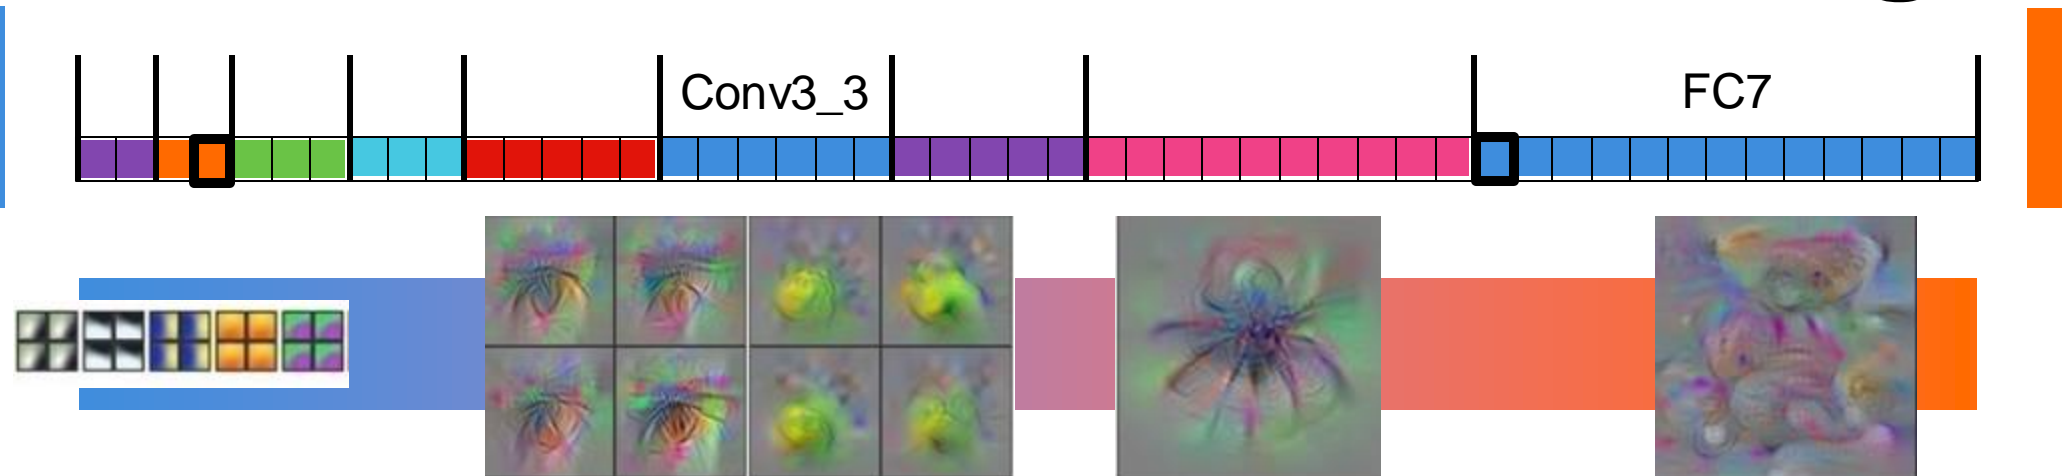


Feature extraction inside DNN?


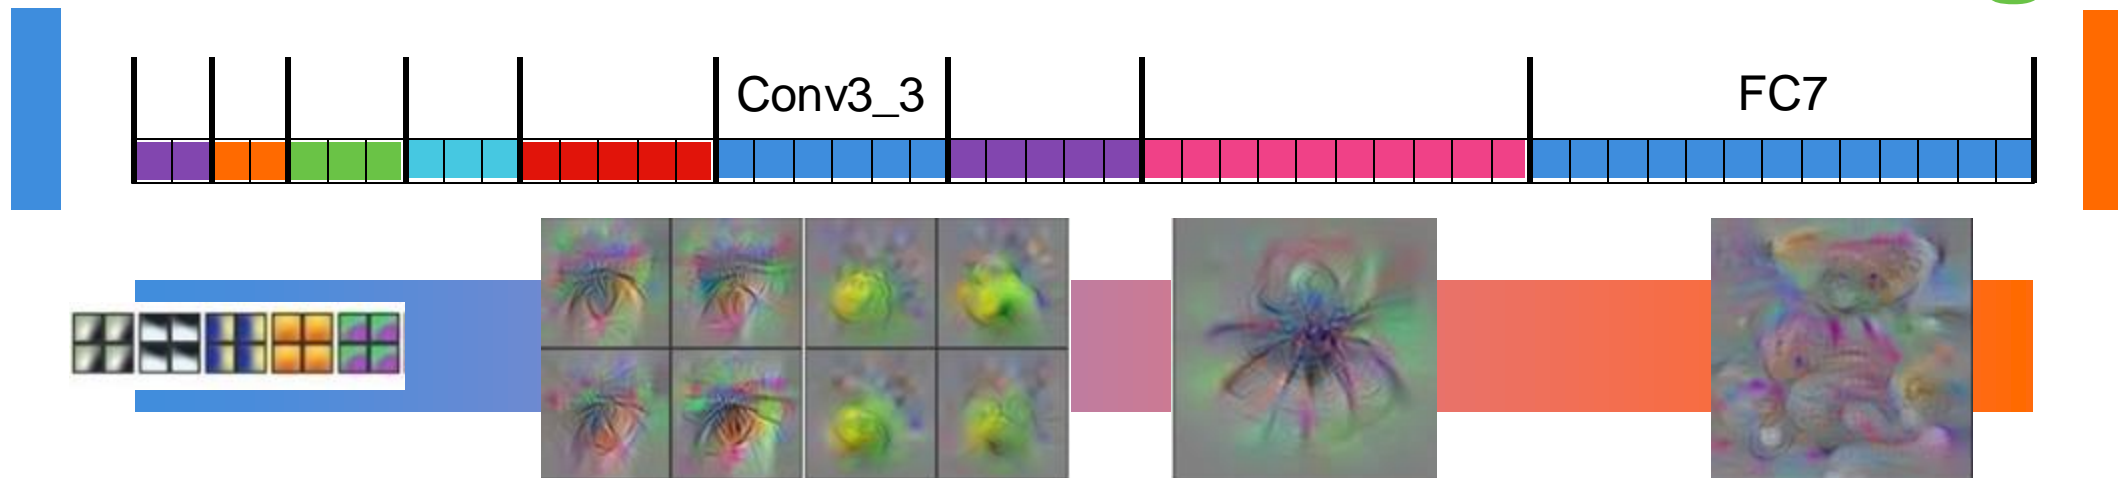
VGG16 dim: 12,416



Feature behavior in Transfer Learning




Standardization: features in same range




Each feature independently

Standardization: features in same range




Each feature independently

Standardization: features in same range



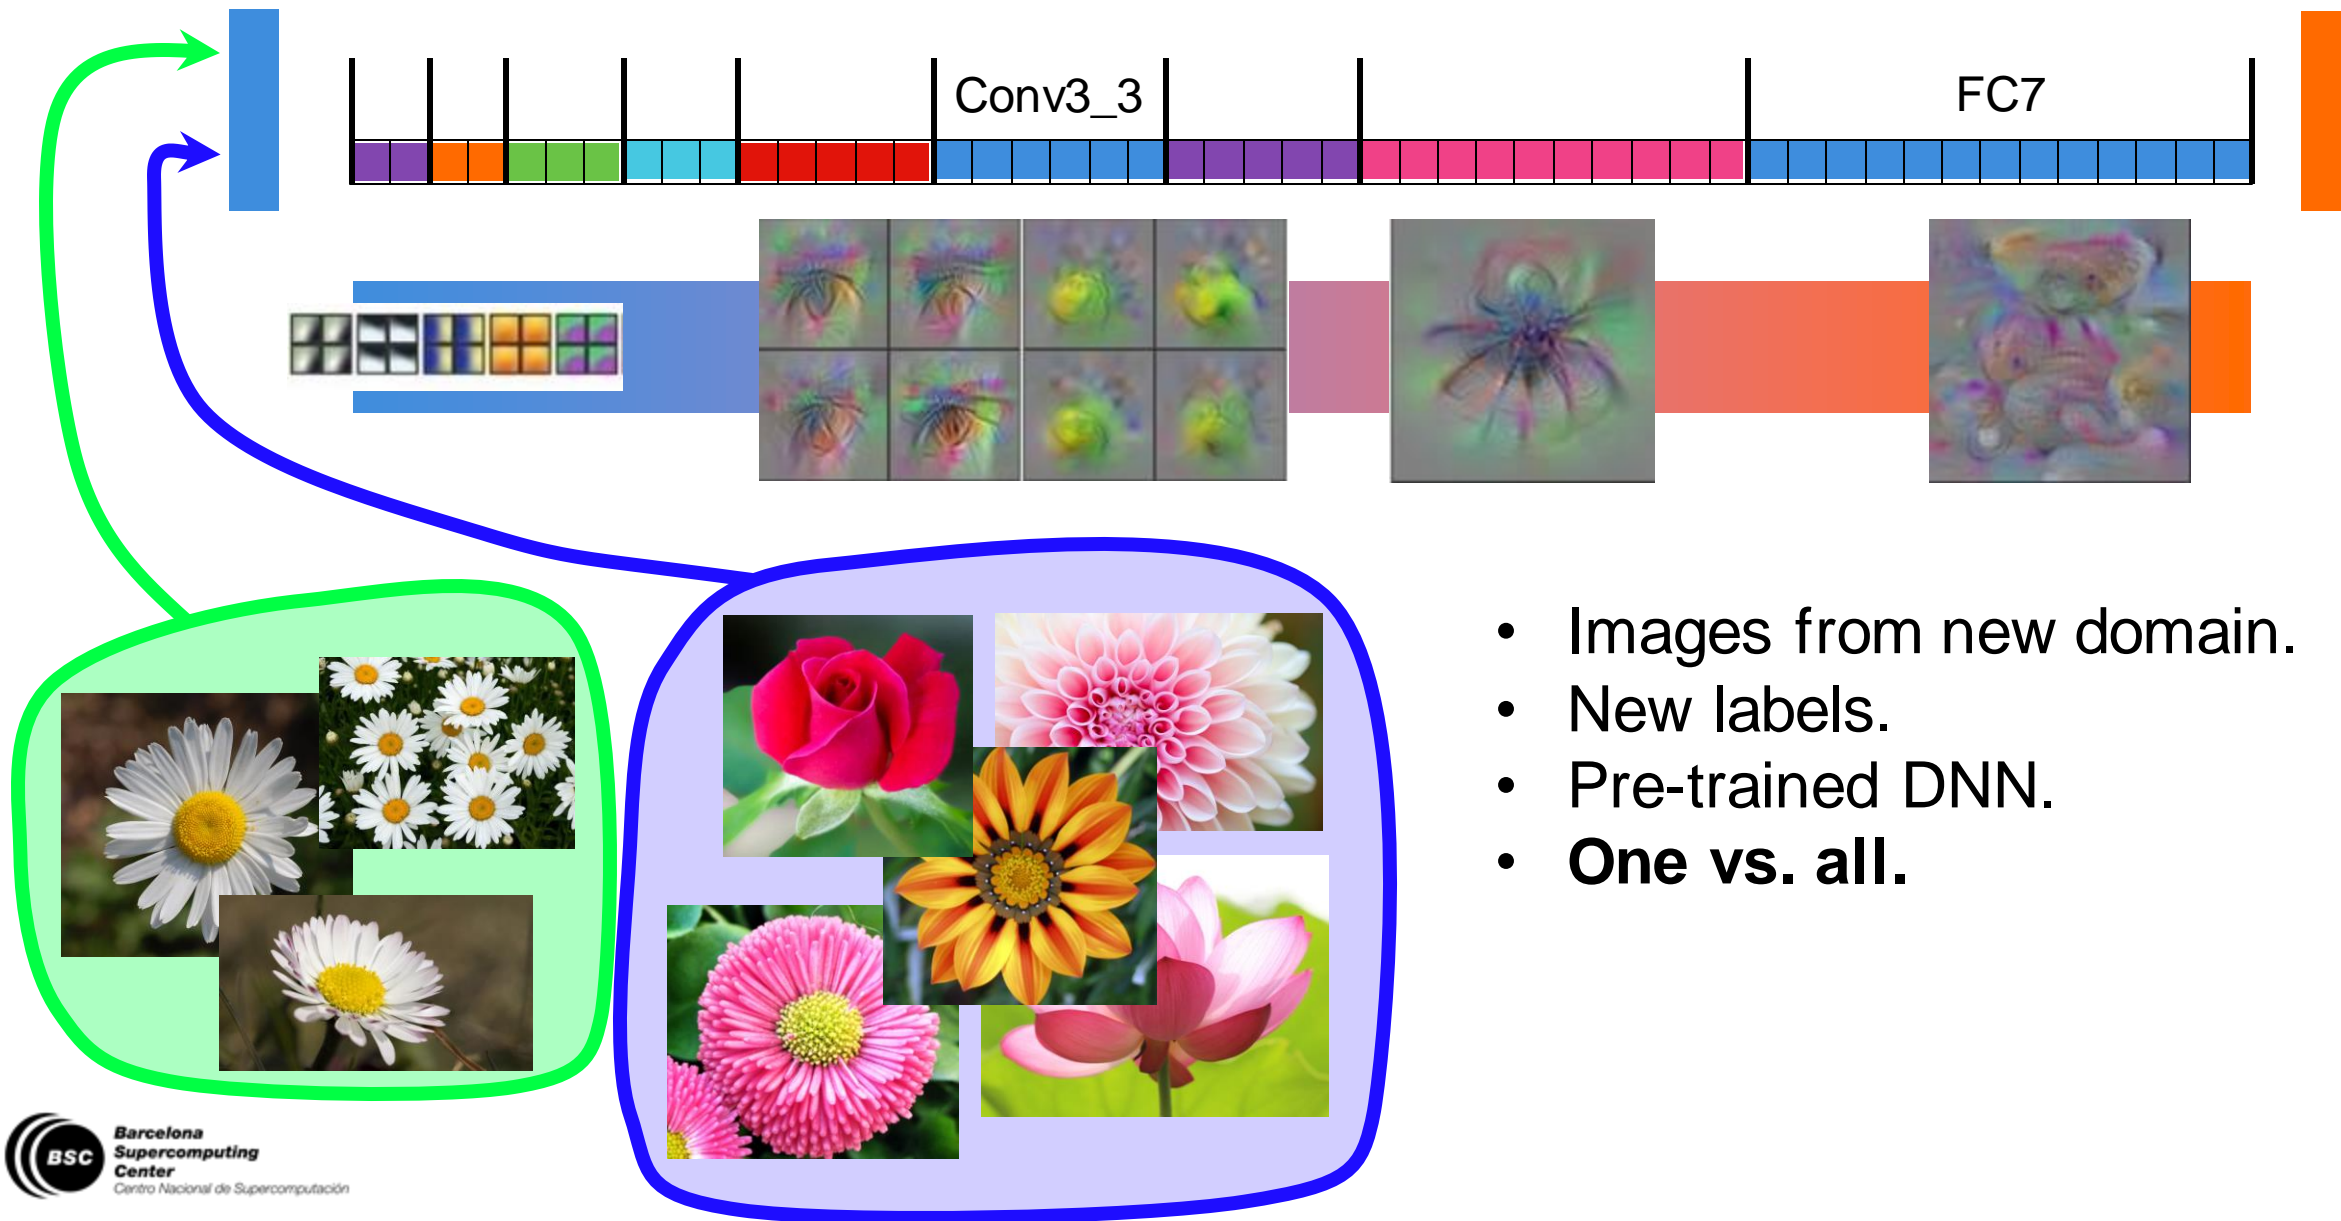
Each feature independently

Standardization: features in same range

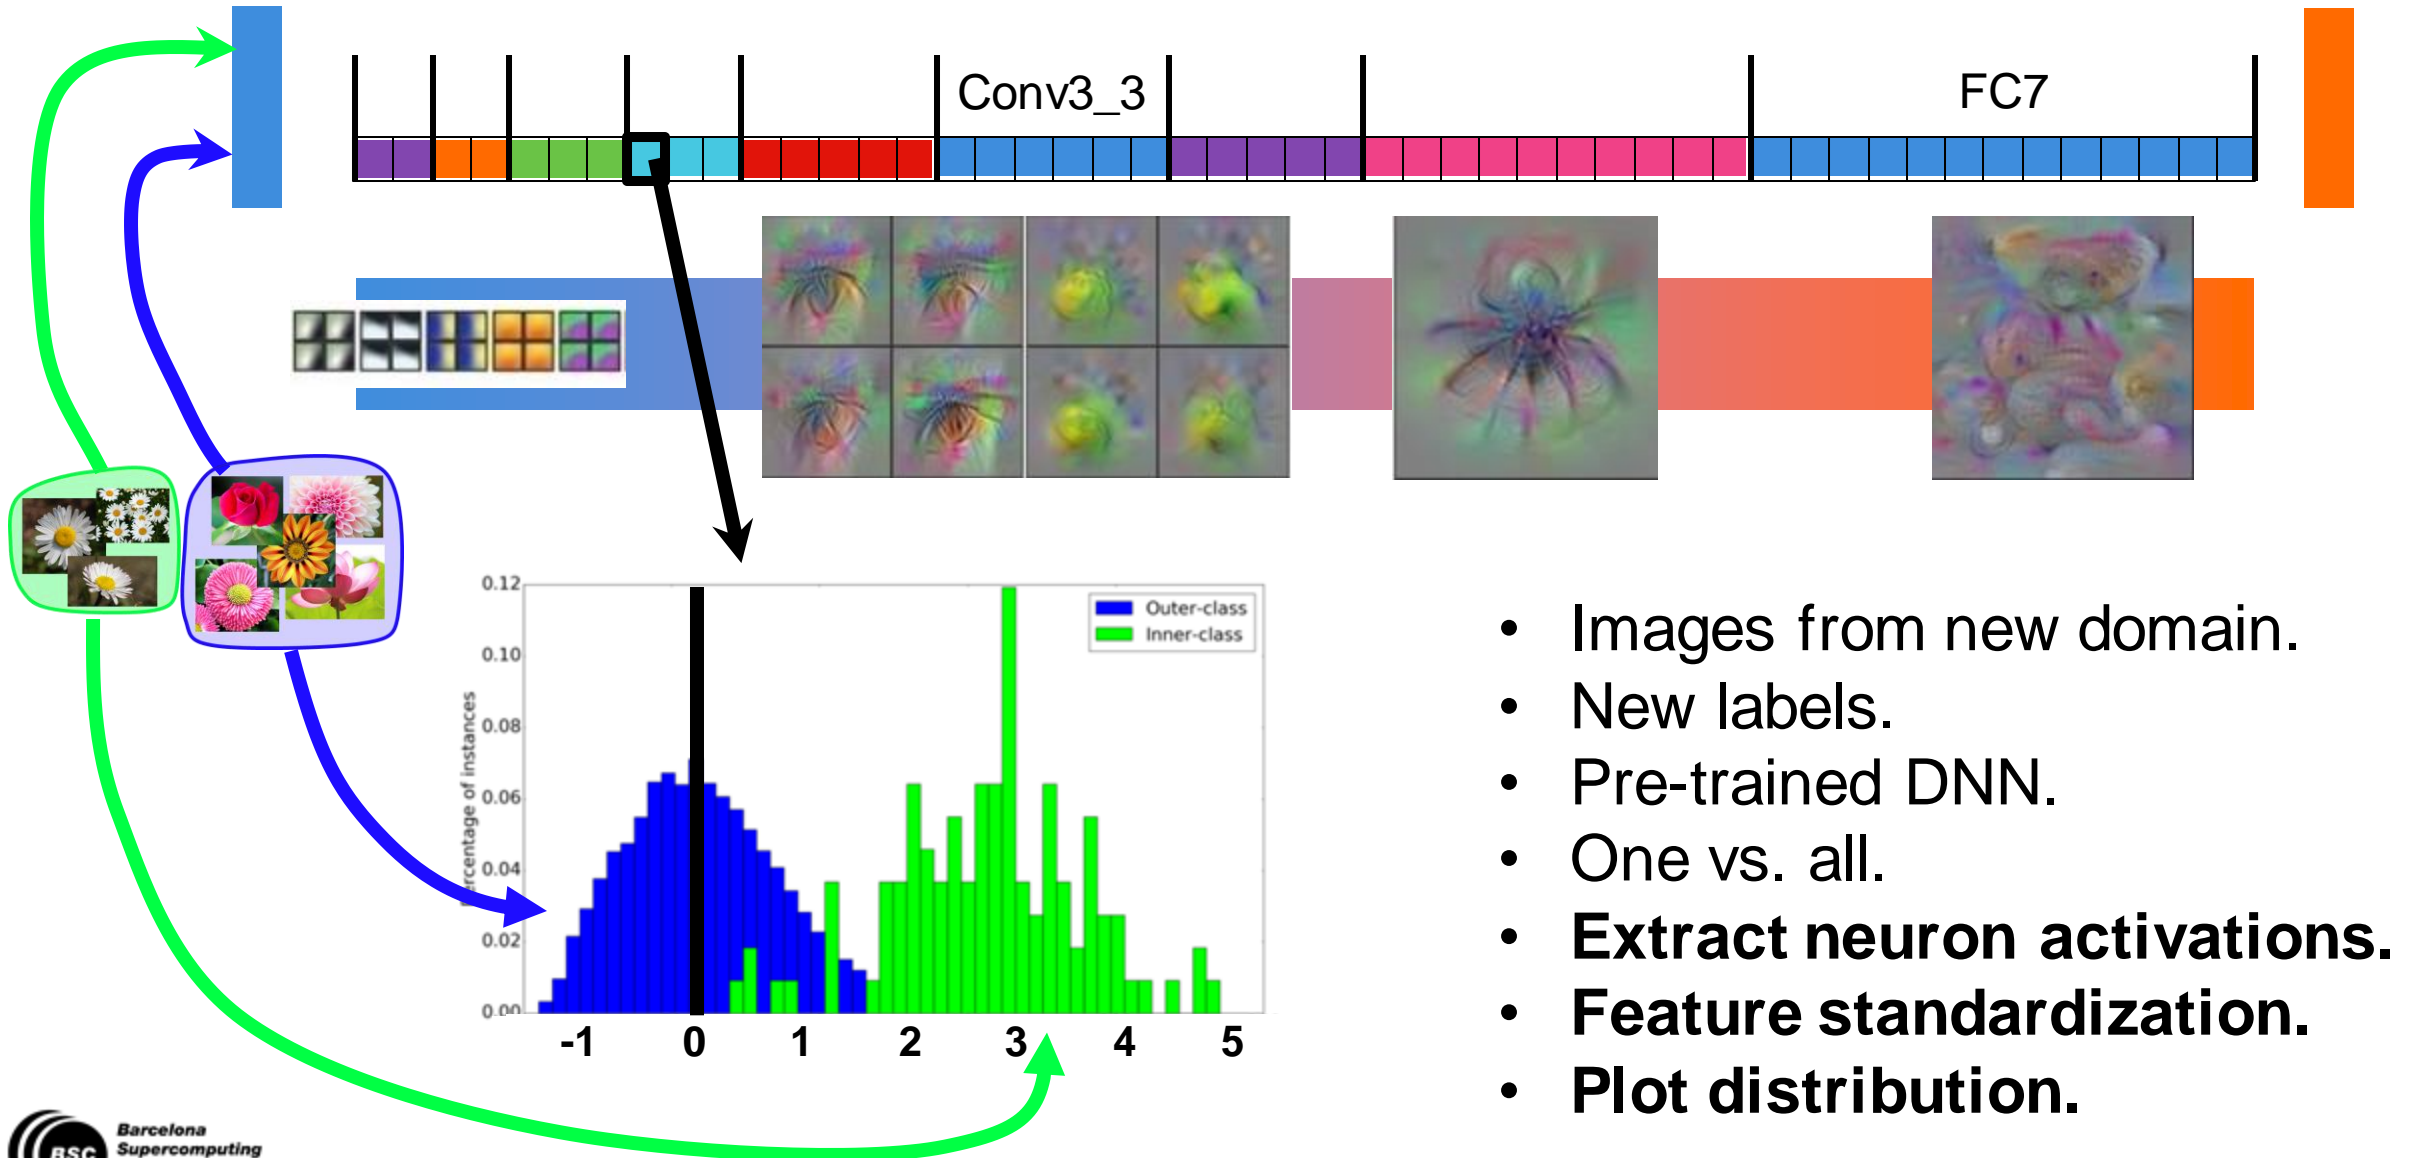


Each feature independently

Standardization: features in context




Standardization: features in context




Standardization: features in context

High activation features for a specific class



Low activation features for a specific class



conv3_4 n202



Greenhouse

fc7 n1779



Cloister

MIT-67 - indoors

fc7 n1946



Gadwall



Brown Pelican



White Pelican




Heermann Gull

CUB-200 - birds

Standardization: features in context

High activation features for a specific class



conv3_4 n202



Greenhouse


fc7 n1779



Cloister

MIT-67 - indoors

Low activation features for a specific class



fc7 n1449



35 alpine sea holly




10 globe thistle




14 spear thistle

Flowers-102


Curse of dimensionality




Curse of dimensionality



Curse of dimensionality

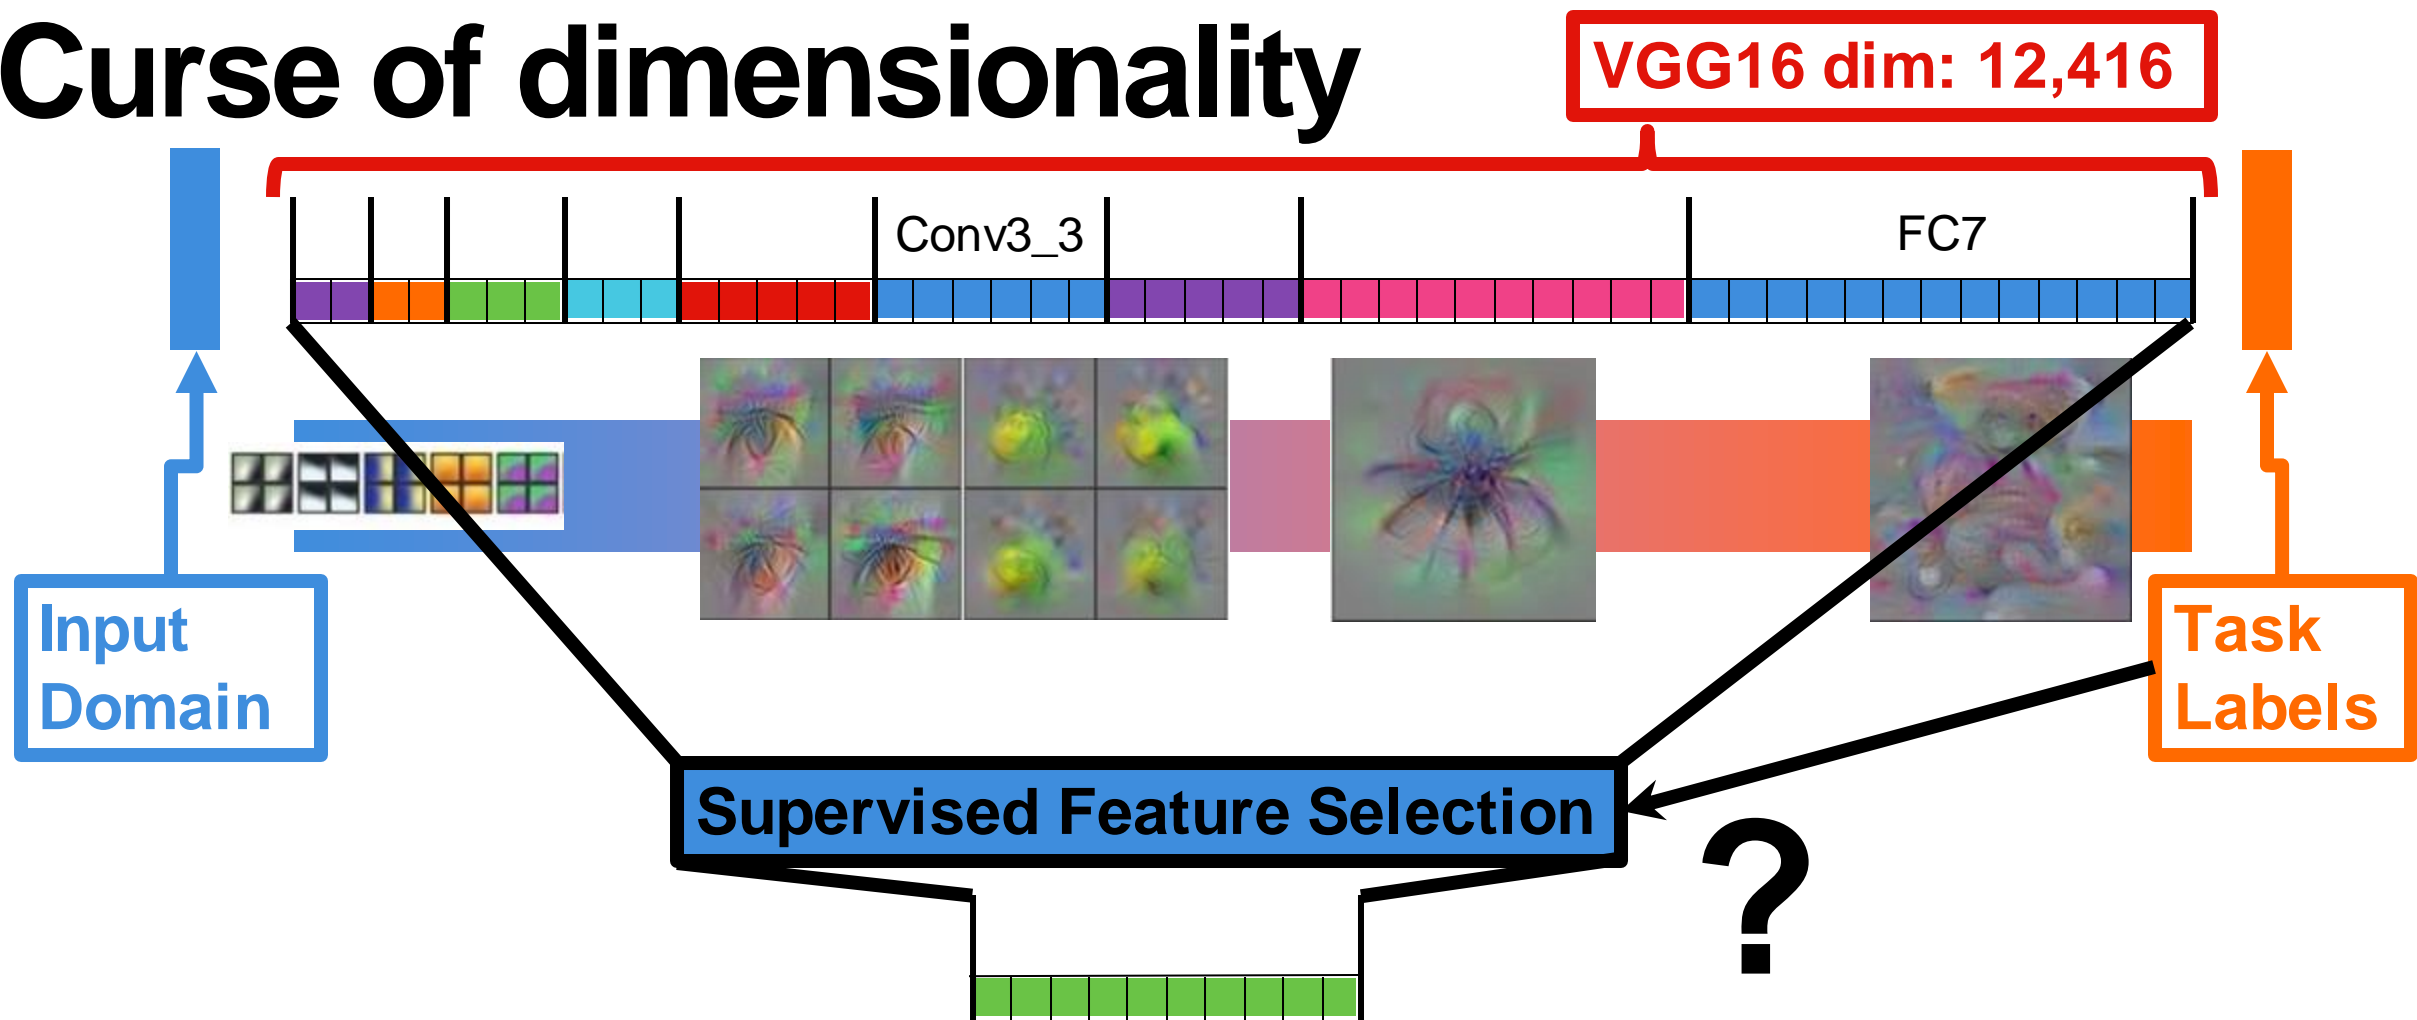


Curse of dimensionality

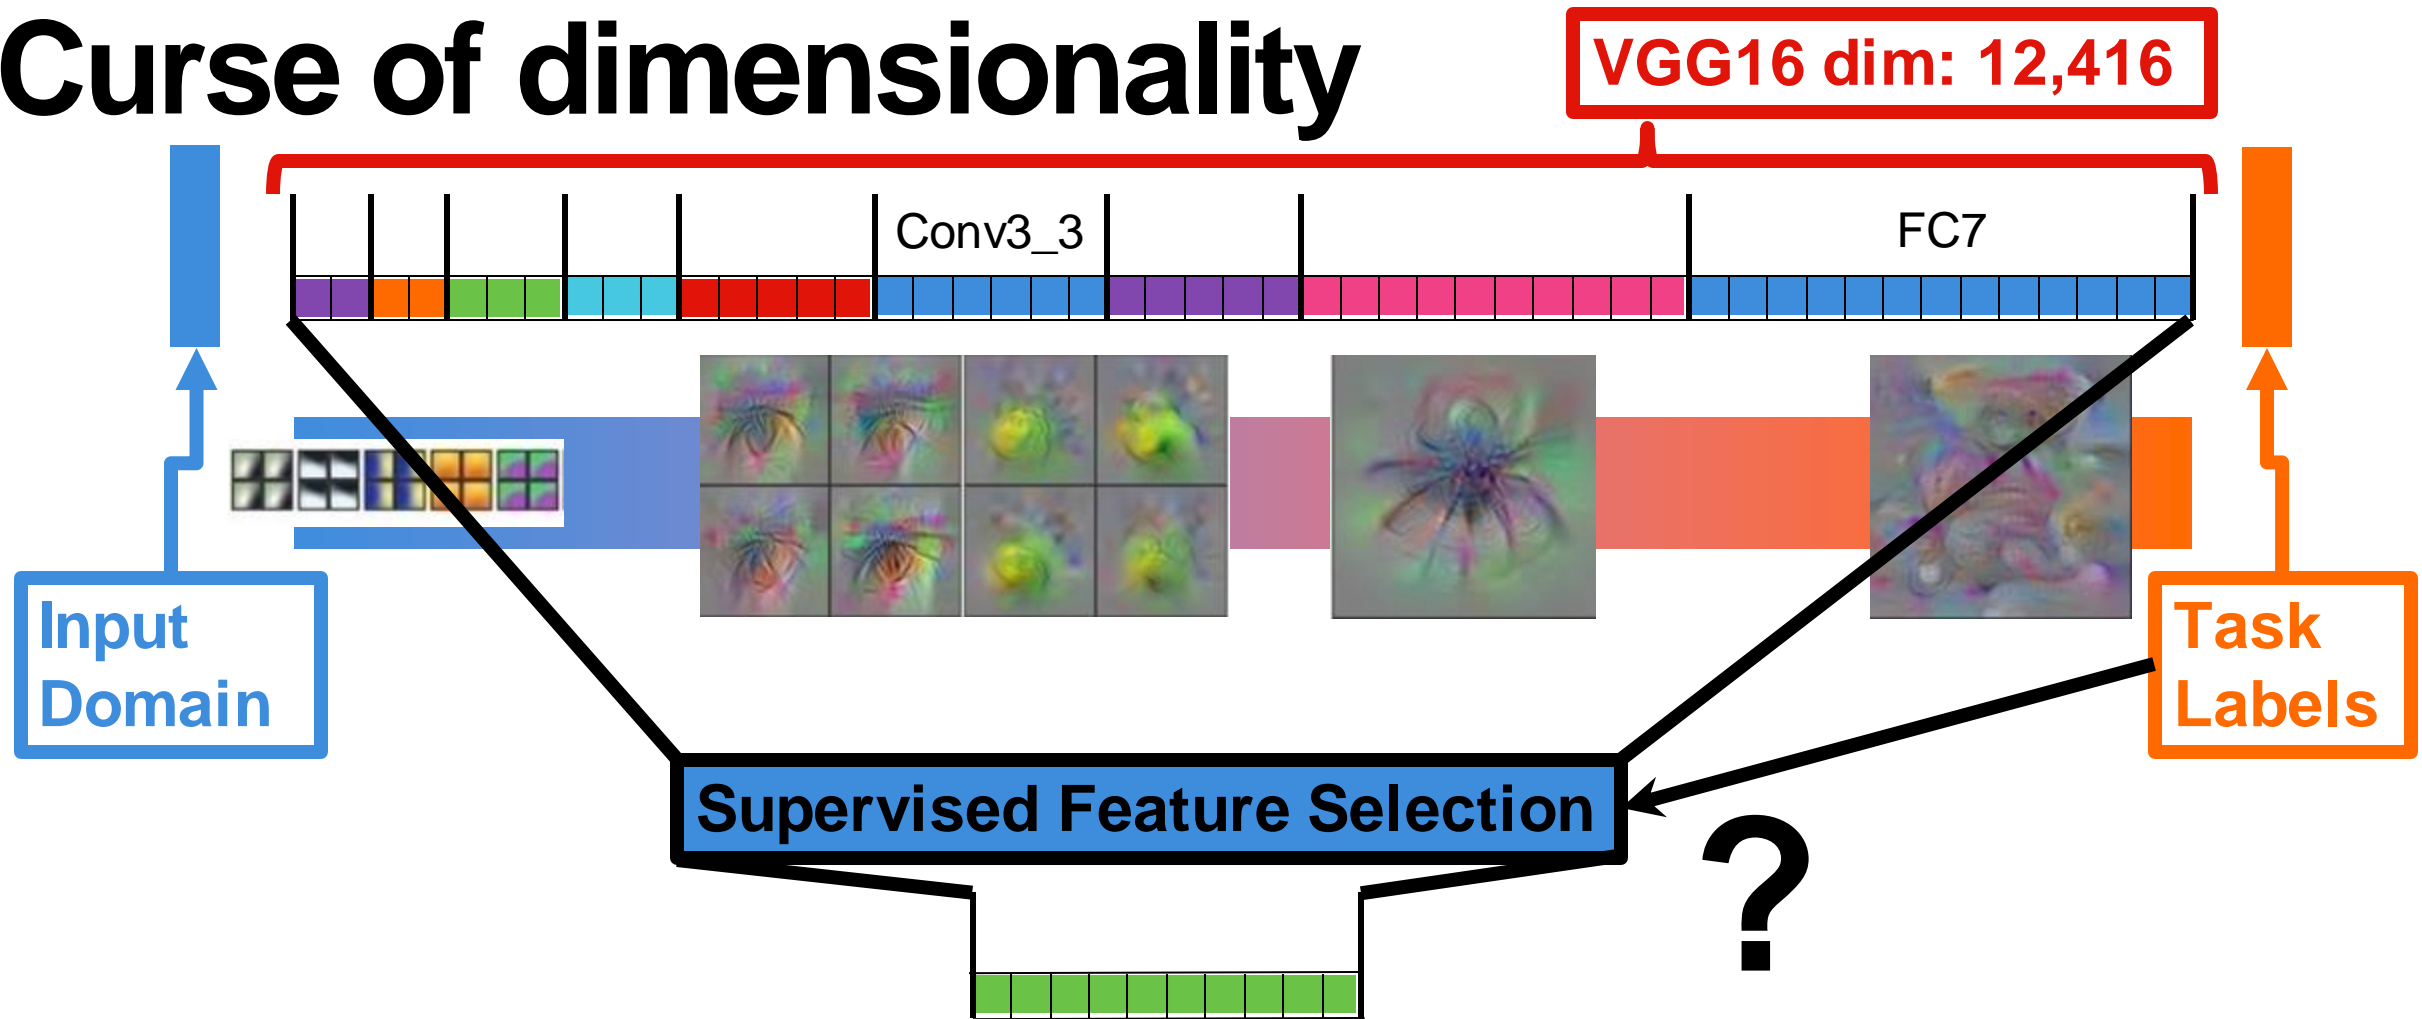


- Unsupervised:
Only takes into account
the **domain**

Curse of dimensionality

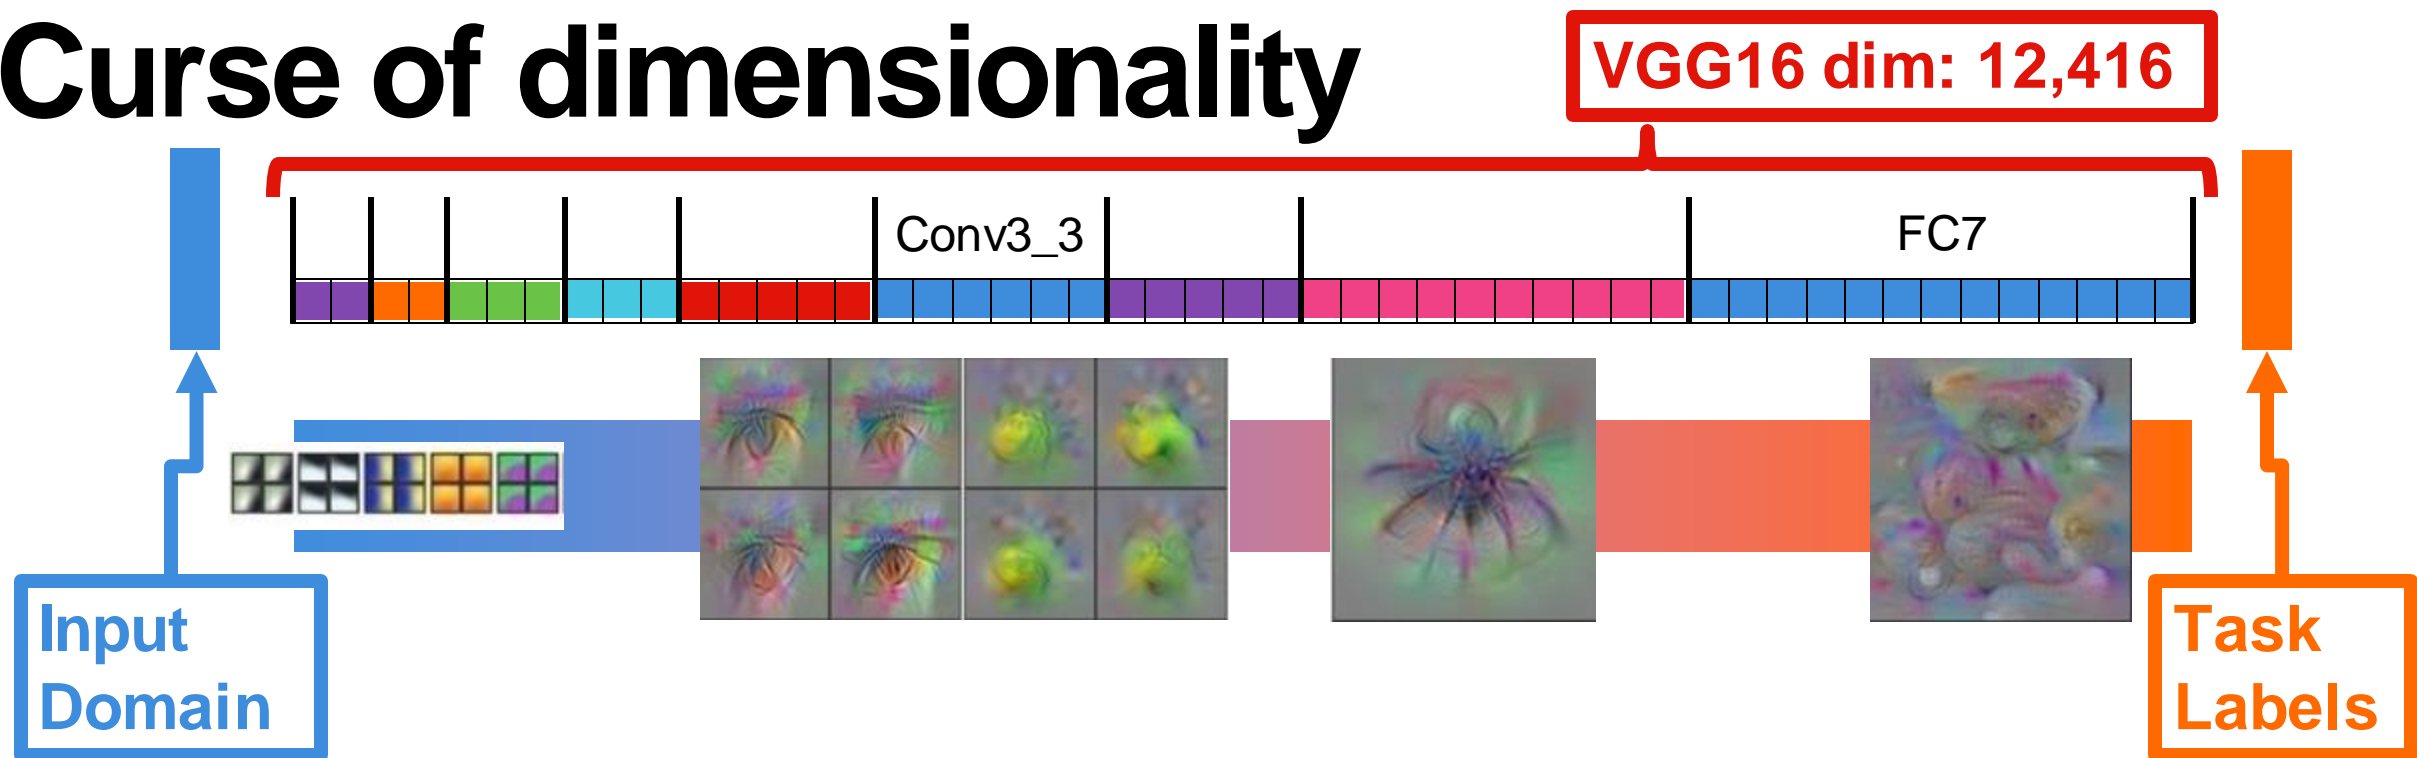


Curse of dimensionality



- High computational **cost!**


Curse of dimensionality



VGG16 dim: 12,416

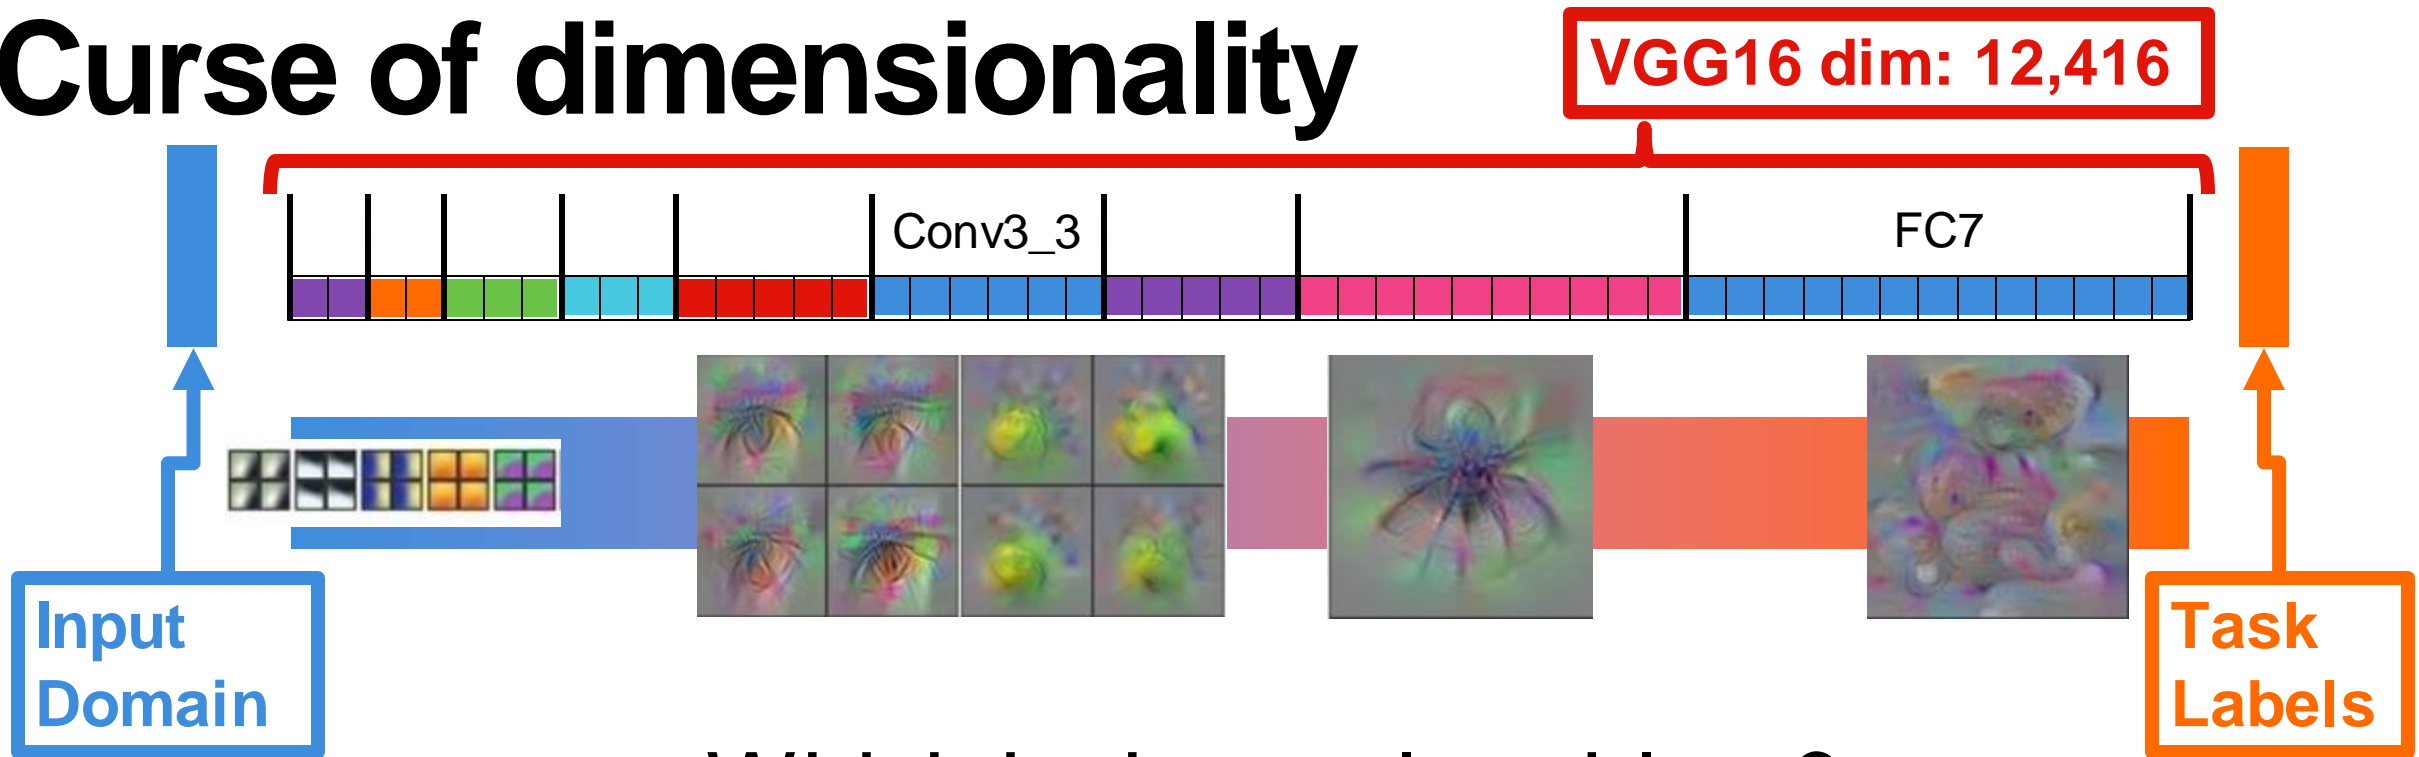


Curse of dimensionality



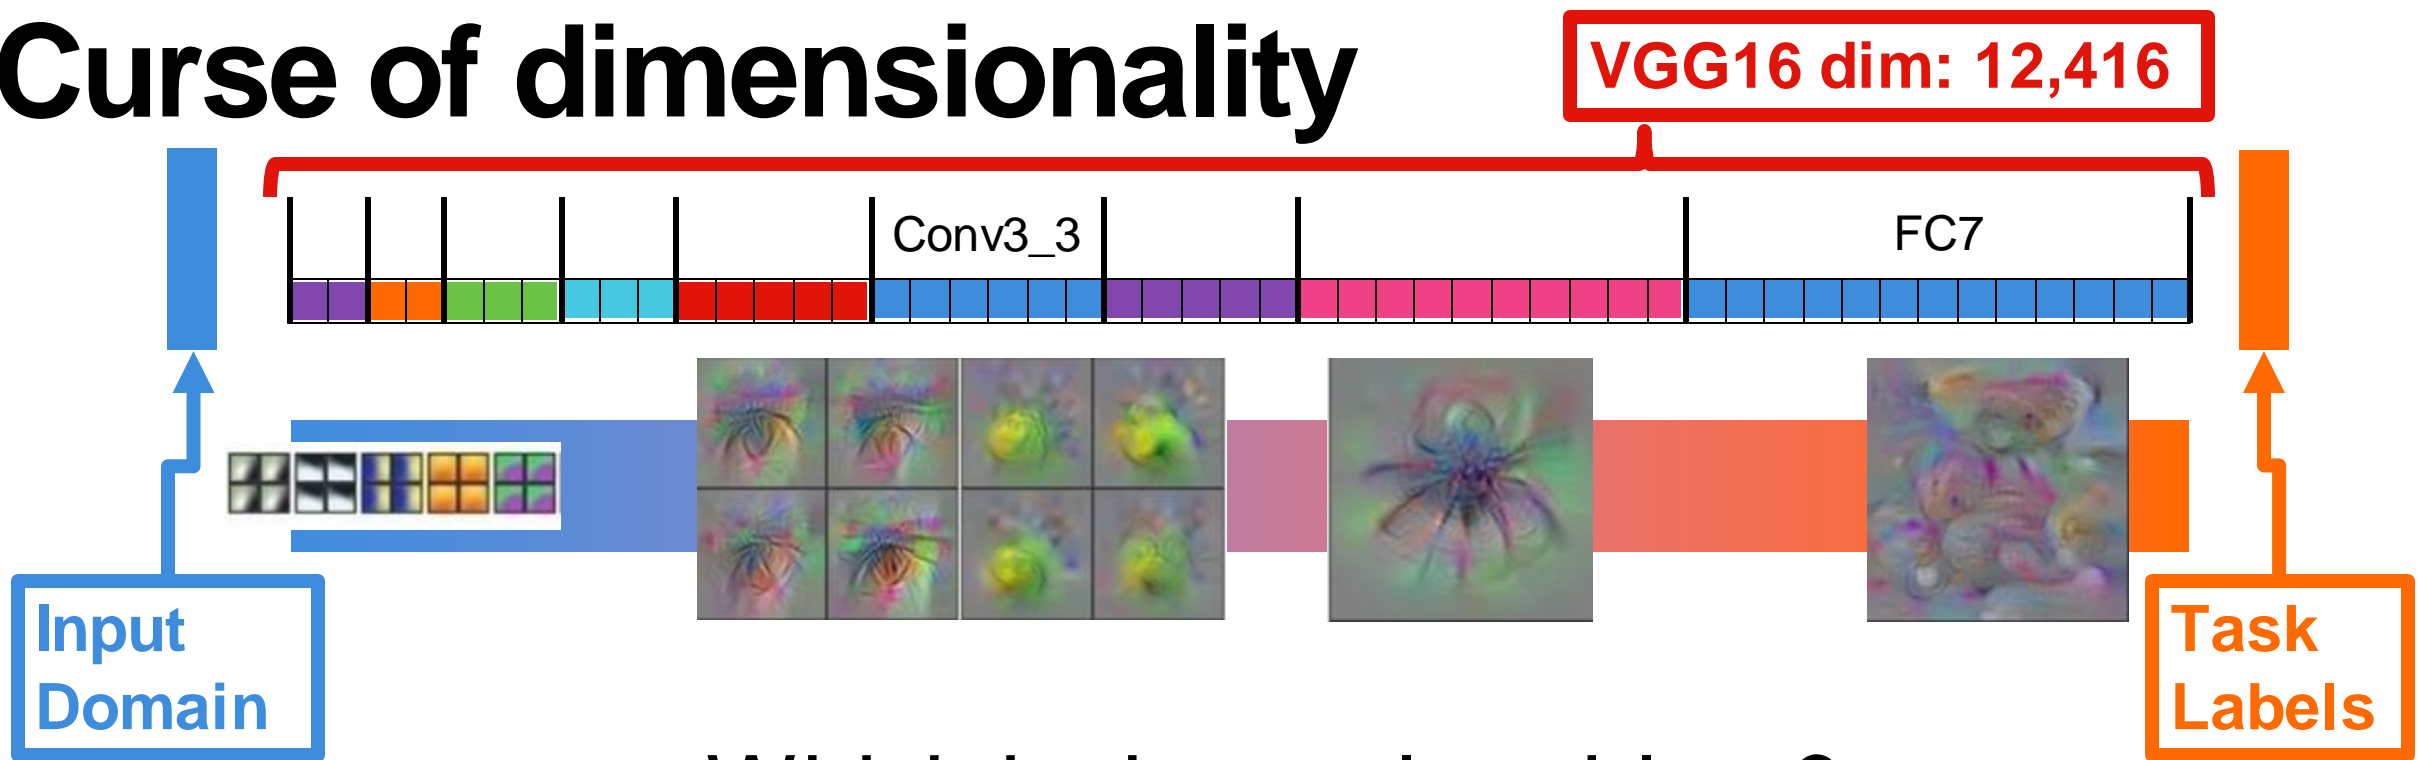
- Which is the real problem?
 - Too many features?
 - Too few images?

Curse of dimensionality



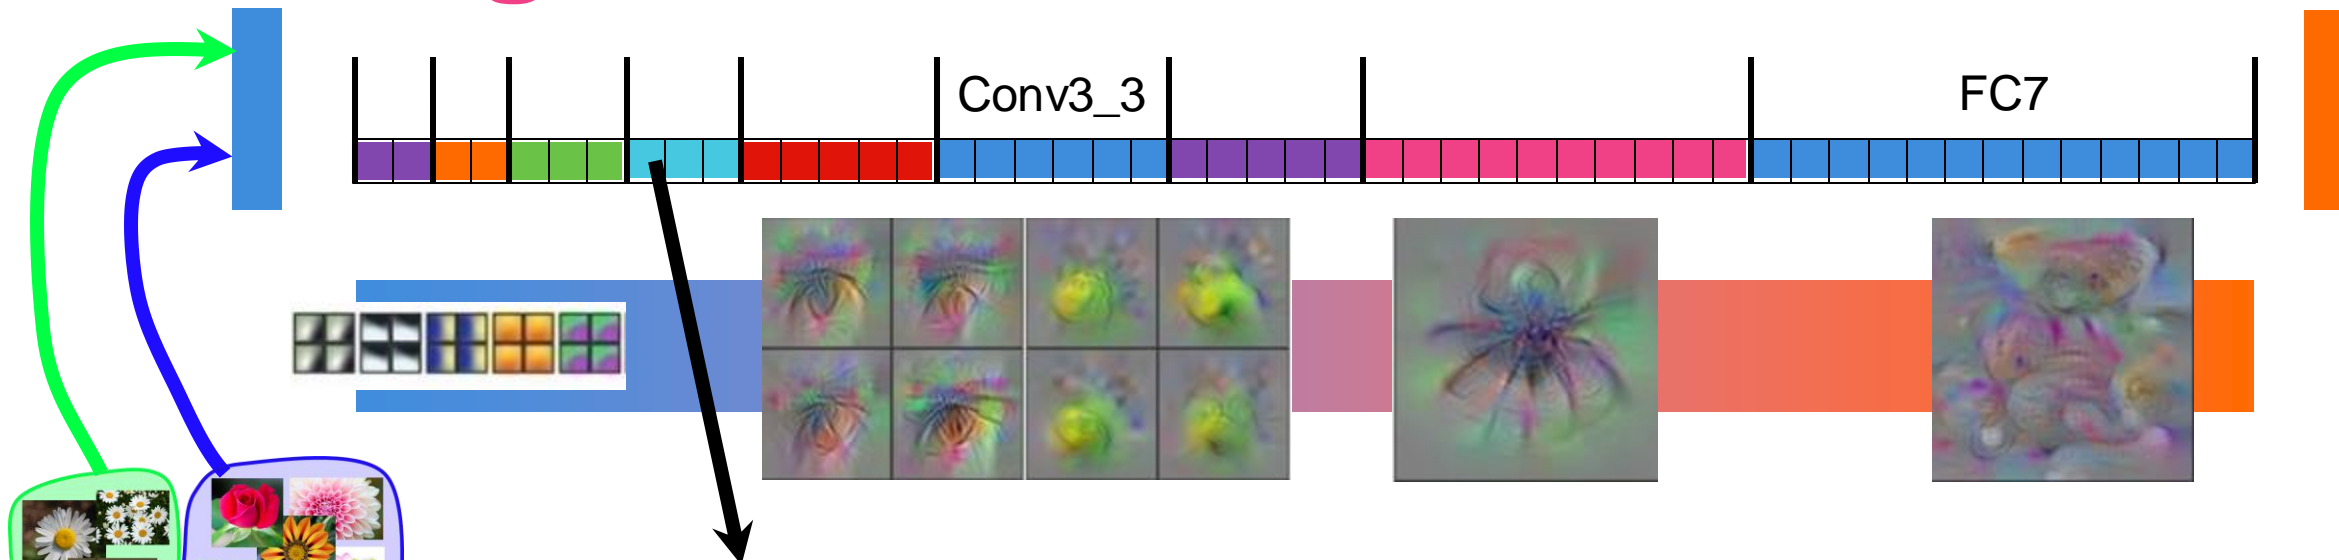
- Which is the real problem?
 - Too many features?
 - ~~Too few images?~~ **A requirement**

Curse of dimensionality





- Which is the real problem?
 - ~~Too many features?~~
 - **Too much information!**
 - ~~Too few images?~~

Meaningful discretization


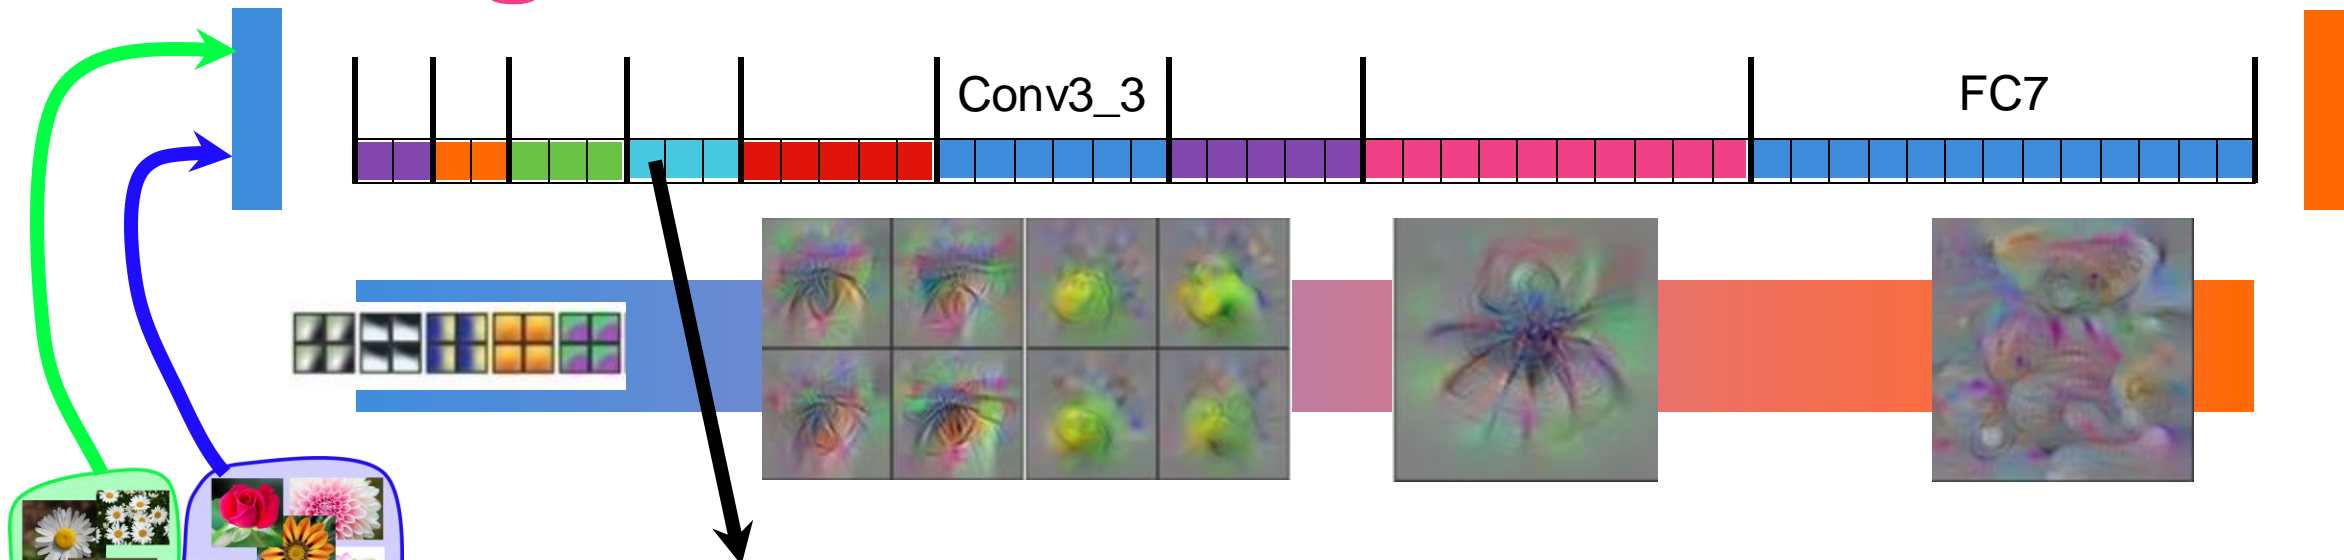


Real
values




Kolmogorov-Smirnov distance

Meaningful discretization



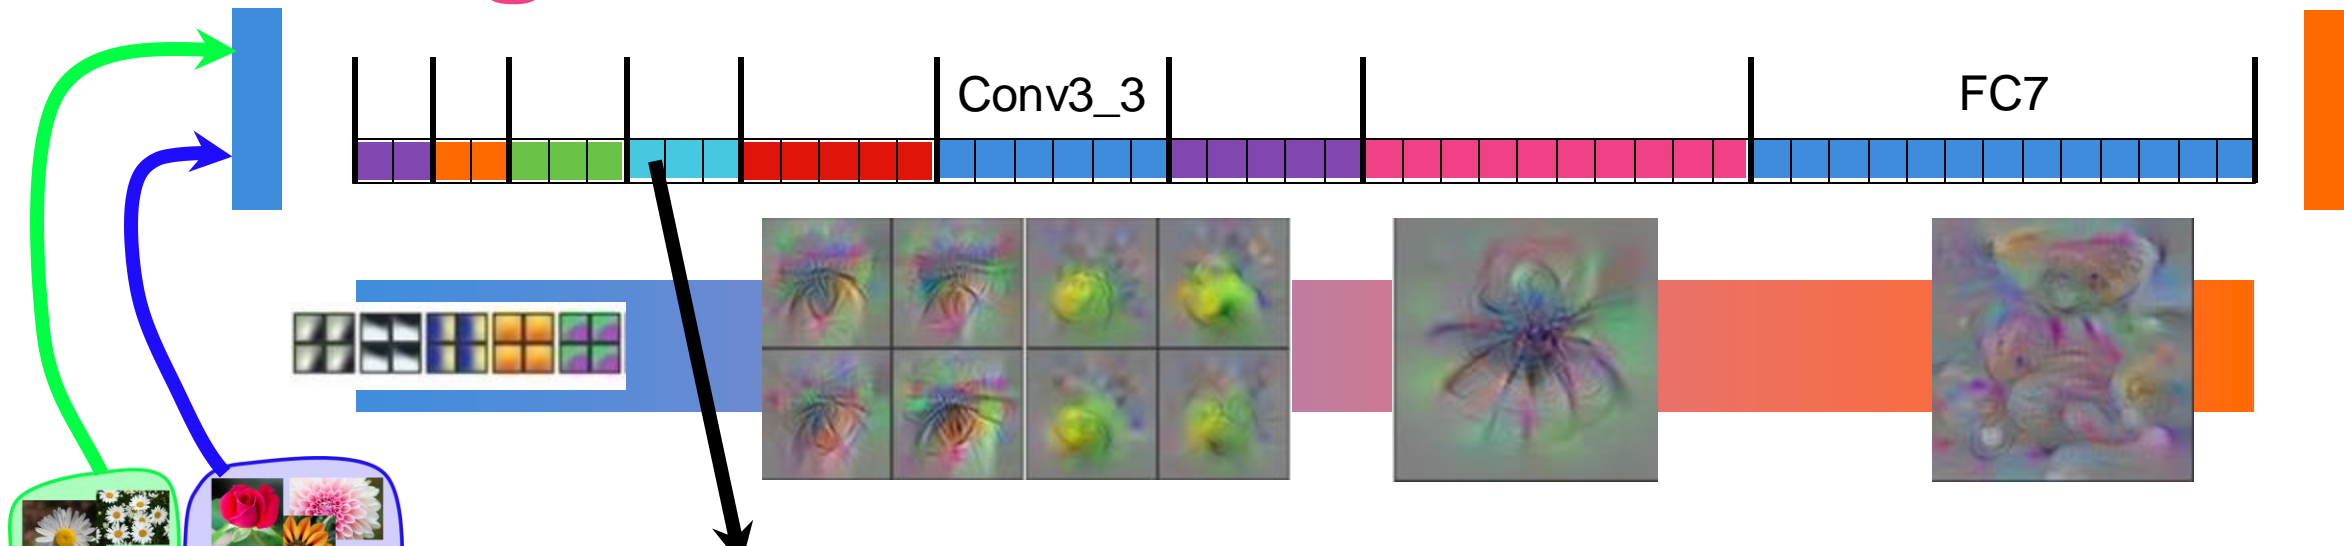
Real
values





If $D_{KS} > \text{threshold}^+$
 $\rightarrow f = 1$

Kolmogorov-Smirnov distance

Meaningful discretization




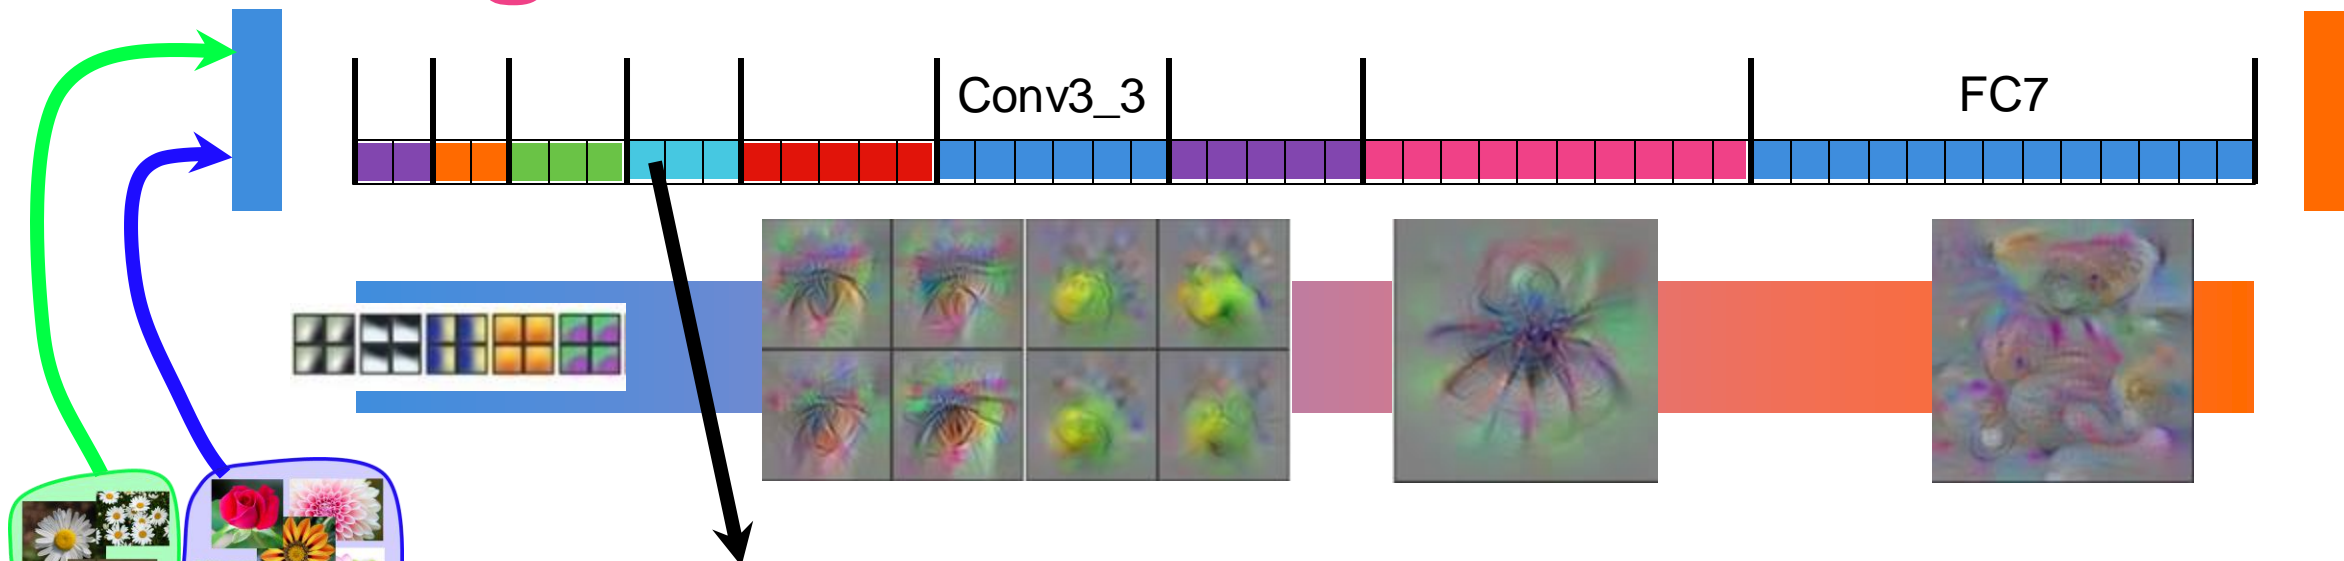
Real
values




If $D_{KS} < \text{threshold}$
 $\rightarrow f = -1$

Kolmogorov-Smirnov distance

Meaningful discretization



Real
values



If $D_{KS} \approx 0$
 $\rightarrow f = 0$


Komogrov-Smirnov distance

Finding appropriate D_{KS} thresholds

Probabilistically

Target labels

Randomized target labels



$f = 0$ ←

Feature is more likely
to behave randomly

Feature is more likely
to have meaning for
some class

→ $f = 1$


Using D_{KS} thresholds

- We have the thresholds, but...
 - They are **expensive** to compute
 - They are **class dependent. Supervised.**
- We would like to have a threshold on **feature value**, not on D_{KS} (feature-class).
- We would like the threshold to be **unsupervised**.

Using D_{KS} thresholds

- Probabilistically

- Find distributions of **feature values** conditioned to D_{KS} thresholds.
- Find the thresholds that best separate them.



mit67

Using D_{KS} thresholds

- Probabilistically

- Find distributions of feature values conditioned to D_{KS} thresholds.
- Find the thresholds that best separate them.
- Solution:

Dataset	ft^+	ft^-
mit67	0.14	-0.23
cub200	0.20	-0.24
flowers102	0.15	-0.24

Negative feature threshold = **-0.25**

Positive feature threshold = **0.15**

Using D_{KS} thresholds

- Probabilistically

- Find distributions of feature values conditioned to D_{KS} thresholds.
- Find the thresholds that best separate them.
- Solution:

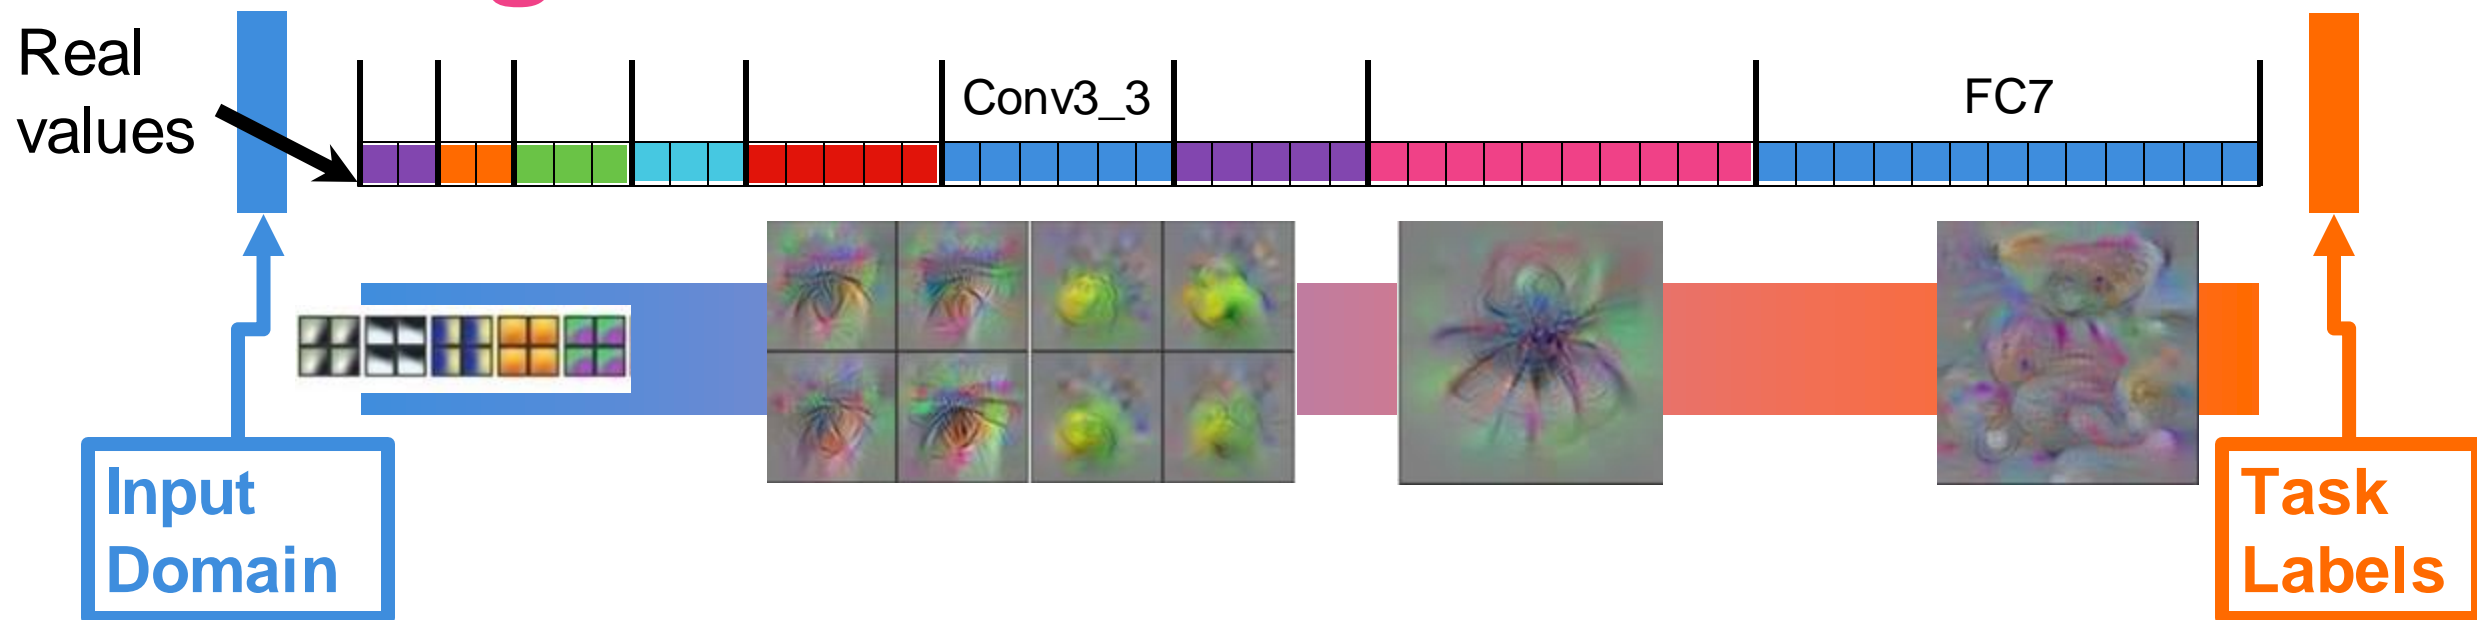
Dataset	ft^+	ft^-
mit67	0.14	-0.23
cub200	0.20	-0.24
flowers102	0.15	-0.24

Negative feature threshold = **-0.25**


Positive feature threshold = **0.15**

Good for all datasets!

Meaningful discretization



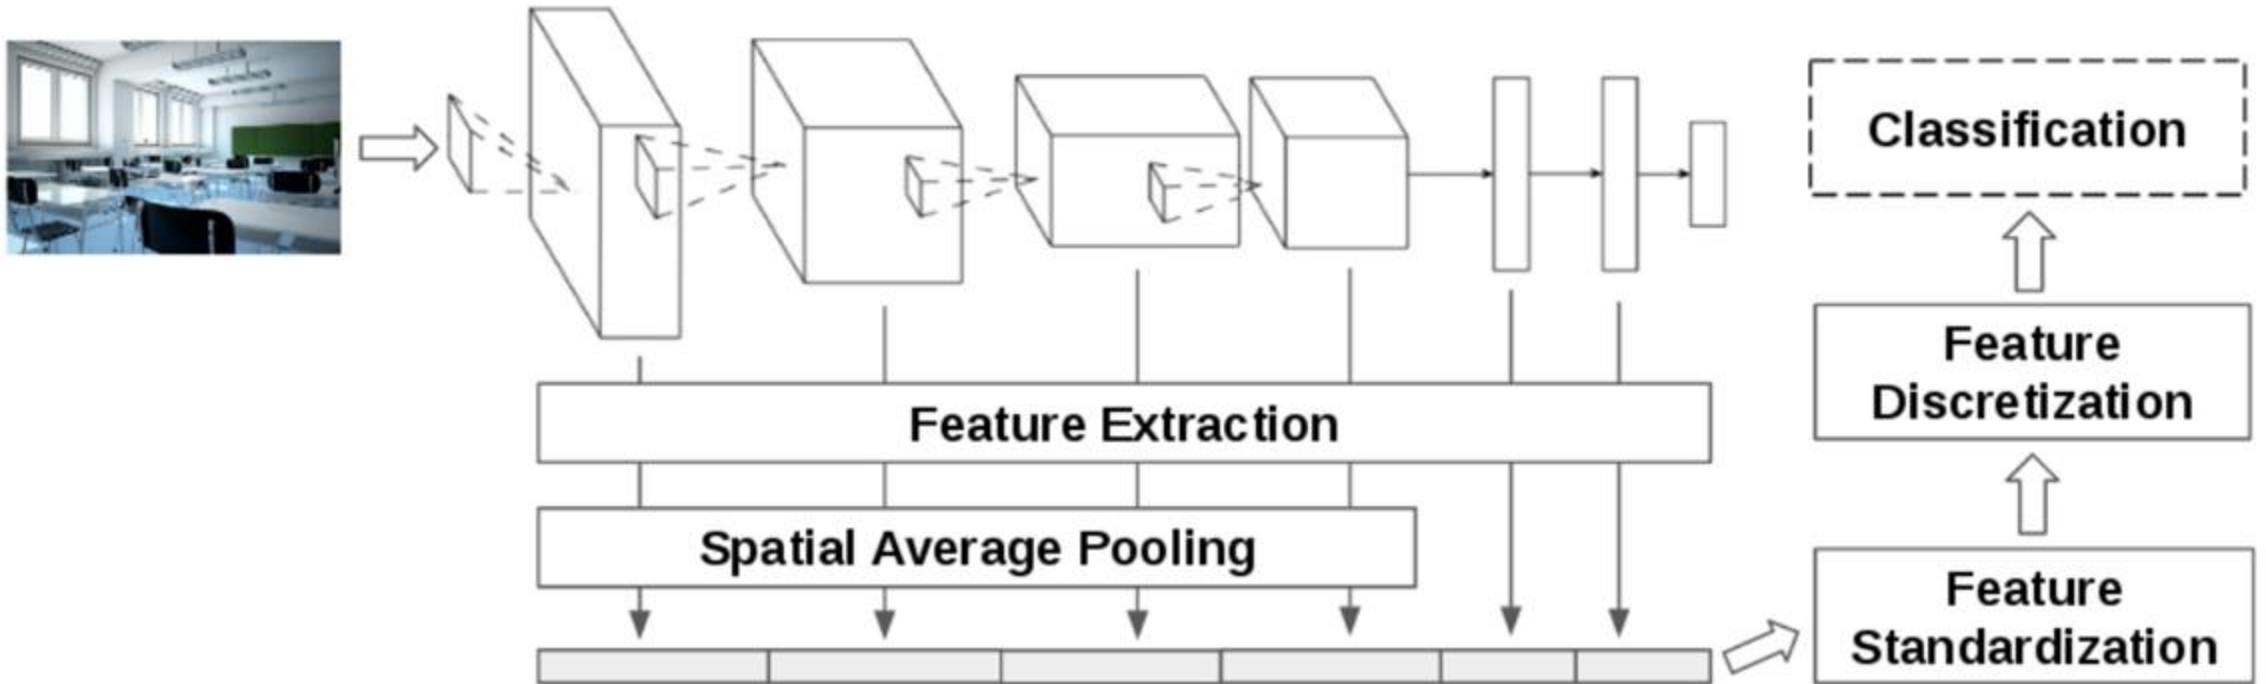
Negative feature threshold = -0.25
Positive feature threshold = 0.15



Full-Network Embedding Recipe

1. Spatial Average Pooling
2. Standardisation
3. Discretization

Full Network embedding



FNE – Small datasets

Dataset	#Images	#Classes	#Images (train)	#Images (test)	#Images per class	#Images per class (train)	#Images per class (test)
mit67	6,700	67	5,360	1,340	100	77 - 83	17 - 23
cub200	11,788	200	5,994	5,794	41 - 60	29 - 30	12 - 30
flowers102	8,189	102	2,040	6,149	40 - 258	20	20 - 238
cats-dogs	7,349	37	3,680	3,669	184 - 200	93 - 100	88 - 100
sdogs	20,580	120	12,000	8,580	150 - 200	100	50 - 100
caltech101	9,146	101	3,060	2,995	31 - 800	30	1 - 50
food101	25,250	101	20,200	5,050	250	200	50
textures	5,640	47	3,760	1,880	120	80	40
wood	438	7	350	88	14 - 179	10 - 142	3 - 37

FNE – Small datasets

Dataset	#Images	#Classes	#Images (train)	#Images (test)	#Images per class	#Images per class (train)	#Images per class (test)
mit67	6,700	67	5,360	1,340	100	77 - 83	17 - 23
cub200	11,788	200	5,994	5,794	41 - 60	29 - 30	12 - 30
flowers102	8,189	102	2,040	6,149	40 - 258	20	20 - 238
cats-dogs	7,349	37	3,680	3,669	184 - 200	93 - 100	88 - 100
sdogs	20,580	120	12,000	8,580	150 - 200	100	50 - 100
caltech101	9,146	101	3,060	2,995	31 - 800	30	1 - 50
food101	25,250	101	20,200	5,050	250	200	50
textures	5,640	47	3,760	1,880	120	80	40
wood	438	7	350	88	14 - 179	10 - 142	3 - 37

**10 – 200
images/class**

FNE - Results: Similar source task

Network pre-trained on **Places2** for mit67 and on **ImageNet** for the rest.

Dataset	mit67	cub200	flowers102	cats-dogs	sdogs	caltech101	food101	textures	wood
Baseline fc6	80.0	65.8	89.5	89.3	78.0	91.4±0.6	61.4±0.2	69.6	70.8±6.6
Baseline fc7	81.7	63.2	87.0	89.6	79.3	89.7±0.3	59.1±0.6	69.0	68.9 ±6.8
Full-network	83.6	65.5	93.3	89.2	78.8	91.4±0.6	67.0±0.7	73.0	74.1±6.9
SotA	86.9 [5]	92.3 [10]	97.0 [5]	91.6 [6]	90.3 [5]	93.4 [31]	77.4 [4]	75.5 [17]	-
ED	✓	✓	✓	✗	✓	✗	✗	✗	-
FT	✓	✓	✓	✓	✓	✓	✓	✗	-

FNE - Results: Similar source task

Network pre-trained on **Places2** for mit67 and on **ImageNet**

Best case scenario!

Dataset	mit67	cub200	flowers102	cats-dogs	sdogs	caltech101	food101	textures	wood
Baseline fc6	80.0	65.8	89.5	89.3	78.0	91.4±0.6	61.4±0.2	69.6	70.8±6.6
Baseline fc7	81.7	63.2	87.0	89.6	79.3	89.7±0.3	59.1±0.6	69.0	68.9 ±6.8
Full-network	83.6	65.5	93.3	89.2	78.8	91.4±0.6	67.0±0.7	73.0	74.1±6.9
SotA	86.9 [5]	92.3 [10]	97.0 [5]	91.6 [6]	90.3 [5]	93.4 [31]	77.4 [4]	75.5 [17]	-
ED	✓	✓	✓	✗	✓	✗	✗	✗	-
FT	✓	✓	✓	✓	✓	✓	✓	✗	-

FNE - Results: Similar source task

Network pre-trained on **Places2** for mit67 and on **ImageNet** for the rest.

Dataset	mit67	cub200	flowers102	cats-dogs	sdogs	caltech101	food101	textures	wood	
Baseline fc6	80.0	65.8	89.5	89.3	78.0	91.4±0.6	61.4±0.2	69.6	70.8±6.6	+2.9
Baseline fc7	81.7	63.2	87.0	89.6	79.3	89.7±0.3	59.1±0.6	69.0	68.9 ±6.8	+4.2
Full-network	83.6	65.5	93.3	89.2	78.8	91.4±0.6	67.0±0.7	73.0	74.1±6.9	
SotA	86.9 [5]	92.3 [10]	97.0 [5]	91.6 [6]	90.3 [5]	93.4 [31]	77.4 [4]	75.5 [17]	-	
ED	✓	✓	✓	✗	✓	✗	✗	✗	-	
FT	✓	✓	✓	✓	✓	✓	✓	✗	-	

FNE - Results: Similar source task

Network pre-trained on **Places2** for mit67 and on **ImageNet** for the rest.

Dataset	mit67	cub200	flowers102	cats-dogs	sdogs	caltech101	food101	textures	wood	
Baseline fc6	80.0	65.8	89.5	89.3	78.0	91.4±0.6	61.4±0.2	69.6	70.8±6.6	+2.9
Baseline fc7	81.7	63.2	87.0	89.6	79.3	89.7±0.3	59.1±0.6	69.0	68.9 ±6.8	+4.2
Full-network	83.6	-0.3	93.3	-0.4	-0.5	91.4±0.6	67.0±0.7	73.0	74.1±6.9	
SotA	86.9 [5]	92.3 [10]	97.0 [5]	91.6 [6]	90.3 [5]	93.4 [31]	77.4 [4]	75.5 [17]	-	
ED	✓	✓	✓	✗	✓	✗	✗	✗	-	
FT	✓	✓	✓	✓	✓	✓	✓	✗	-	

FNE - Results: Similar source task

Network pre-trained on **Places2** for mit67 and on **ImageNet** for the rest.

Dataset	mit67	cub200	flowers102	cats-dogs	sdogs	caltech101	food101	textures	wood	
Baseline fc6	80.0	65.8	89.5	89.3	78.0	91.4±0.6	61.4±0.2	69.6	70.8±6.6	+2.9
Baseline fc7	81.7	63.2	87.0	89.6	79.3	89.7±0.3	59.1±0.6	69.0	68.9 ±6.8	+4.2
Full-network	83.6	-0.3	93.3	-0.4	-0.5	91.4±0.6	67.0±0.7	73.0	74.1±6.9	
SotA	86.9 [5]	92.3 [10]	97.0 [5]	91.6 [6]	90.3 [5]	93.4 [31]	77.4 [4]	75.5 [17]	-	
ED	✓	✓	✓	✗	✓	✗	✗	✗	-	
FT	✓	✓	✓	✓	✓	✓	✓	✗	-	

FNE - Results: Dissimilar source task

Network pre-trained on **ImageNet** for mit67 and on **Places2** for the rest.

Dataset	mit67	cub200	flowers102	cats-dogs	caltech101	textures	wood
Baseline fc7	72.2	23.6	73.3	38.7	72.0	55.8	65.3
Full-network	75.5	35.5	88.7	56.2	80.0	65.1	74.0

FNE - Results: Dissimilar source task

Most frequent real-world scenario!

Network pre-trained on **ImageNet** for mit67 and on **Places2** for the rest.

Dataset	mit67	cub200	flowers102	cats-dogs	caltech101	textures	wood
Baseline fc7	72.2	23.6	73.3	38.7	72.0	55.8	65.3
Full-network	75.5	35.5	88.7	56.2	80.0	65.1	74.0



FNE - Results: Dissimilar source task

Network pre-trained on **ImageNet** for mit67 and on **Places2** for the rest.

Dataset	mit67	cub200	flowers102	cats-dogs	caltech101	textures	wood
Baseline fc7	72.2	23.6	73.3	38.7	72.0	55.8	65.3
Full-network	75.5	35.5	88.7	56.2	80.0	65.1	74.0
	+3.3	+11.9	+15.4	+17.5	+8.0	+9.3	+10.6

Simple solutions

- DNN last layer features + SVM
(Feature extraction)
 - Similar task and domain
- Add one or several NN layers +
Fine-tuning pre-trained layers
 - Enough data
- Full Network Embedding
 - Robust to different task and domain
 - Works with little data



Questions?

Dario Garcia Gasulla dario.garcia@bsc.es
Armand Vilalta