



A bit of History



A bit of History

McCulloch & Pitts / Hebb
Rosenblatt's Perceptron
Minsky & Papert - XOR
Backpropagation Algorithm

- 1943 Warren McCulloch & Walter Pitts:
 - How To: From neurons to complex thought
 - Binary threshold activations
- 1949 Howard Hebb:
 - Neurons that fire together wire together
 - Weights: Learning and memory



A bit of History

McCulloch & Pitts / Hebb

Rosenblatt's Perceptron

Minsky & Papert - XOR

Backpropagation Algorithm

[Rosenblatt, 58]

[Mark I

Perceptron]

[Perceptrons]

[1]

1948, Rosenblatt applied *Hebb's* learning to *McCulloch & Pitts* design

$$f(x) = \begin{cases} 1 & \text{if } w \cdot x + b > 0 \\ 0 & \text{otherwise} \end{cases}$$

w real-valued weights

\cdot dot product

b real scalar constant

The Mark I Perceptron. A visual classifier with:

- 400 photosensitive receptors (sensory units)
- 512 stepping motors (association units, trainable)
- 8 output neurons (response units)

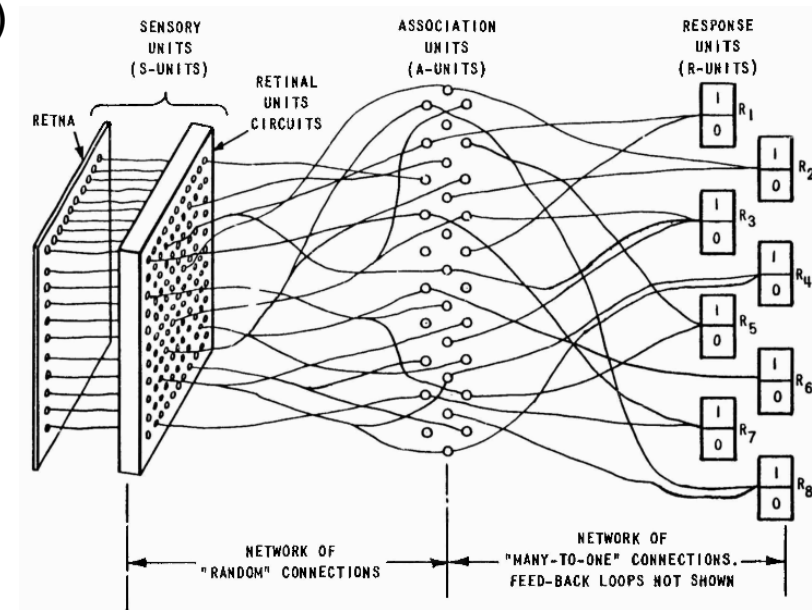


Figure 1 ORGANIZATION OF THE MARK I PERCEPTRON


A bit of History

McCullock & Pitts / Hebb
Rosenblatt's Perceptron
Minsky & Papert - XOR
Backpropagation Algorithm

Rosenblatt acknowledged a set of limitations in the Perceptron machine.

Minsky & Papert did too in “*Perceptrons: an introduction to computational geometry*”, including:

- A multilayer perceptron (MLP) is needed for learning basic functions like XOR
- MLP cannot be trained.



This had a huge impact on the public, resulting in a drastic cut in funding of NNs until the mid 80s

1st AI WINTER

A bit of History

McCullock & Pitts / Hebb
Rosenblatt's Perceptron
Minsky & Papert - XOR
Backpropagation Algorithm

How can we optimize neuron weights which are not directly connected to the error measure?

Backpropagation algorithm:

Use the chain rule to find the derivative of cost with respect to any variable.

In other words, find the contribution of each weight to the overall error.


First proposed for training MLPs by *Werbos* in '74.
Rediscovered by *Rumelhart, Hinton and Williams* in '85.

End of NNs Winter

Training with backprop

1. Forward pass from input to output
2. Error measurement (loss function)
3. Find gradients towards minimizing error layer by layer (backward pass)

Feedforward Neural Networks



Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate

Other optimization methods

Regularization

Normalizing inputs

Vanishing/Exploding Gradients

Weights initialization

Computing the gradients using all available training data would require huge amounts of memory.

Stochastic Gradient Descent: Iteratively update weights using random samples (hence, *stochastic*)

Each feedforward/backward cycle (a **step**) processes a random **batch** of images.

- Typical batch sizes: Powers of 2.
- Batch size = 1 --> Full stochastic (slower)
- Batch size = dataset_size --> Deterministic (bad generalization)

An **epoch** is the processing of the whole dataset once. It corresponds to processing as many batches as:

$$\text{dataset_size} / \text{batch_size}$$

Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate

Other optimization methods


Regularization

Normalizing inputs

Vanishing/Exploding Gradients

Weights initialization

Activation functions transform the output of a layer to a given range. If the function is non-linear, the net can learn non-linear patterns (e.g., XOR).



- Zero gradient in most of $f(x)$. Saturates!
- Max gradient is 0.25 or 1. Vanishing!
- Does not saturate
- Does not vanish
- Faster
- May die

ReLU is a safe choice in most cases

Undying alternatives: Leaky ReLU, ELU, ...

Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate

Other optimization methods

Regularization

Normalizing inputs



Vanishing/Exploding Gradients

Weights initialization

Gradient descent is a simple and straight-forward optimization algorithm to update weights towards a min.

Learning rate determines how much we move in that direction. With the wrong LR you may end up in local minima or saddle points, or be too slow.

SGD will overshoot unless we keep decreasing the LR.



Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate

Other optimization methods

Regularization

Normalizing inputs

Vanishing/Exploding Gradients

Weights initialization

Momentum: Include a fraction of the previous gradient. Keeps the general direction so far.

Nesterov: Compute current gradient considering where the previous gradient took you. (RNNs?)

Adagrad: Parameter-wise LR considering past updates. Good for infrequent patterns (GloVe). Vanishing LR due to growing history.

Adadelata: Adagrad with a decaying average over history. Typically set around 0.9.


Adam: Adadelata + Momentum

Feedforward Neural Networks

SGD, Epochs, Batches and Steps
Activation functions
SGD learning rate
Other optimization methods
Regularization
Normalizing inputs
Vanishing/Exploding Gradients
Weights initialization

Why do we need regularization?

Because the difference between **Machine Learning** and **Optimization** is called **Generalization**



Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate

Other optimization methods

Regularization

Normalizing inputs

Vanishing/Exploding Gradients

Weights initialization




Generalization

Polynomial regression

1 $h(x) = w_1x + b$

2 $h(x) = w_3x^3 + w_2x^2 + w_1x + b$

3 $h(x) = w_{14}x^{14} + w_{13}x^{13} + \dots + w_1x + b$



Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate

Other optimization methods

Regularization

Normalizing inputs

Vanishing/Exploding Gradients

Weights initialization

Generalization

Polynomial regression

Training Error

Huge

Small




Tiny

Model Generalization

Bad

Good

Horrible



Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate

Other optimization methods

Regularization

Normalizing inputs

Vanishing/Exploding Gradients

Weights initialization


Generalization

What **policy** can we use to **improve model generalization**?



Occam's Razor

when you have **two competing hypotheses** that make the **same predictions**, the **simpler one is the better**



Machine Learning

given **two models** that have a **similar performance**, It's better to **choose the simpler one**

Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate

Other optimization methods

Regularization

Normalizing inputs




Vanishing/Exploding Gradients

Weights initialization

Model Complexity

What **policy** can we use to **improve model generalization**?

$$\text{Cost function} = \text{Training Error} + \text{Model Complexity}$$



Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate

Other optimization methods

Regularization

Normalizing inputs

Vanishing/Exploding Gradients

Weights initialization

Model Complexity

$$h(x) = w_3x^3 + w_2x^2 + w_1x + w_0 \quad \text{VS} \quad h(x) = 0x^3 + 0x^2 + w_1x + w_0$$

Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate

Other optimization methods

Regularization

Normalizing inputs

Vanishing/Exploding Gradients

Weights initialization

Model Complexity

$$h(x) = w_3x^3 + w_2x^2 + w_1x + w_0 \quad \text{VS} \quad h(x) = 0x^3 + 0x^2 + w_1x + w_0$$

$$h(x) = w_3x^3 + w_2x^2 + w_1x + w_0 \quad \text{VS} \quad h(x) = 0x^3 + w_2x^2 + 0x + w_0$$

Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate

Other optimization methods

Regularization

Normalizing inputs

Vanishing/Exploding Gradients

Weights initialization

Model Complexity

$$h(x) = w_3x^3 + w_2x^2 + w_1x + w_0 \quad \text{VS} \quad h(x) = 0x^3 + 0x^2 + w_1x + w_0$$

$$h(x) = w_3x^3 + w_2x^2 + w_1x + w_0 \quad \text{VS} \quad h(x) = 0x^3 + w_2x^2 + 0x + w_0$$

$$h(x) = 0x^3 + 0x^2 + w_1x + w_0 \quad \text{VS} \quad h(x) = 0x^3 + w_2x^2 + 0x + w_0$$

Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate

Other optimization methods

Regularization

Normalizing inputs

Vanishing/Exploding Gradients

Weights initialization

Model Complexity

$$h(x) = w_3x^3 + w_2x^2 + w_1x + w_0 \quad \text{VS} \quad h(x) = 0x^3 + 0x^2 + w_1x + w_0$$

$$h(x) = w_3x^3 + w_2x^2 + w_1x + w_0 \quad \text{VS} \quad h(x) = 0x^3 + w_2x^2 + 0x + w_0$$

$$h(x) = 0x^3 + 0x^2 + w_1x + w_0 \quad \text{VS} \quad h(x) = 0x^3 + w_2x^2 + 0x + w_0$$

$$h(x) = 0x^3 + 0x^2 + w_1x + w_0 \quad \text{VS} \quad h(x) = 0x^3 + w_2x^2 + 0x + 0$$

Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate

Other optimization methods

Regularization

Normalizing inputs

Vanishing/Exploding Gradients

Weights initialization

Model Complexity

$$h(x) = w_3x^3 + w_2x^2 + w_1x + w_0 \quad \text{VS} \quad h(x) = 0x^3 + 0x^2 + w_1x + w_0$$

$$h(x) = w_3x^3 + w_2x^2 + w_1x + w_0 \quad \text{VS} \quad h(x) = 0x^3 + w_2x^2 + 0x + w_0$$

$$h(x) = 0x^3 + 0x^2 + w_1x + w_0 \quad \text{VS} \quad h(x) = 0x^3 + w_2x^2 + 0x + w_0$$

$$h(x) = 0x^3 + 0x^2 + w_1x + w_0 \quad \text{VS} \quad h(x) = 0x^3 + w_2x^2 + 0x + 0$$

? ? ?

Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate

Other optimization methods

Regularization

Normalizing inputs

Vanishing/Exploding Gradients

Weights initialization

Model Complexity

$$h(x) = 0x^3 + 0x^2 + w_1x + w_0 \quad \text{VS} \quad h(x) = 0x^3 + w_2x^2 + 0x + 0$$

ℓ_0 complexity: Number of non-zero coefficients

ℓ_1 "lasso" complexity: $\sum_{i=0}^d |w_i|$, for coefficients w_0, \dots, w_d

ℓ_2 "ridge" complexity: $\sum_{i=0}^d w_i^2$, for coefficients w_0, \dots, w_d

Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate

Other optimization methods

Regularization

Normalizing inputs

Vanishing/Exploding Gradients

Weights initialization

Model Complexity

$$h(x) = 0x^3 + 0x^2 + w_1x + w_0$$

VS

$$h(x) = 0x^3 + w_2x^2 + 0x + 0$$

$$w_0 = 1.3 \quad w_1 = -1.2$$

$$w_2 = 2.2$$

ℓ_0 complexity

$$|\{w_1, w_0\}| = 2$$

VS

$$|\{w_2\}| = 1$$

ℓ_1 complexity

$$|1.3| + |-1.2| = 2.5$$

VS

$$|2.2| = 2.2$$

ℓ_2 complexity

$$1.3^2 + (-1.2)^2 = 3.13$$

VS

$$2.2^2 = 4.84$$

Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate

Other optimization methods

Regularization

Normalizing inputs

Vanishing/Exploding Gradients

Weights initialization

L1 / L2 Regularization

$$\text{Cost function} = \text{Loss} + \frac{\lambda}{m} \sum_{i=0}^m |w_i|$$

$$\text{Cost function} = \text{Loss} + \frac{\lambda}{2m} \sum_{i=0}^m w_i^2$$

Regularization parameter $\rightarrow \lambda$

What **complexities** do these methods use?

ℓ_1 "lasso" complexity: $\sum_{i=0}^d |w_i|$, for coefficients w_0, \dots, w_d

ℓ_2 "ridge" complexity: $\sum_{i=0}^d w_i^2$, for coefficients w_0, \dots, w_d

Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate

Other optimization methods


Regularization

Normalizing inputs


Vanishing/Exploding Gradients

Weights initialization

L1 / L2 Regularization



L1 Penalty



L2 Penalty

Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate

Other optimization methods



Regularization

Normalizing inputs

Vanishing/Exploding Gradients

Weights initialization

L1 / L2 Regularization



Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate

Other optimization methods


Regularization

Normalizing inputs

Vanishing/Exploding Gradients

Weights initialization

Dropout



Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate

Other optimization methods

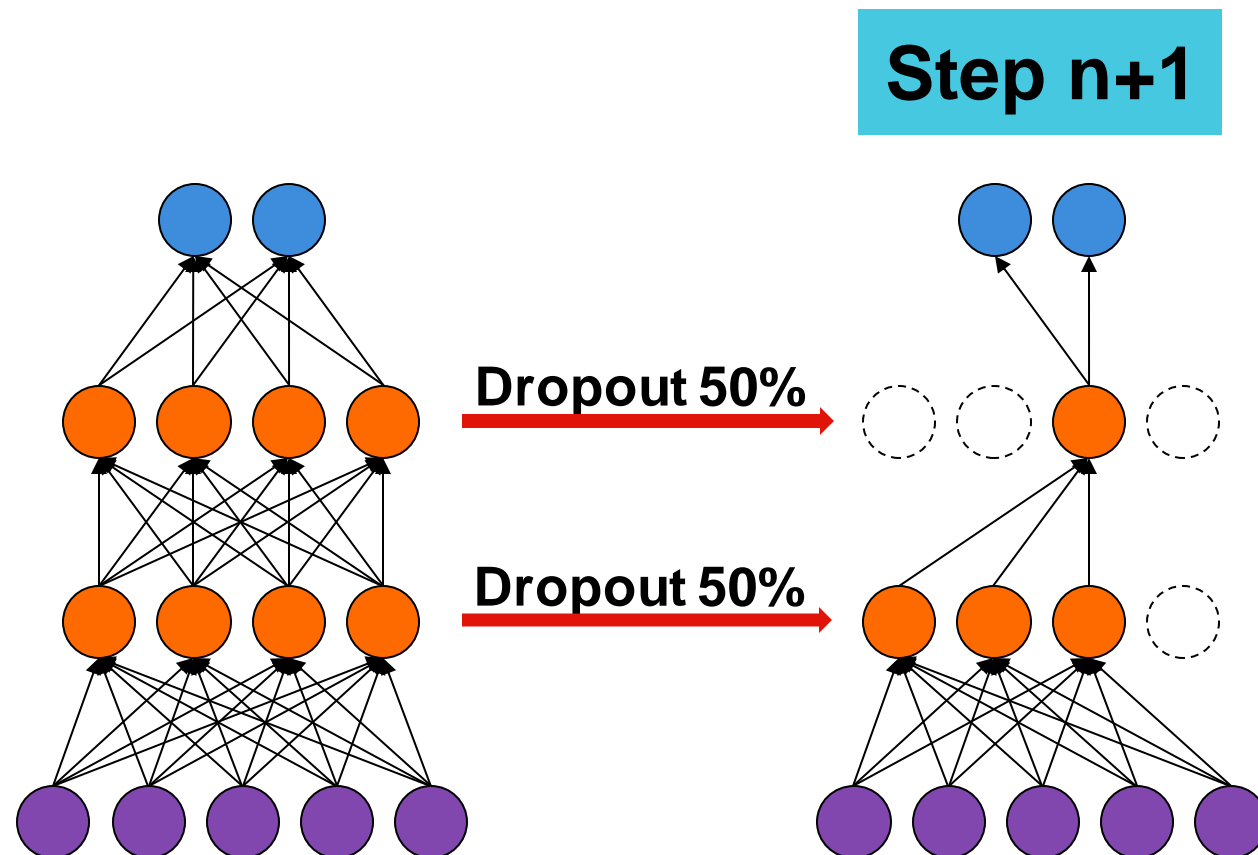
Regularization

Normalizing inputs

Vanishing/Exploding Gradients

Weights initialization

Dropout



Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate

Other optimization methods


Regularization

Normalizing inputs

Vanishing/Exploding Gradients

Weights initialization

Dropout



Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate

Other optimization methods


Regularization

Normalizing inputs

Vanishing/Exploding Gradients

Weights initialization

Dropout



Before drop-out:

$$a_0^{[0]} = g \left(w_{00}^{[0]}x_0 + w_{10}^{[0]}x_1 + w_{20}^{[0]}x_2 + w_{30}^{[0]}x_3 + w_{40}^{[0]}x_4 + b_0^{[0]} \right)$$

After drop-out: $a_0^{[0]} = 0$

What **complexity** does this method use?

ℓ_0 complexity: Number of non-zero coefficients

Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate

Other optimization methods


Regularization

Normalizing inputs

Vanishing/Exploding Gradients

Weights initialization

Early Stopping



Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate

Other optimization methods


Regularization

Normalizing inputs

Vanishing/Exploding Gradients

Weights initialization

Early Stopping



Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate

Other optimization methods


Regularization

Normalizing inputs

Vanishing/Exploding Gradients

Weights initialization

Early Stopping



What **complexity** does this method use?

ℓ_1 "lasso" complexity: $\sum_{i=0}^d |w_i|$, for coefficients w_0, \dots, w_d

ℓ_2 "ridge" complexity: $\sum_{i=0}^d w_i^2$, for coefficients w_0, \dots, w_d

Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate


Other optimization methods

Regularization

Normalizing inputs

Vanishing/Exploding Gradients

Weights initialization



$$x = \frac{x - \mu}{\sigma^2}$$

Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate


Other optimization methods

Regularization

Normalizing inputs

Vanishing/Exploding Gradients

Weights initialization



$$x = \frac{x - \mu}{\sigma^2}$$

Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate


Other optimization methods

Regularization

Normalizing inputs

Vanishing/Exploding Gradients

Weights initialization



$$x = \frac{x - \mu}{\sigma^2}$$

Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate

Other optimization methods



Regularization

Normalizing inputs

Vanishing/Exploding Gradients

Weights initialization

Why **input normalization** matters?



Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate

Other optimization methods



Regularization

Normalizing inputs

Vanishing/Exploding Gradients


Weights initialization

Why **input normalization** matters?




Feedforward Neural Networks

SGD, Epochs, Batches and Steps
Activation functions
SGD learning rate
Other optimization methods
Regularization
Normalizing inputs
Vanishing/Exploding Gradients
Weights initialization



Feedforward Neural Networks

SGD, Epochs, Batches and Steps
Activation functions
SGD learning rate
Other optimization methods
Regularization
Normalizing inputs
Vanishing/Exploding Gradients
Weights initialization



Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate


Other optimization methods

Regularization

Normalizing inputs


Vanishing/Exploding Gradients

Weights initialization



Feedforward Neural Networks

SGD, Epochs, Batches and Steps
Activation functions
SGD learning rate
Other optimization methods
Regularization
Normalizing inputs
Vanishing/Exploding Gradients
Weights initialization



Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate


Other optimization methods

Regularization

Normalizing inputs

Vanishing/Exploding Gradients

Weights initialization



Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate

Other optimization methods


Regularization

Normalizing inputs

Vanishing/Exploding Gradients

Weights initialization

Why **gradients** get **smaller and smaller** at each layer on backpropagation?



$$\nabla_{W^0} \mathcal{J}$$



$$\nabla_{W^1} \mathcal{J}$$



$$\nabla_{W^2} \mathcal{J}$$



$$\nabla_{W^3} \mathcal{J}$$



Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate

Other optimization methods


Regularization

Normalizing inputs


Vanishing/Exploding Gradients

Weights initialization

Why **gradients** get **smaller and smaller** at each layer on backpropagation?



Well, this is related to the derivatives **chain rule**.



$$\nabla_{W^0} \mathcal{J}$$



$$\nabla_{W^1} \mathcal{J}$$



$$\nabla_{W^2} \mathcal{J}$$



$$\nabla_{W^3} \mathcal{J}$$



Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate

Other optimization methods


Regularization

Normalizing inputs

Vanishing/Exploding Gradients

Weights initialization

Why **gradients** get **smaller and smaller** at each layer on backpropagation?



$$\nabla_{w_{21}^L} \mathcal{J} = \frac{\partial \mathcal{J}}{\partial w_{21}^L} = \frac{\partial \mathcal{J}}{\partial a_1^L} \times \frac{\partial a_1^L}{\partial z_1^L} \times \frac{\partial z_1^L}{\partial w_{21}^L}$$

$$\nabla_{w_{ij}^L} \mathcal{J} = a \times b \times c$$

$$\nabla_{w_{ij}^{L-1}} \mathcal{J} = a \times b \times c \times d \times e \times f$$

$$\nabla_{w_{ij}^{L-2}} \mathcal{J} = a \times b \times c \times d \times e \times f \times g \dots$$

Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate

Other optimization methods

Regularization

Normalizing inputs

Vanishing/Exploding Gradients

Weights initialization

Why **gradients** get **smaller and smaller** at each layer on backpropagation?

$$\nabla_{w_{ij}^L} \mathcal{J} = a \times b \times c$$

$$\nabla_{w_{ij}^{L-1}} \mathcal{J} = a \times b \times c \times d \times e \times f$$

$$\nabla_{w_{ij}^{L-2}} \mathcal{J} = a \times b \times c \times d \times e \times f \times g \dots$$


Terms values < 1.0

+

#Multiplying terms $\uparrow\uparrow\uparrow$

||

$$\nabla_{W^0} \mathcal{J}$$



$$\nabla_{W^1} \mathcal{J}$$



$$\nabla_{W^2} \mathcal{J}$$



$$\nabla_{W^3} \mathcal{J}$$



Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate

Other optimization methods

Regularization

Normalizing inputs

Vanishing/Exploding Gradients

Weights initialization

What about **exploding gradients**?

$$\nabla_{w_{ij}^L} \mathcal{J} = a \times b \times c$$

$$\nabla_{w_{ij}^{L-1}} \mathcal{J} = a \times b \times c \times d \times e \times f$$

$$\nabla_{w_{ij}^{L-2}} \mathcal{J} = a \times b \times c \times d \times e \times f \times g \dots$$

Terms values > 1.0

+

#Multiplying terms ↑↑↑

||

$$\nabla_{W^0} \mathcal{J}$$

↑↑↑↑

$$\nabla_{W^1} \mathcal{J}$$

↑↑↑

$$\nabla_{W^2} \mathcal{J}$$

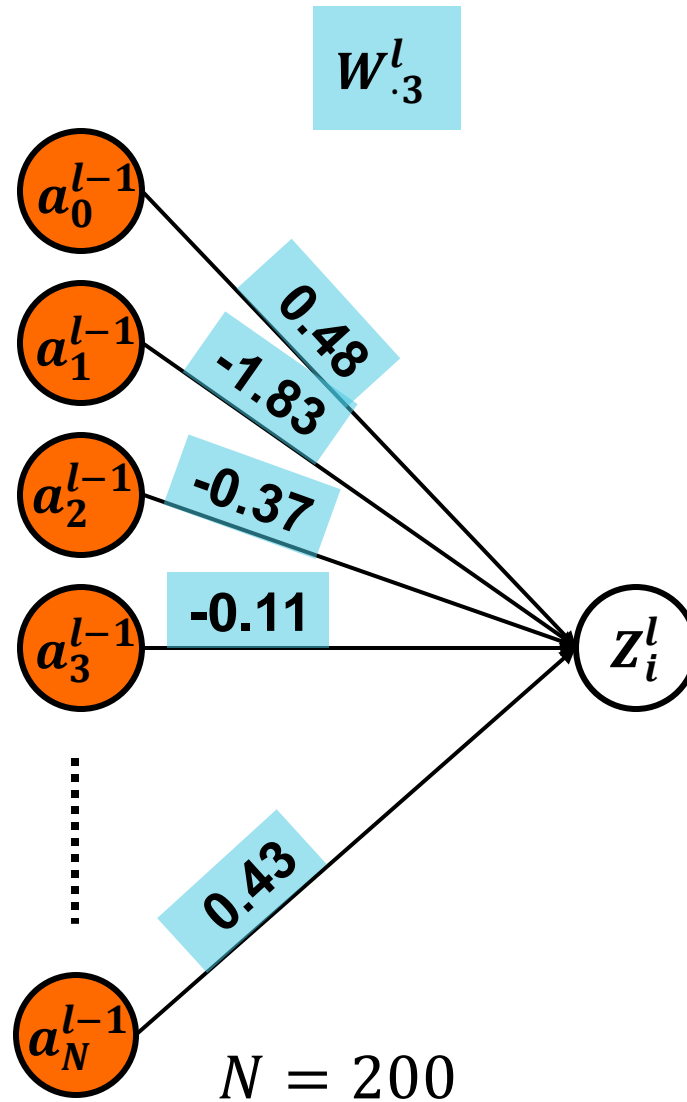
↑↑

$$\nabla_{W^3} \mathcal{J}$$


↑

Feedforward Neural Networks

SGD, Epochs, Batches and Steps
Activation functions
SGD learning rate
Other optimization methods
Regularization
Normalizing inputs
Vanishing/Exploding Gradients
Weights initialization

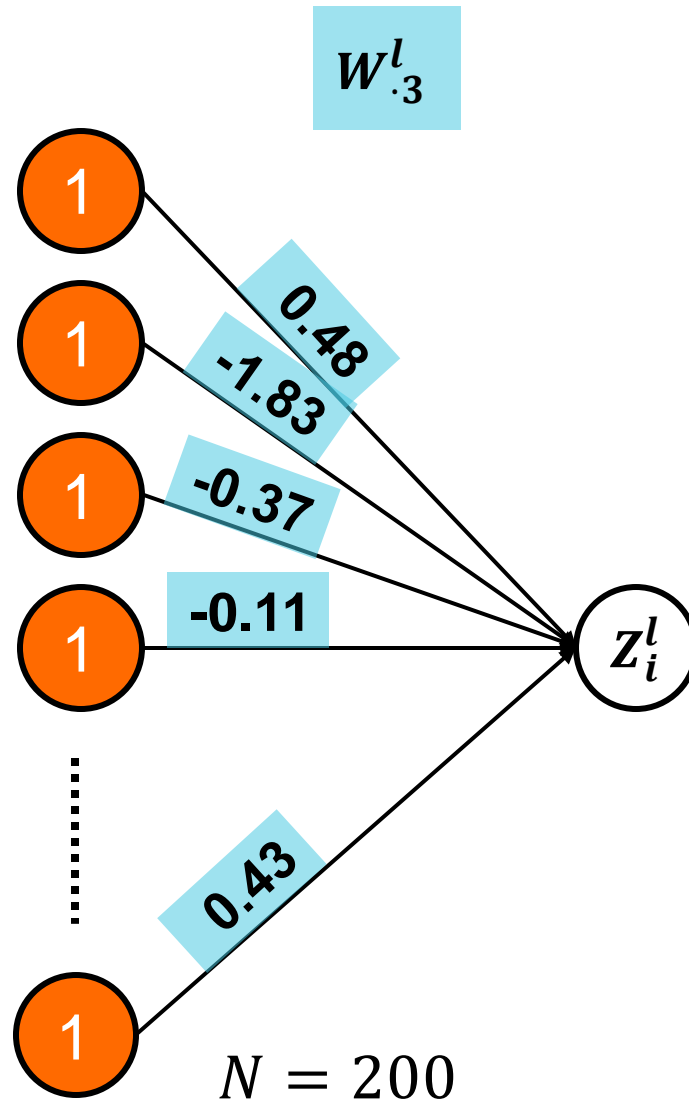


Normal Distribution
 $N(\mu = 0, \sigma^2 = 1)$




Feedforward Neural Networks

SGD, Epochs, Batches and Steps
Activation functions
SGD learning rate
Other optimization methods
Regularization
Normalizing inputs
Vanishing/Exploding Gradients
Weights initialization

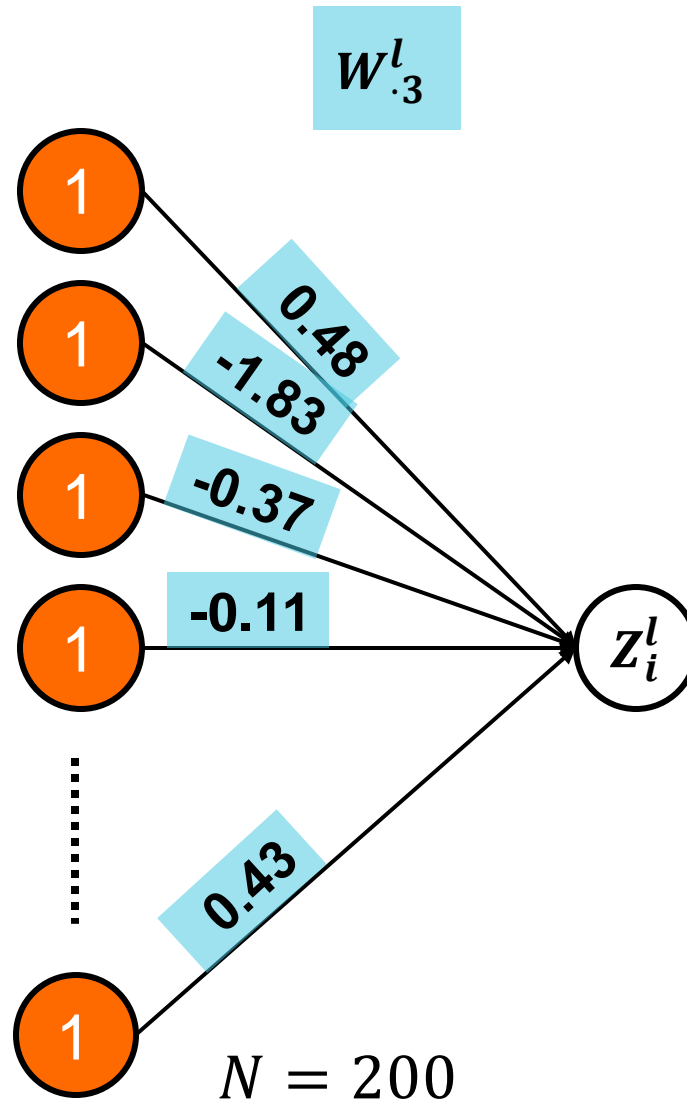


Normal Distribution
 $N(\mu = 0, \sigma^2 = 1)$



Feedforward Neural Networks

SGD, Epochs, Batches and Steps
Activation functions
SGD learning rate
Other optimization methods
Regularization
Normalizing inputs
Vanishing/Exploding Gradients
Weights initialization



Sum of **independent random variables** that are **normally distributed**:

$$X \sim N(\mu_X, \sigma_X^2)$$


$$Y \sim N(\mu_Y, \sigma_Y^2)$$

$$Z = X + Y$$

$$Z \sim N(\mu_X + \mu_Y, \sigma_X^2 + \sigma_Y^2)$$

Feedforward Neural Networks

SGD, Epochs, Batches and Steps
Activation functions
SGD learning rate
Other optimization methods
Regularization
Normalizing inputs
Vanishing/Exploding Gradients
Weights initialization



Sum of **independent random variables** that are **normally distributed**:

$$X \sim N(\mu_X, \sigma_X^2)$$

$$Y \sim N(\mu_Y, \sigma_Y^2)$$

$$Z = X + Y$$


$$Z \sim N(\mu_X + \mu_Y, \sigma_X^2 + \sigma_Y^2)$$

$$\mu(z_i^l) = 0.0$$


$$\sigma(z_i^l) = \sqrt{200} \approx 14.14$$

Feedforward Neural Networks

SGD, Epochs, Batches and Steps
Activation functions
SGD learning rate
Other optimization methods
Regularization
Normalizing inputs
Vanishing/Exploding Gradients
Weights initialization



$$\text{sigmoid}(x) = \frac{e^x}{e^x + 1}$$




$$\mu(z_i^l) = 0.0$$

$$\sigma(z_i^l) = \sqrt{200} \approx 14.14$$

Feedforward Neural Networks


SGD, Epochs, Batches and Steps
Activation functions
SGD learning rate
Other optimization methods
Regularization
Normalizing inputs
Vanishing/Exploding Gradients
Weights initialization



Alternative weights
initializations?


Feedforward Neural Networks

SGD, Epochs, Batches and Steps
Activation functions
SGD learning rate
Other optimization methods
Regularization
Normalizing inputs
Vanishing/Exploding Gradients
Weights initialization




Truncated Normal Initialization

Normal Distribution
 $N(\mu = 0, \sigma^2 = 1)$




Feedforward Neural Networks

SGD, Epochs, Batches and Steps
Activation functions
SGD learning rate
Other optimization methods
Regularization
Normalizing inputs
Vanishing/Exploding Gradients
Weights initialization



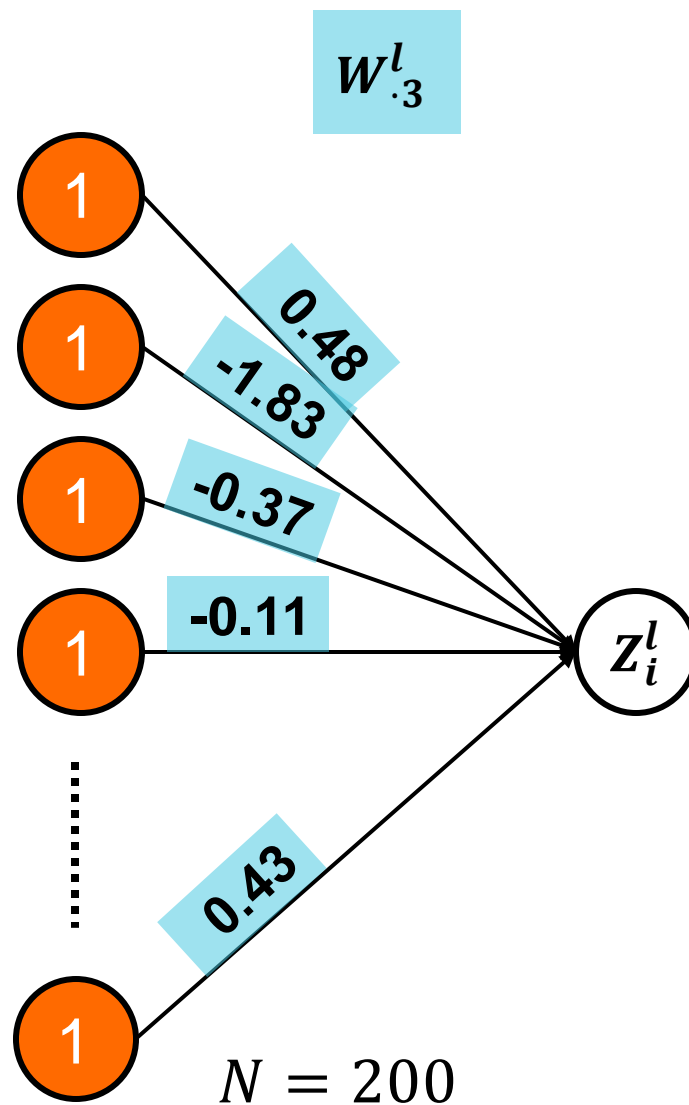
Truncated Normal Initialization

Normal Distribution
 $N(\mu = 0, \sigma^2 = 1)$




Feedforward Neural Networks

SGD, Epochs, Batches and Steps
Activation functions
SGD learning rate
Other optimization methods
Regularization
Normalizing inputs
Vanishing/Exploding Gradients
Weights initialization




Truncated Normal Initialization

Normal Distribution
 $N(\mu = 0, \sigma^2 = 1)$




Feedforward Neural Networks

SGD, Epochs, Batches and Steps
Activation functions
SGD learning rate
Other optimization methods
Regularization
Normalizing inputs
Vanishing/Exploding Gradients
Weights initialization



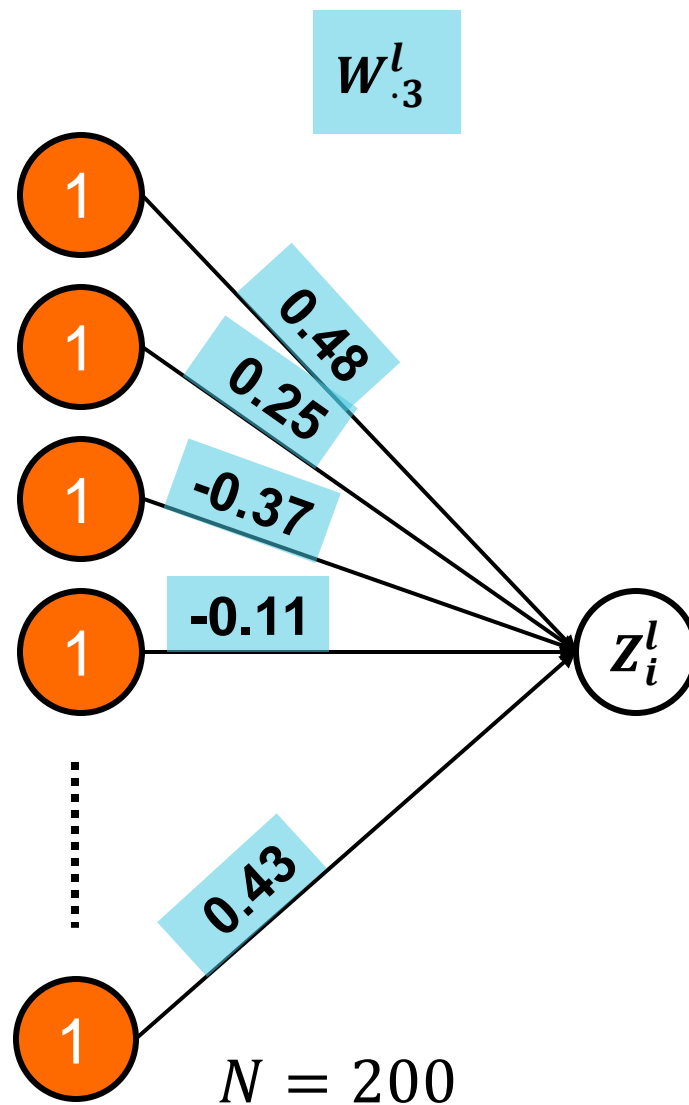
Truncated Normal Initialization

Normal Distribution
 $N(\mu = 0, \sigma^2 = 1)$




Feedforward Neural Networks

SGD, Epochs, Batches and Steps
Activation functions
SGD learning rate
Other optimization methods
Regularization
Normalizing inputs
Vanishing/Exploding Gradients
Weights initialization



Truncated Normal Initialization

Normal Distribution
 $N(\mu = 0, \sigma^2 = 1)$



Feedforward Neural Networks

SGD, Epochs, Batches and Steps

Activation functions

SGD learning rate


Other optimization methods

Regularization

Normalizing inputs

Vanishing/Exploding Gradients

Weights initialization



Xavier / Glorot Initialization

$$\mu(z_i^l) = 0.0$$


$$\sigma(z_i^l) = 1.0$$

$$\sigma^2(W_{\cdot 3}^l) = \frac{1}{N} = 0.005$$

$$\sigma(W_{\cdot 3}^l) = \sqrt{1/N} \approx 0.0707$$


Feedforward Neural Networks

SGD, Epochs, Batches and Steps
Activation functions
SGD learning rate
Other optimization methods
Regularization
Normalizing inputs
Vanishing/Exploding Gradients
Weights initialization




Xavier / Glorot Initialization

Normal Distribution
 $N(\mu = 0, \sigma^2 = 0.005)$




Practical Aspects



Practical Aspects

Train / Validation / Test Workflow



Train

Set used to **train & control bias**

Validation


Set used to **control variance**

Test

Set used to **estimate** the **generalization error** of the final model

Practical Aspects

Train / Validation / Test Workflow



Train

Set used to **train & control bias**

Validation


Set used to **control variance**

OPTIONAL

Set used to **estimate** the **generalization error** of the final model

Practical Aspects

Train / Validation / Test Workflow



Train

Set used to **train & control bias**

Test

Set used to **control variance**

OPTIONAL

Set used to **estimate** the **generalization error** of the final model

Practical Aspects

Train / Validation / Test Workflow

Machine Learning



How many **images** do we **need** for each set?

Train

As many as possible

Validation

The **minimum** amount to appropriately **represent each class**

Test

The **minimum** amount to appropriately **represent each class**

Practical Aspects

Train / Validation / Test Workflow

Machine Learning



Typical dataset size: 1.000 - 30.000

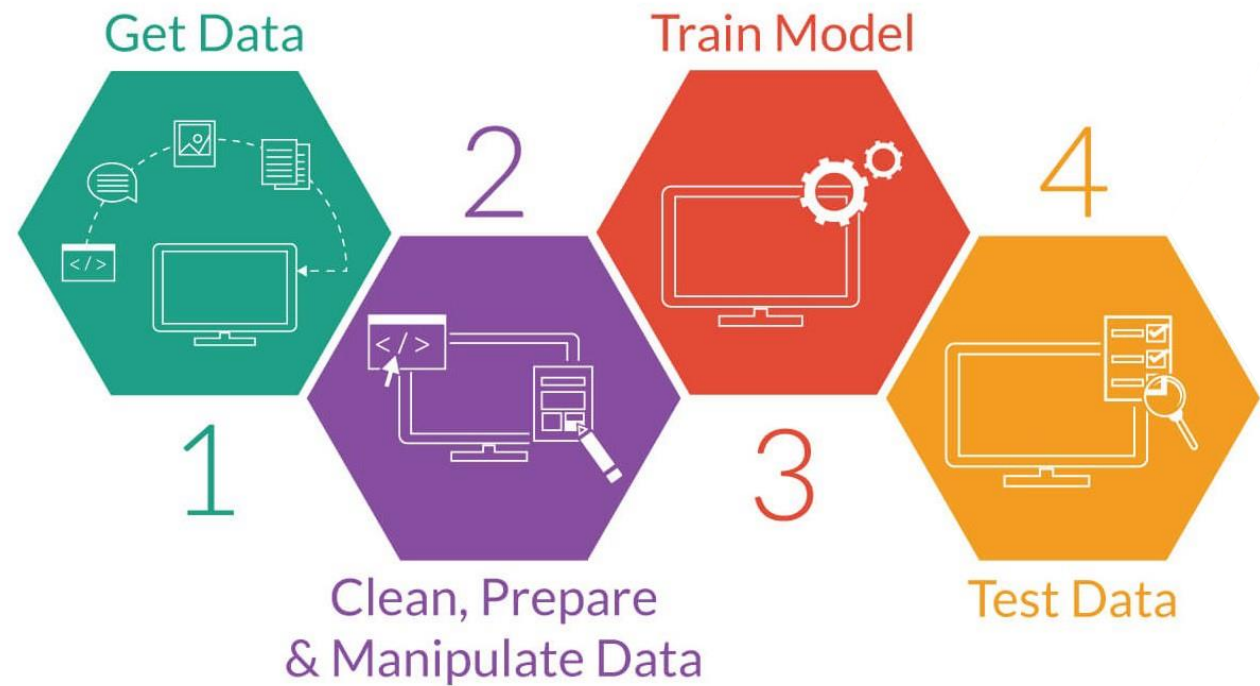
Deep Learning



Typical dataset size: 30.000 - 10.000.000

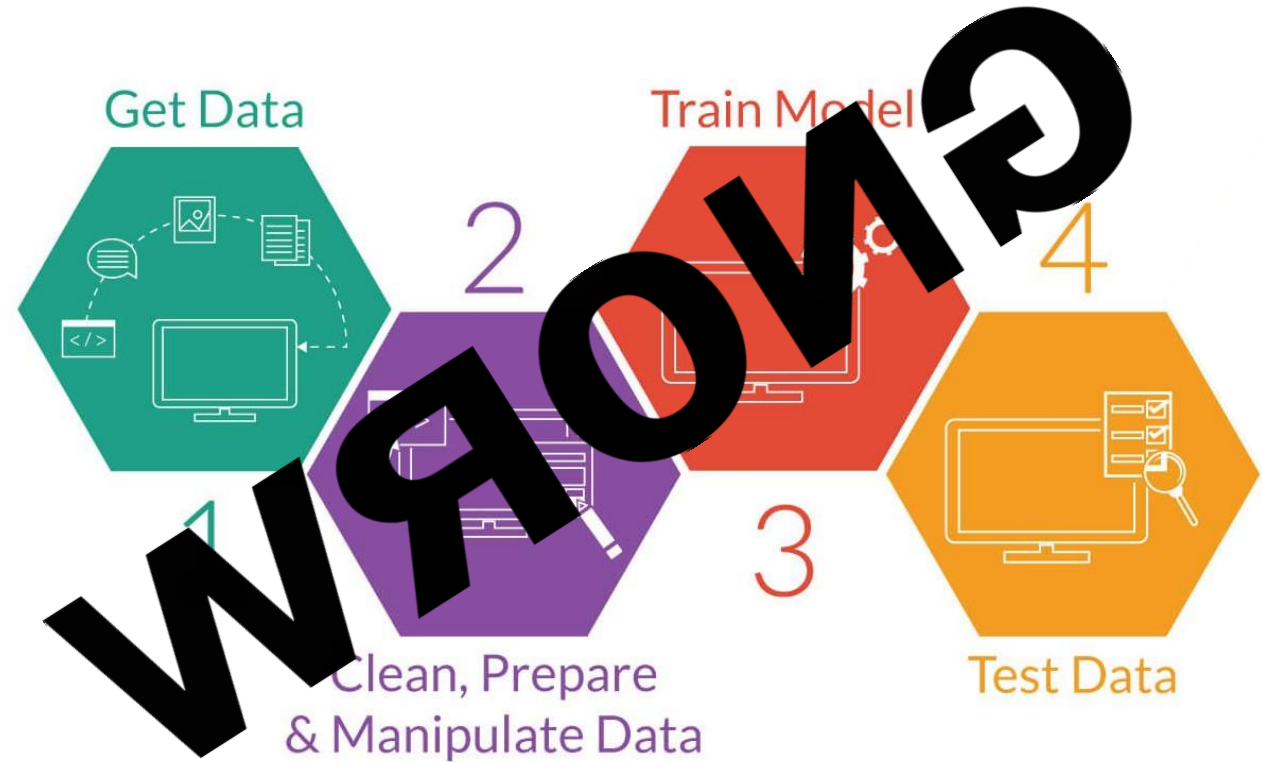
Practical Aspects

Train / Validation / Test Workflow



Practical Aspects

Train / Validation / Test
Workflow




Practical Aspects

Train / Validation / Test
Workflow

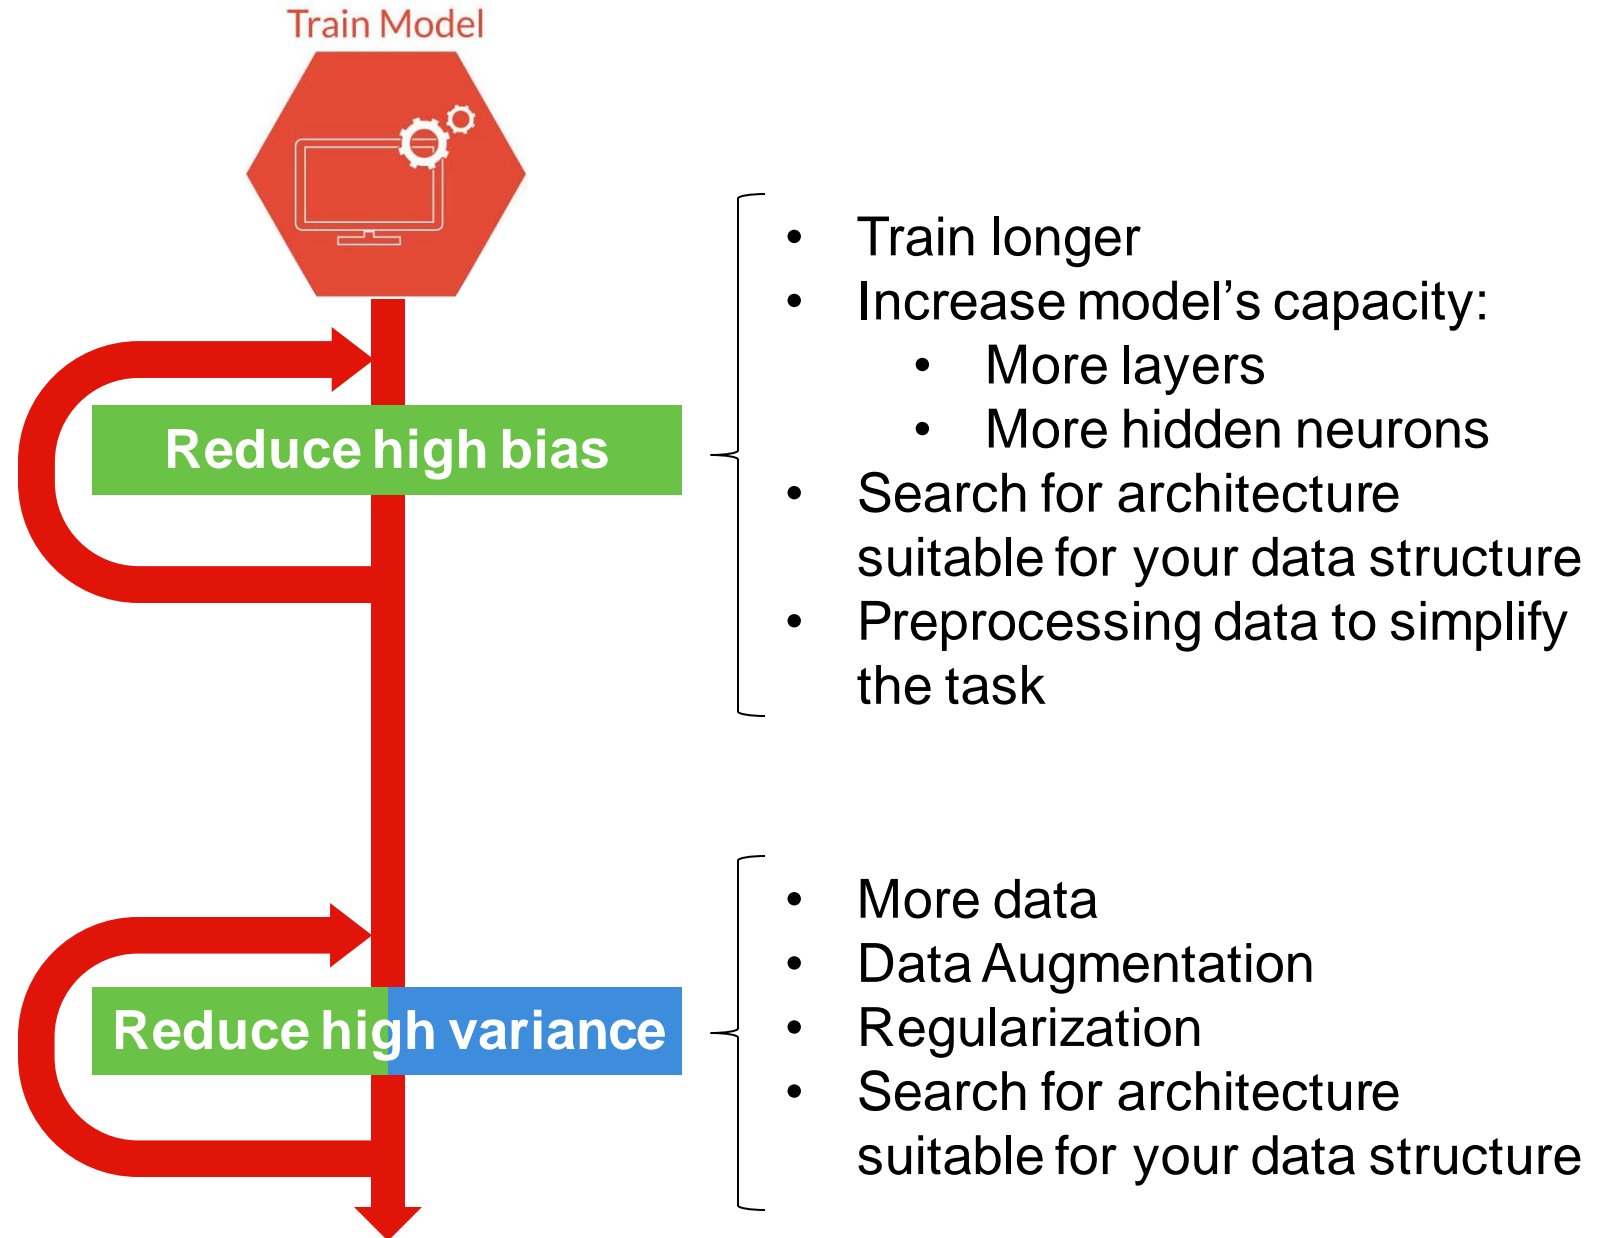
Lots of hyper-parameters:

- Network architecture:
 - # layers
 - # neurons
 - Activation function
- Learning rate
- Optimization algorithm
- ...




Practical Aspects

Train / Validation / Test
Workflow



Convolutional Neural Networks



Convolutional Neural Networks

Limited connectivity

Convolution & weight sharing
Filters

Kernel size, stride and padding

Convolutional volumes

Pooling layers

Convolutional architectures



CNNs from the inside

CNN Applications

Some data has spatial correlations that could be exploited (in 1D, 2D, 3D, ...):

- *Near-by* data points are more relevant than *far-away*.

If we sparsify connectivity with a consistent purpose, we may **reduce complexity** and ease the learning of **more coherent patterns**



Convolutional Neural Networks

Limited connectivity

Convolution & weight sharing

Filters

Kernel size, stride and padding

Convolutional volumes

Pooling layers

Convolutional architectures

CNNs from the inside

CNN Applications

Sparse connectivity is nice, but we still want to apply filters everywhere.

Each limited connectivity pattern (a **kernel**) will get **convolved** all over the image, generating a number of values.

Notice each kernel generates a 2D matrix of values.

1 <small>x1</small>	1 <small>x0</small>	1 <small>x1</small>	0	0
0 <small>x0</small>	1 <small>x1</small>	1 <small>x0</small>	1	0
0 <small>x1</small>	0 <small>x0</small>	1 <small>x1</small>	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved
Feature

In practice we have sets of neurons **sharing weights**

Convolutional Neural Networks

Limited connectivity

Convolution & weight sharing

Filters

Kernel size, stride and padding

Convolutional volumes

Pooling layers


Convolutional architectures

CNNs from the inside

CNN Applications

Convolution kernels can do all sorts of things on an image:

Input image



Edge detection

$$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$$



Sharpen

$$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$



Gaussian blur
 3×3

$$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$



Let's let the model learn them

Convolutional Neural Networks

Limited connectivity

Convolution & weight sharing

Filters

Kernel size, stride and padding

Convolutional volumes

Pooling layers


Convolutional architectures

CNNs from the inside

CNN Applications

Kernel size: Size of the receptive field of convolutional neurons. Typically 3x3, 5x5, 7x7

Stride: Number of steps while convolving filter.



Stride 1 the most common. Larger strides can replace pooling.

Padding: Border added to center conv. everywhere

- No padding: Dimensionality reduced
- Most common, zero equal/same padding

$$OutputSize = \frac{InputSize - KernelSize + 2 * Padding}{Stride} + 1$$

Convolutional Neural Networks

Limited connectivity

Convolution & weight sharing
Filters

Kernel size, stride and padding

Convolutional volumes

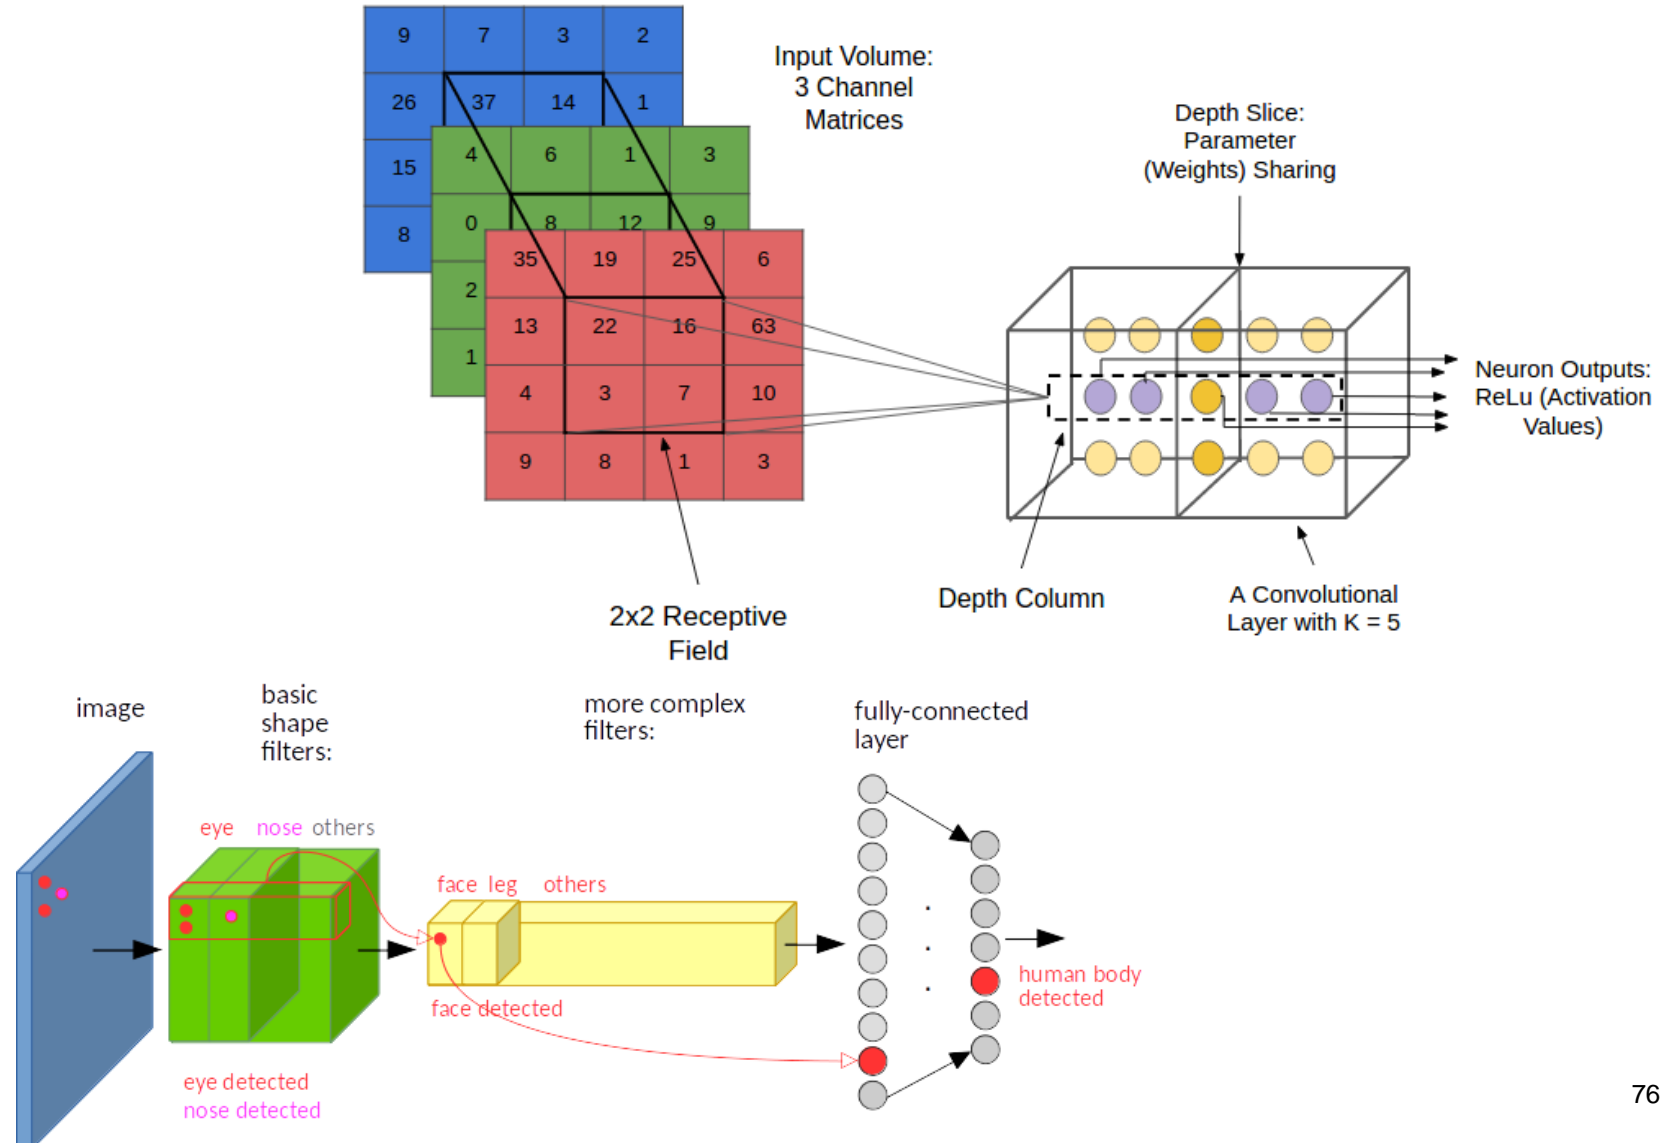
Pooling layers

Convolutional architectures

CNNs from the inside

CNN Applications

- In a typical 2D CNN, conv filters are 3D (full depth).
- Each filter convolved generates a 2D plane of data.
- Depth provides all the neural views on a part of data



Convolutional Neural Networks

Limited connectivity

Convolution & weight sharing

Filters

Kernel size, stride and padding

Convolutional volumes

Pooling layers



Convolutional architectures

CNNs from the inside

CNN Applications

Pooling:

- Small spatial invariance
- Dimensionality reduction (along x and y only)
- Never applied full depth!
- Parameter free layer
- Hyperparams:
 - Size & Stride
- Loss in precision
- Max >> Avg



Convolutional Neural Networks

Limited connectivity

Convolution & weight sharing

Filters

Kernel size, stride and padding

Convolutional volumes

Pooling layers

Convolutional architectures

CNNs from the inside

CNN Applications

The first influential architecture was **AlexNet**:


- 5 layers using convs, pools, *ReLU*, 2 dense, and *dropout*.
- 62M parameters

VGG16/19 extends the (conv-pool)*dense design:

- Smaller, 3x3 filters, but more
- 138M parameters

Some design principles: KISS, be repetitive & pyramidal

Bigger is not better!



Convolutional Neural Networks

Limited connectivity

Convolution & weight sharing

Filters

Kernel size, stride and padding

Convolutional volumes

Pooling layers

Convolutional architectures

CNNs from the inside

CNN Applications

But deeper should never be worse!

- In theory, yes. In practice, identity is hard to learn

ResNet: Learning zero is easier than learning id.

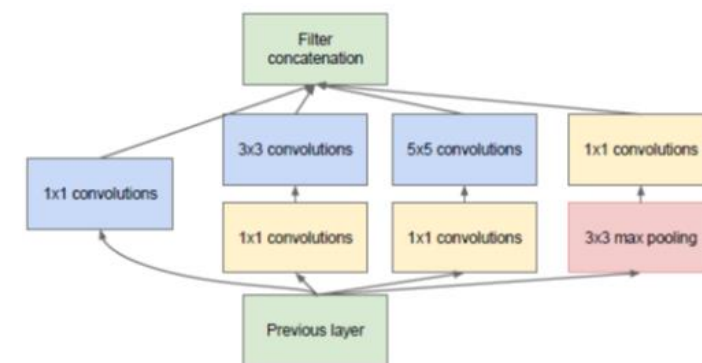



- We can now train a 1K layer net

DenseNet: link all to all

- Use depth concats
- 1x1 convs to make it feasible

Inception: how to fix filter size?

- Let the net decide which is best
- Avg. Pooling instead of dense



[Inception,15]

[ResNet,16]

[Huang,16]

[Xu,WWW]

Convolutional Neural Networks

Limited connectivity

Convolution & weight sharing

Filters

Kernel size, stride and padding

Convolutional volumes

Pooling layers


Convolutional architectures

CNNs from the inside

CNN Applications

Different architectures that can be done...

- Convolution – Transposed convolution (pixel-wise)



Convolutional Neural Networks

Limited connectivity

Convolution & weight sharing

Filters

Kernel size, stride and padding

Convolutional volumes

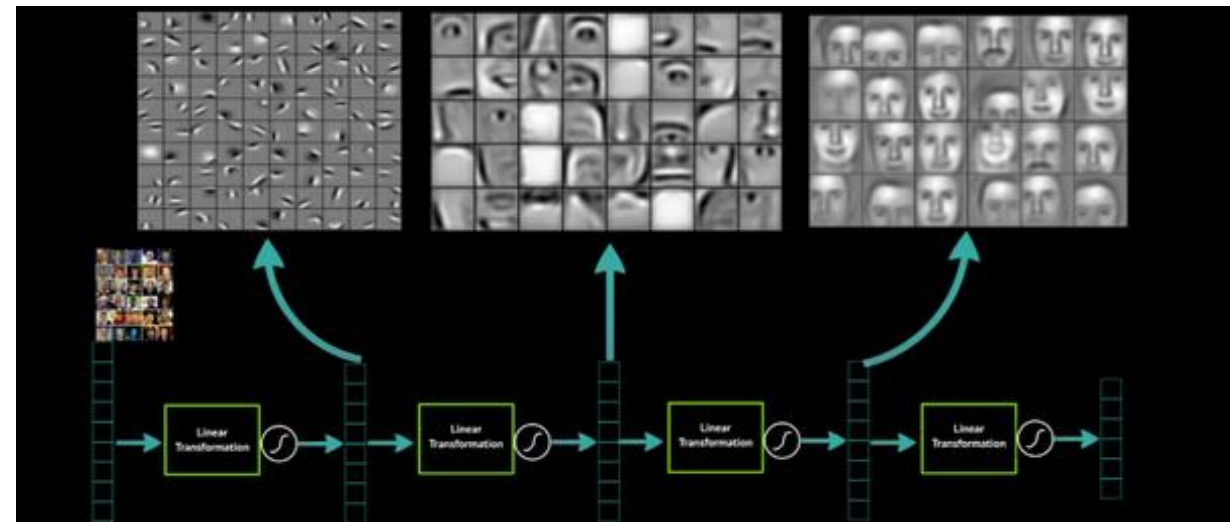


Pooling layers

Convolutional architectures

CNNs from the inside

CNN Applications

What do filters learn?



Convolutional Neural Networks

Limited connectivity

Convolution & weight sharing

Filters

Kernel size, stride and padding

Convolutional volumes

Pooling layers

Convolutional architectures

CNNs from the inside

CNN Applications

Style transfer



there is a cat sitting on a shelf .



a plate with a fork and a piece of cake .



a black and white photo of a window .



a young boy standing on a parking lot next to cars .



a wooden table and chairs arranged in a room .



a kitchen with stainless steel appliances .



this is a herd of cattle out in the field .



a car is parked in the middle of nowhere .



a ferry boat on a marina with a group of people .



a little boy with a bunch of friends on the street .



a giraffe is standing next to a fence in a field .
(hallucination)



the two birds are trying to be seen in the water .
(counting)



a parked car while driving down the road .
(contradiction)



the handlebars are trying to ride a bike rack .
(nonsensical)



a woman and a bottle of wine in a garden .
(gender)

Multimodal pipelines

Convolutional Neural Networks

Limited connectivity

Convolution & weight sharing

Filters

Kernel size, stride and padding

Convolutional volumes

Pooling layers

Convolutional architectures


CNNs from the inside

CNN Applications

Image colorization



Image segmentation





**Barcelona
Supercomputing
Center**

Centro Nacional de Supercomputación

Thanks

mail@domain.com